Lecture 7
Basics for Machine Learning and A Special Emphasis on CNN

Lin ZHANG, PhD
School of Software Engineering
Tongji University
Fall 2023

傍晚，小街路面上沁出微雨后的湿润，和煦的细风吹来，抬头看看天边的晚霞，嗯，明天又是一个好天气。走到水果摊旁，挑了个根蒂蜷缩、敲起来声音浊响的青绿西瓜，一边满心期待着皮薄肉厚瓤甜的爽落感，一边愉快地想着：这学期狠下了功夫，基础概念弄得清清楚楚，算法作业也是信手拈来，这门课成绩一定差不了！

摘自《机器学习》（周志华著，2016）

Outline

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection

A little history about AI

人工智能

1956年，**麦卡锡**召集哈佛大学、麻省理工学院、IBM公司、贝尔实验室的研究人员召开**达特茅斯会议**正式提出“**人工智能**”

2006年达特茅斯会议当事人重聚，左起：**摩尔、麦卡锡、明斯基、赛弗里奇、所罗门诺夫**

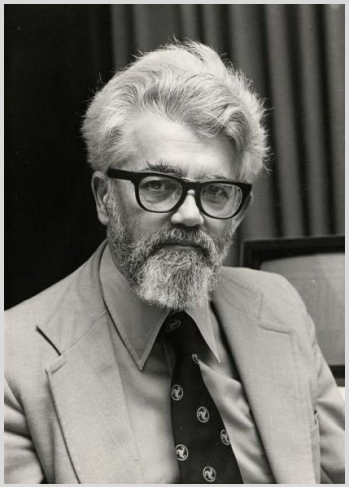

John McCarthy

人工智能之父

人工智能是指计算机系统具备的能力，该能力可以履行原本只有依靠人类智慧才能完成的复杂任务

A little history about AI

人工智能

人工智能的最终目标

探讨智能形成的基本机理，研究利用自动机模拟人的思维过程

人工智能的近期目标

研究如何使计算机去做那些靠人的智力才能做的工作

A little history about AI

人工智能的研究范式及历程

符号主义：采用知识表达和逻辑符号系统来模拟人类的智能，试图对智能进行宏观研究（**Knowledge-driven)**

联接主义：始于W.S. McCulloch和皮兹（Pitts）的先驱工作，直到目前的深度学习，是微观意义上的探索（**Data driven)**

生物启发的智能：依赖于生物学、脑科学、生命科学和心理学等学科的发现，将机理变为可计算的模型（**Biology mechanism driven)**

1950-1960	二者独立并驾齐驱
1960-1970	符号主义：专家系统和知识工程为主流
1970-1980	符号主义滞步，日本第五代计算机失败，联接主义蓬勃发展
1990-2015	联接主义占据主导；同时模糊逻辑取得重大进展
2016后	生物启发的智能一跨模态的信息处理

A little history about AI

人工智能的研究领域划分

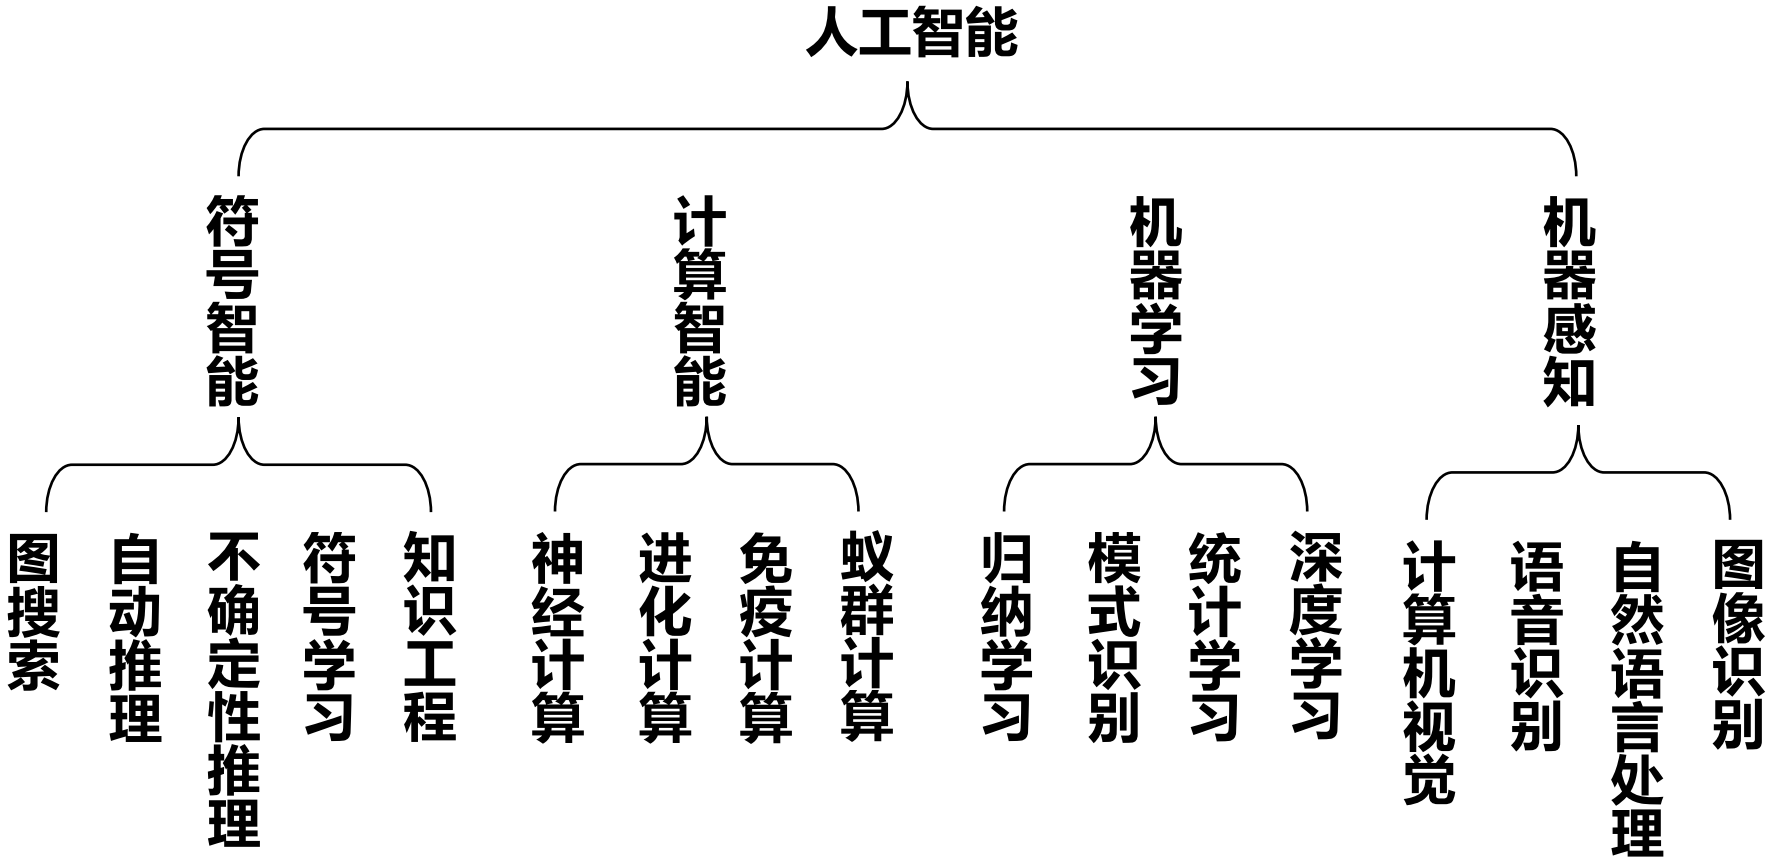

A little history about AI

人工智能产业发展加速明显

自然语言处理（NLP）：
微软Skype Translator实现同声传译

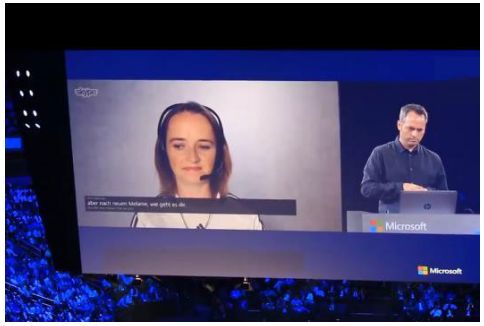

计算机视觉（CV）：
格灵深瞳的视频监控可智能识别犯罪

计算机视觉（CV）：
Face++的人脸识别云服务

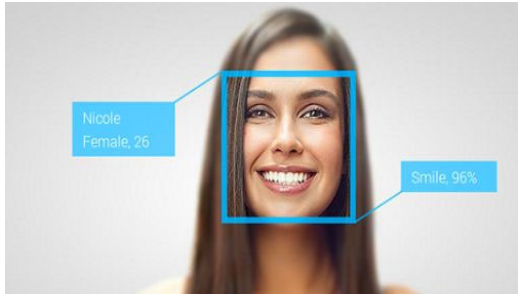

感知、规划和决策:
Google无人驾驶汽车

A little history about AI

人工智能成为世界焦点

人工智能目前已经成为世界各国关注的焦点。2017年7月，中国政府发布了“新一代人工智能发展规划”

人工智能是开启未来智能世界的秘钥，是未来科技发展的战略制高点；谁掌握人工智能，谁就将成为未来核心技术的掌控者

What is machine learning?

- Gives "computers the ability to learn without being explicitly programmed" (Arthur Samuel in 1959)

Arthur Lee Samuel
(December 5, 1901 – July 29, 1990)

- It explores the study and construction of algorithms that can learn from and make predictions on data
- It is employed in a range of computing tasks where designing and programming explicit algorithms with good performance is difficult or unfeasible

[1] Samuel, Arthur L., Some Studies in Machine Learning Using the Game of Checkers, IBM Journal of Research and Development, 1959

Supervised VS Unsupervised

- Supervised learning
  - It will infer a function from labeled training data
  - The training data consists of a set of training examples
  - Each example is a pair consisting of an input object (typically a vector) and a desired output value (also called the supervisory signal)
- Unsupervised learning
  - Trying to find hidden structure in unlabeled data
  - Since the examples given to the learner are unlabeled, there is no error or reward signal to evaluate a potential solution
  - Such as PCA, K-means (a clustering algorithm)

About sample

- Attribute (feature), attribute value, label, and example

features　　labels

{青绿，蜷缩，浊响：好瓜}

feature values　label value

one example

Training, testing, and validation

- Training sample and training set

  A training set comprising m training samples,

$$D = \{(\boldsymbol{x}_1, y_1), (\boldsymbol{x}_2, y_2), ..., (\boldsymbol{x}_m, y_m)\}$$

  where $\boldsymbol{x}_i = (x_{i1}, x_{i2}, ..., x_{id}) \in \boldsymbol{\chi}$ is the feature vector of ith sample and $y_i \in \psi$ is its label

  By training, our aim is to find a mapping,

$$f : \boldsymbol{\chi} \mapsto \psi$$

  based on D

  If ψ comprises discrete values, such a prediction task is called "**classification**"; if it comprises real numbers, such a prediction task is called "**regression**"

Training, testing, and validation

- Training sample and training set
- Test set
  - A test set is a set of data that is independent of the training data, but that follows the same probability distribution as the training data
  - Used only to assess the performance of a fully specified classifier

Training, testing, and validation

- Training sample and training set
- Test set
- Validation set
  - In order to avoid overfitting, when any classification parameter needs to be adjusted, it is necessary to have a validation set; it is used for model selection
  - The training set is used to train the candidate algorithms, while the validation set is used to compare their performances and decide which one to take

Overfitting, Generalization, and Capacity

- Overfitting
  - It occurs when a statistical model describes random error or noise instead of the underlying relationship
  - It generally occurs when a model is excessively complex, such as having too many parameters relative to the number of observations
  - A model that has been overfit will generally have poor predictive performance, as it can exaggerate minor fluctuations in the data

Overfitting, Generalization, and Capacity

- Overfitting
- Generalization
  - Refers to the performance of the learned model on new, previously unseen examples, such as the test set

Overfitting, Generalization, and Capacity

- Overfitting
- Generalization

Example: Linear regression (housing prices)

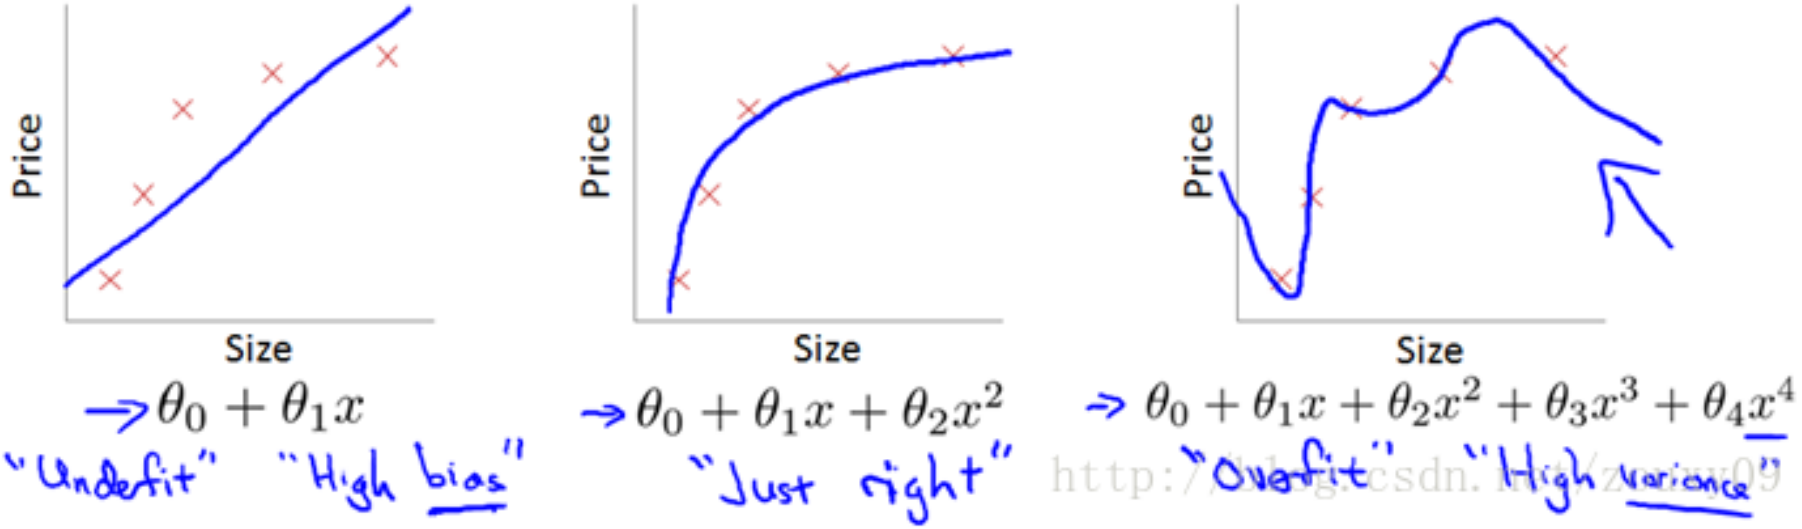

Overfitting, Generalization, and Capacity

- Overfitting
- Generalization

Example: Logistic regression

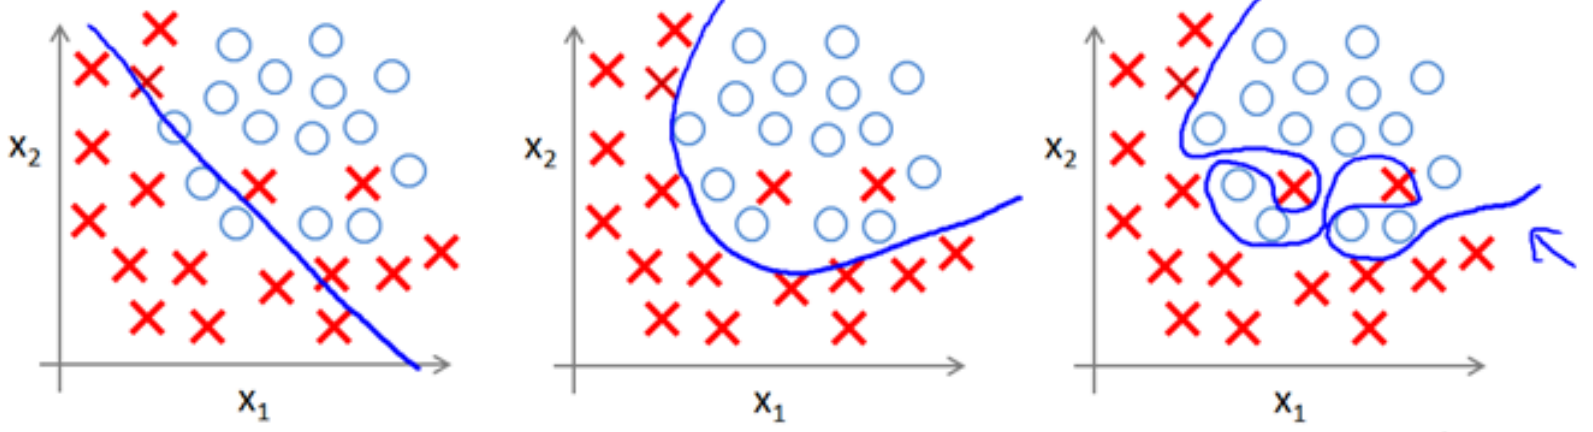

$h_\theta(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$

(g = sigmoid function)

"Underfit"

$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$

$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$

"Overfit"

Overfitting, Generalization, and Capacity

- Overfitting
- Generalization
- Capacity
  - Measures the complexity, expressive power, richness, or flexibility of a classification algorithm
  - Ex, DCNN (deep convolutional neural networks) is powerful since its capacity is very large

$$y^* = b + \omega x, \quad y^* = b + \omega_1 x_1 + \omega_2 x_2, \quad y^* = b + \sum_{i=1}^{10} \omega_i x_i$$

higher capacity

Overfitting, Generalization, and Capacity

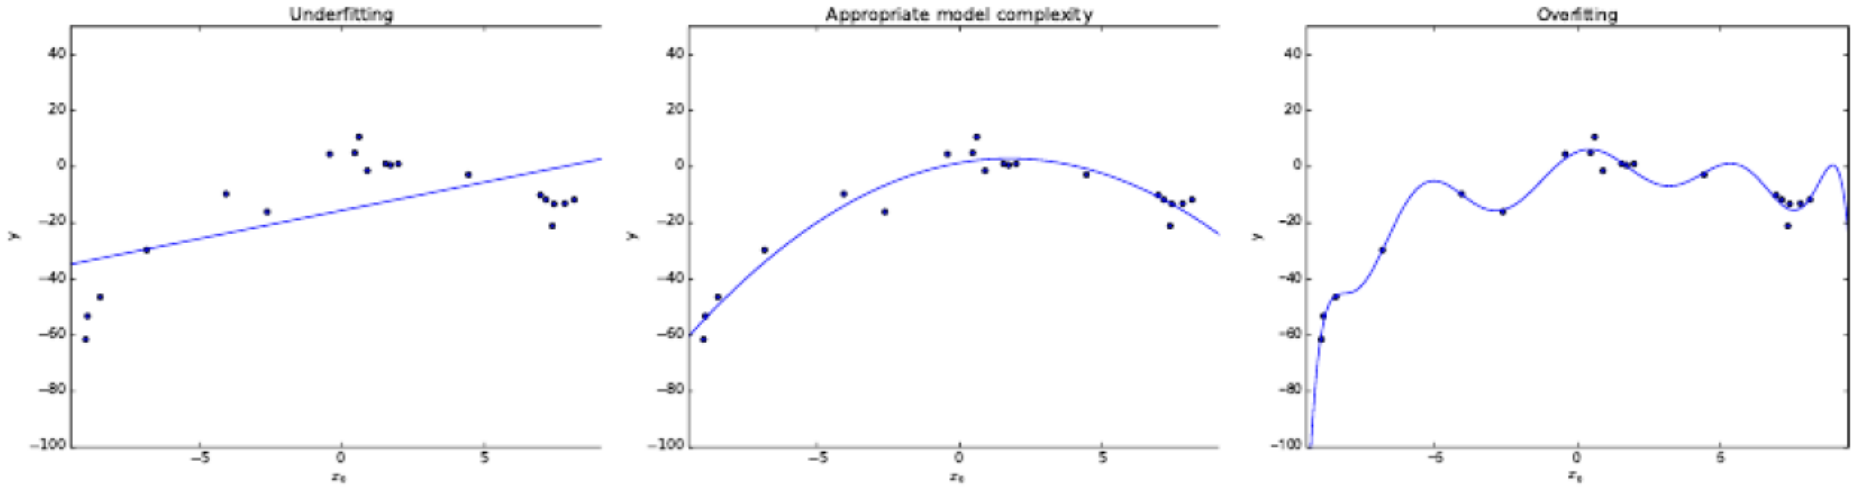

higher capacity

Performance Evaluation

Given a sample set (training, validation, or test)

$$D = \{(\boldsymbol{x}_1, y_1), (\boldsymbol{x}_2, y_2), \ldots, (\boldsymbol{x}_m, y_m)\}$$

To assess the performance of the learner f, we need to compare the prediction $f(\boldsymbol{x})$ and its ground-truth label y

For regression task, the most common performance measure is MSE (mean squared error),

$$E(f; D) = \frac{1}{m}\sum_{i=1}^{m}(f(\boldsymbol{x}_i) - y_i)^2$$

Performance Evaluation (for classification)

- Error rate
  - The ratio of the number of misclassified samples to the total number of samples

$$E(f;D)=\frac{1}{m}\sum_{i=1}^{m}\mathbf{1}\left(f(\boldsymbol{x}_i)\neq y_i\right)$$

- Accuracy
  - It is derived from the error rate

$$acc(f;D)=\frac{1}{m}\sum_{i=1}^{m}\mathbf{1}\left(f(\boldsymbol{x}_i)=y_i\right)=1-E(f;D)$$

Performance Evaluation (for classification)

- Precision and Recall

Ground truth	Prediction	
	positive	negative
positive	True Positive (TP)	False Negative (FN)
negative	False Positive (FP)	True Negative (TN)

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

Performance Evaluation (for classification)

- Precision and Recall
  - Often, there is an inverse relationship between precision and recall, where it is possible to increase one at the cost of reducing the other
  - Usually, PR-curve is not monotonic

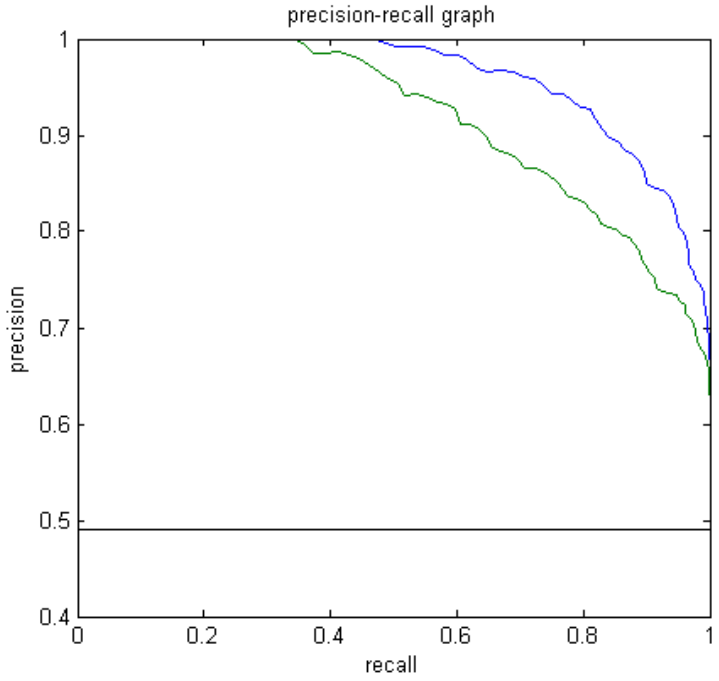

Performance Evaluation (for classification)

- Precision-recall should be used together; it is meaningless to use only one of them
- However, in many cases, people want to know explicitly which algorithm is better; we can use F-measure

$$F_\beta = \frac{(1+\beta^2)\times P \times R}{(\beta^2 \times P) + R}$$

Performance Evaluation (for classification)

- To derive a single performance measure

  Varying threshold, we can have a series of (P, R) pairs,

$$(P_1, R_1), (P_2, R_2), ..., (P_n, R_n)$$

  Then,

$$P_{macro} = \frac{1}{n}\sum_{i=1}^{n} P_i \qquad R_{macro} = \frac{1}{n}\sum_{i=1}^{n} R_i$$

$$F_{\beta-macro} = \frac{\left(1+\beta^2\right)\times P_{macro} \times R_{macro}}{\left(\beta^2 \times P_{macro}\right) + R_{macro}}$$

Class-imbalance Issue

- Problem definition
  - It is the problem in machine learning where the total number of a class of data is far less than the total number of another class of data
  - This problem is extremely common in practice
- Why is it a problem?
  - Most machine learning algorithms work best when the number of instances of each classes are roughly equal
  - When the number of instances of one class far exceeds the other, problems arise

Class-imbalance Issue

- How to deal with this issue?
  - Modify the cost function
  - Under-sampling, throwing out samples from majority classes
  - Oversampling, creating new virtual samples for minority classes
    - » Just duplicating the minority classes could lead the classifier to overfitting to a few examples
    - » Instead, use some algorithm for oversampling, such as SMOTE (synthetic minority over-sampling techniqe)[1]

[1] N.V. Chawla *et al*., SMOTE: Synthetic Minority Over-sampling Technique, J. Artificial Intelligence Research 16: 321-357, 2002

Class-imbalance Issue

- Minority oversampling by SMOTE[1]

Add new minority class instances by:

- For each minority class instance c
  - neighbours = Get KNN(5)
  - n = Random pick one from neighbours
  - Create a new minority class r instance using c's feature vector and the feature vector's difference of n and c multiplied by a random number
    - » i.e. r.feats = c.feats + (n.feats - c.feats) * rand(0,1)

[1] N.V. Chawla *et al.*, SMOTE: Synthetic Minority Over-sampling Technique, J. Artificial Intelligence Research 16: 321-357, 2002

Outline

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection

Linear regression

- Our goal in linear regression is to predict a target continuous value y from a vector of input values $\boldsymbol{x} \in R^d$; we use a linear function h as the model
- At the training stage, we aim to find $h(\boldsymbol{x})$ so that we have $h(\boldsymbol{x}_i) \approx y_i$ for each training sample $(\boldsymbol{x}_i, y_i)$
- We suppose that h is a linear function, so

$$h_{(\boldsymbol{\theta},b)}(\boldsymbol{x}) = \boldsymbol{\theta}^T \boldsymbol{x} + b, \boldsymbol{\theta} \in R^{d\times 1}$$

Rewrite it,

$$\boldsymbol{\theta}' = \begin{pmatrix} \boldsymbol{\theta} \\ b \end{pmatrix}, \boldsymbol{x}' = \begin{pmatrix} \boldsymbol{x} \\ 1 \end{pmatrix}$$

$$\boldsymbol{\theta}^T \boldsymbol{x} + b = \boldsymbol{\theta}'^T \boldsymbol{x}' \equiv h_{\boldsymbol{\theta}'}(\boldsymbol{x}')$$

Later, we simply use $h_{\boldsymbol{\theta}}(\boldsymbol{x}) = \boldsymbol{\theta}^T \boldsymbol{x}, \boldsymbol{\theta} \in R^{(d+1)\times 1}, \boldsymbol{x} \in R^{(d+1)\times 1}$

Linear regression

- Then, our task is to find a choice of $\boldsymbol{\theta}$ so that $h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)$ is as close as possible to y_i

The cost function can be written as,

$$L(\boldsymbol{\theta}) = \frac{1}{2}\sum_{i=1}^{m}\left(\boldsymbol{\theta}^T \boldsymbol{x}_i - y_i\right)^2$$

Then, the task at the training stage is to find

$$\boldsymbol{\theta}^* = \arg\min_{\boldsymbol{\theta}} \frac{1}{2}\sum_{i=1}^{m}\left(\boldsymbol{\theta}^T \boldsymbol{x}_i - y_i\right)^2$$

For this special case, it has a closed-form optimal solution

Here we use a more general method, **gradient descent** method

Linear regression

- Gradient descent
  - It is a first-order optimization algorithm
  - To find a local minimum of a function, one takes steps proportional to the negative of the gradient of the function at the current point
  - One starts with a guess $\boldsymbol{\theta}_0$ for a local minimum of $L(\boldsymbol{\theta})$ and considers the sequence such that

$$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \alpha \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta})_{|\boldsymbol{\theta}=\boldsymbol{\theta}_{n-1}}$$

where α is called as **learning rate**

Linear regression

- Gradient descent

Linear regression

- Gradient descent

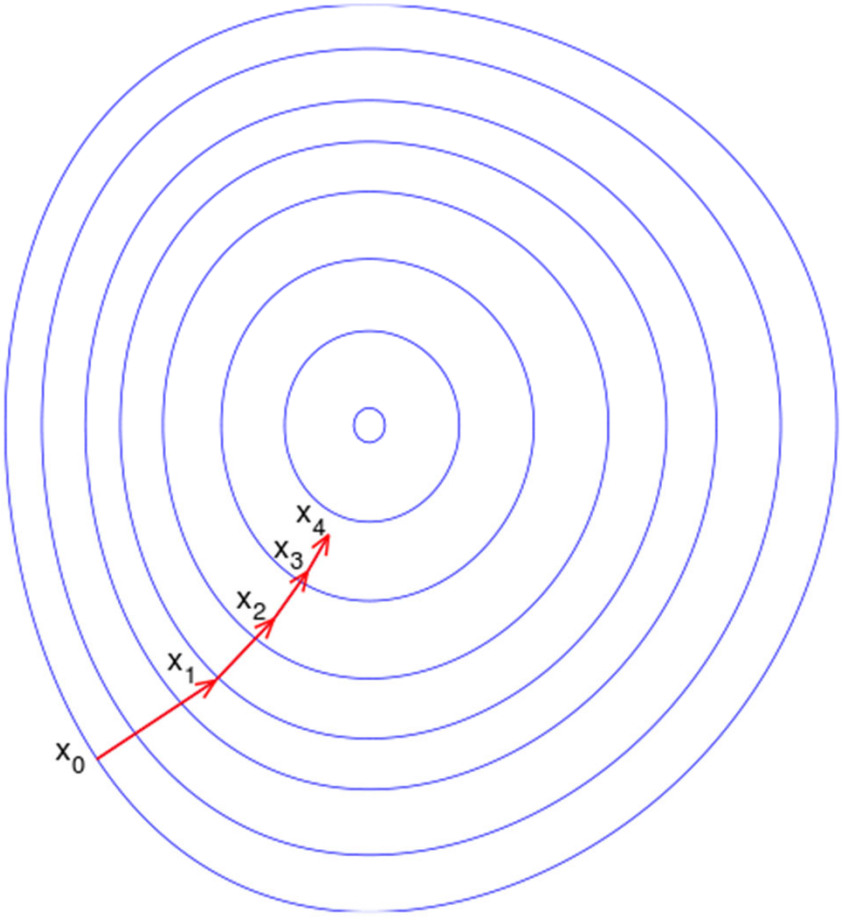

Linear regression

- Gradient descent

Repeat until convergence ($L(\boldsymbol{\theta})$ will not reduce anymore)
{

$$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \alpha \nabla_{\theta} L(\boldsymbol{\theta})_{|\theta=\theta_{n-1}}$$

}

GD is a general optimization solution; for a specific problem, the key step is how to compute gradient

Linear regression

- Gradient of the cost function of linear regression

$$L(\boldsymbol{\theta}) = \frac{1}{2}\sum_{i=1}^{m}\left(\boldsymbol{\theta}^T \boldsymbol{x}_i - y_i\right)^2$$

The gradient is,

$$\nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \begin{bmatrix} \frac{\partial L(\boldsymbol{\theta})}{\partial \theta_1} \\ \frac{\partial L(\boldsymbol{\theta})}{\partial \theta_2} \\ \vdots \\ \frac{\partial L(\boldsymbol{\theta})}{\partial \theta_{d+1}} \end{bmatrix} \quad \text{where,} \quad \frac{\partial L(\boldsymbol{\theta})}{\partial \theta_j} = \sum_{i=1}^{m}\left(h_{\boldsymbol{\theta}}(\boldsymbol{x}_i) - y_i\right)x_{ij}$$

Linear regression

- Some variants of gradient descent
  - The ordinary gradient descent algorithm looks at every sample in the **entire** training set on every step; it is also called as **batch gradient descent**
  - **Stochastic gradient descent (SGD)** repeatedly run through the training set, and each time when we encounter a training sample, we update the parameters according to the gradient of the error w.r.t that single training sample only

Repeat until convergence
{
  for i = 1 to m (m is the number of training samples)
  {
    $$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \alpha\left(\boldsymbol{\theta}_{n-1}^T \boldsymbol{x}_i - y_i\right)\boldsymbol{x}_i$$
  }
}

Linear regression

- Some variants of gradient descent
  - The ordinary gradient descent algorithm looks at every sample in the **entire** training set on every step; it is also called as **batch gradient descent**
  - **Stochastic gradient descent (SGD)** repeatedly run through the training set, and each time when we encounter a training sample, we update the parameters according to the gradient of the error w.r.t that single training sample only
  - **Minibatch SGD**: it works identically to SGD, except that it uses more than one training samples to make each estimate of the gradient

Linear regression

- More concepts
  - *m* Training samples can be divided into *N* minibatches
  - When the training sweeps all the batches, we say we complete one **epoch** of training process; for a typical training process, several epochs are usually required

epochs = 10;
numMiniBatches = *N*;
while *epochIndex*< *epochs* && *not convergent*
{
  reshuffle minibatches
  for *n* = 1 to *numMiniBatches*
  {
    //update the model parameters based on minibatch $\mathcal{B}_n$
    $\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \alpha \mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right)$
  }
}

$$\mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right) = \frac{dL\left(\mathcal{B}_n; \boldsymbol{\theta}_{n-1}\right)}{d\boldsymbol{\theta}_{n-1}}$$

Outline

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection

Logistic regression

- Logistic regression is used for binary classification
- It squeezes the linear regression $\boldsymbol{\theta}^T \boldsymbol{x}$ into the range (0, 1) ; thus the prediction result can be interpreted as probability
- At the testing stage

  The probability that the testing sample $\boldsymbol{x}$ is positive is represented as $h_{\boldsymbol{\theta}}(\boldsymbol{x}) = \dfrac{1}{1+\exp(-\boldsymbol{\theta}^T \boldsymbol{x})}$

  The probability that the testing sample $\boldsymbol{x}$ is negative is represented as $1\text{-}h_{\boldsymbol{\theta}}(\boldsymbol{x})$

Function $\sigma(z) = \dfrac{1}{1+\exp(-z)}$ is called as **sigmoid** or **logistic** function

Logistic regression

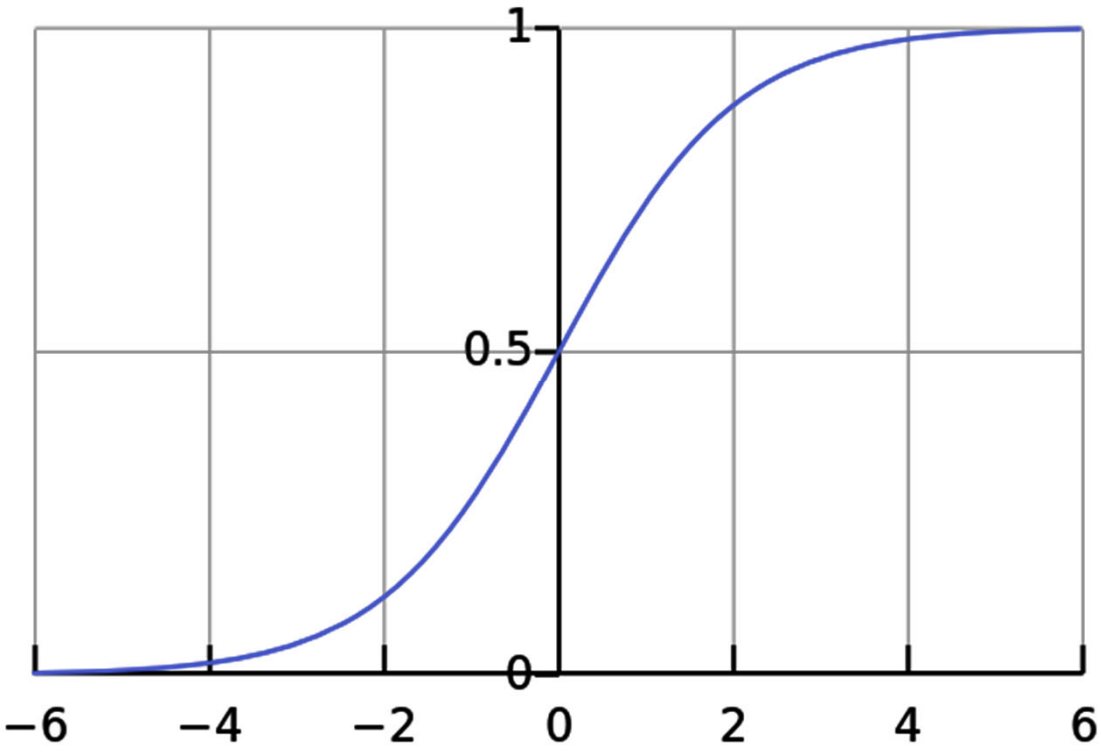

The shape of sigmoid function

One property of the sigmoid function

$$\sigma'(z) = \sigma(z)(1-\sigma(z))$$

Can you verify?

Logistic regression

- The hypothesis model can be written neatly as

$$P(y \mid \boldsymbol{x};\boldsymbol{\theta}) = \left(h_{\boldsymbol{\theta}}(\boldsymbol{x})\right)^{y} \left(1-h_{\boldsymbol{\theta}}(\boldsymbol{x})\right)^{1-y}, y \in \{0,1\}$$

- Our goal is to search for a value $\boldsymbol{\theta}$ so that $h_{\boldsymbol{\theta}}(\boldsymbol{x})$ is large when $\boldsymbol{x}$ belongs to "1" class and small when $\boldsymbol{x}$ belongs to "0" class

Thus, given a training set with binary labels $\{(\boldsymbol{x}_i, y_i) :| i = 1,\ldots,m, y_i \in \{0,1\}\}$, we want to maximize,

$$\prod_{i=1}^{m} \left(h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right)^{y_i} \left(1-h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right)^{1-y_i}$$

Equivalent to maximize,

$$\sum_{i=1}^{m} y_i \log\left(h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right) + (1-y_i)\log\left(1-h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right)$$

Logistic regression

- Thus, the cost function for the logistic regression is (we want to minimize),

$$L(\boldsymbol{\theta}) = -\sum_{i=1}^{m} y_i \log\left(h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right) + (1-y_i)\log\left(1-h_{\boldsymbol{\theta}}(\boldsymbol{x}_i)\right)$$

To solve it with gradient descent, gradient needs to be computed,

$$\nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}) = \sum_{i=1}^{m} \boldsymbol{x}_i \left(h_{\boldsymbol{\theta}}(\boldsymbol{x}_i) - y_i\right)$$

Outline

- Basic concepts
- Linear model
  - Linear regression
  - Logistic regression
  - Softmax regression
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection

Softmax regression

- Softmax operation
  - It squashes a K-dimensional vector $\mathbf{z}$ of arbitrary real values to a K-dimensional vector $\sigma(\mathbf{z})$ of real values in the range (0, 1). The function is given by,

$$\sigma(\mathbf{z})_j = \frac{\exp(\mathbf{z}_j)}{\sum_{k=1}^{K} \exp(\mathbf{z}_k)}$$

  - Since the components of the vector $\sigma(\mathbf{z})$ sum to one and are all strictly between 0 and 1, they represent a categorical probability distribution

Softmax regression

- For multiclass classification, given a test input $\boldsymbol{x}$, we want our hypothesis to estimate $p(y = k \mid \boldsymbol{x})$ for each value $k=1,2,\ldots,K$

Softmax regression

- The hypothesis should output a K-dimensional vector giving us K estimated probabilities. It takes the form,

$$h_\phi(\boldsymbol{x}) = \begin{bmatrix} p(y=1 \mid \boldsymbol{x};\phi) \\ p(y=2 \mid \boldsymbol{x};\phi) \\ \vdots \\ p(y=K \mid \boldsymbol{x};\phi) \end{bmatrix} = \frac{1}{\sum_{j=1}^{K} \exp\left((\boldsymbol{\theta}_j)^T \boldsymbol{x}\right)} \begin{bmatrix} \exp\left((\boldsymbol{\theta}_1)^T \boldsymbol{x}\right) \\ \exp\left((\boldsymbol{\theta}_2)^T \boldsymbol{x}\right) \\ \vdots \\ \exp\left((\boldsymbol{\theta}_K)^T \boldsymbol{x}\right) \end{bmatrix}$$

where $\phi = \left[\boldsymbol{\theta}_1, \boldsymbol{\theta}_2, \ldots, \boldsymbol{\theta}_K\right] \in R^{(d+1)\times K}$

Softmax regression

- In softmax regression, for each training sample we have,

$$p\left(y_i = k \mid \boldsymbol{x}_i; \phi\right) = \frac{\exp\left(\left(\boldsymbol{\theta}_k\right)^T \boldsymbol{x}_i\right)}{\sum_{j=1}^{K} \exp\left(\left(\boldsymbol{\theta}_j\right)^T \boldsymbol{x}_i\right)}$$

At the training stage, we want to maximize $p\left(y_i = k \mid \boldsymbol{x}_i; \phi\right)$ for each training sample for the correct label k

Softmax regression

- Cost function for softmax regression

$$L(\phi) = -\sum_{i=1}^{m}\sum_{k=1}^{K} 1\{y_i = k\} \log \frac{\exp\left((\boldsymbol{\theta}_k)^T \boldsymbol{x}_i\right)}{\sum_{j=1}^{K} \exp\left((\boldsymbol{\theta}_j)^T \boldsymbol{x}_i\right)}$$

where 1{.} is an indicator function

- Gradient of the cost function

$$\nabla_{\theta_k} L(\phi) = -\sum_{i=1}^{m} \left[\boldsymbol{x}_i \left(1\{y_i = k\} - p\left(y_i = k \mid \boldsymbol{x}_i; \phi\right)\right)\right]$$

Can you verify?

Cross entropy

- After the softmax operation, the output vector can be regarded as a discrete probability density function
- For multiclass classification, the ground-truth label for a training sample is usually represented in one-hot form, which can also be regarded as a density function

  For example, we have 10 classes, and the ith training sample belongs to class 7, then $y_i = [0\ 0\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 0]$

- Thus, at the training stage, we want to minimize

$$\sum_i dist(h(\boldsymbol{x}_i;\boldsymbol{\theta}), y_i)$$

How to define *dist*? Cross entroy is a common choice

Cross entropy

- Information entropy is defined as the average amount of information produced by a probabilistic stochastic source of data $H(X) = -\sum_i p(x_i) \log p(x_i)$

  where X is a random variable and x_i is the event in the sample space represented by X

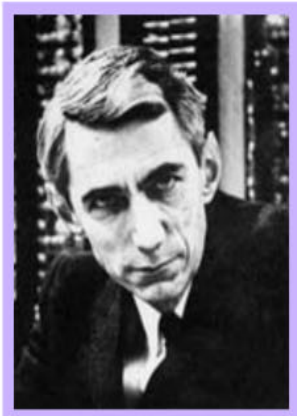

Claude Shannon
1916–2001

After graduating from Michigan and MIT, Shannon joined the AT&T Bell Telephone laboratories in 1941. His paper 'A Mathematical Theory of Communication' published in the *Bell System Technical Journal* in 1948 laid the foundations for modern information theory. This paper introduced the word 'bit', and his concept that information could be sent as a stream of 1s and 0s paved the way for the communications revolution. It is said that von Neumann recommended to Shannon that he use the term entropy, not only because of its similarity to the quantity used in physics, but also because "nobody knows what entropy really is, so in any discussion you will always have an advantage".

Cross entropy

- Information entropy is defined as the average amount of information produced by a probabilistic stochastic source of data $H(X) = -\sum_i p(x_i) \log p(x_i)$
- Cross entropy can measure the difference between two distributions

$$H(p, q) = -\sum_i p(x_i) \log q(x_i)$$

where p is the ground-truth, q is the prediction result, and x_i is the class index

- For multiclass classification, the last layer usually is a softmax layer and the loss is the 'cross entropy'

Cross entropy

Example:

Suppose that the label of one sample is [1 0 0]

For model 1, the output of the last softmax layer is [0.5 0.4 0.1]

Its cross entropy is (base 10),

$$H(p,q) = -\sum_i p(x_i)\log q(x_i) = -\left(1*\log 0.5 + 0*\log 0.4 + 0*\log 0.1\right) \approx 0.3$$

For model 2, the output of the last softmax layer is [0.8 0.1 0.1]

$$H(p,q) = -\sum_i p(x_i)\log q(x_i) = -\left(1*\log 0.8 + 0*\log 0.1 + 0*\log 0.1\right) \approx 0.1$$

Model 2 is better

Outline

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- DCNN for object detection

Neural networks

- Neural networks are proposed by Walter Harry Pitts and Warren Sturgis McCulloch in 1943

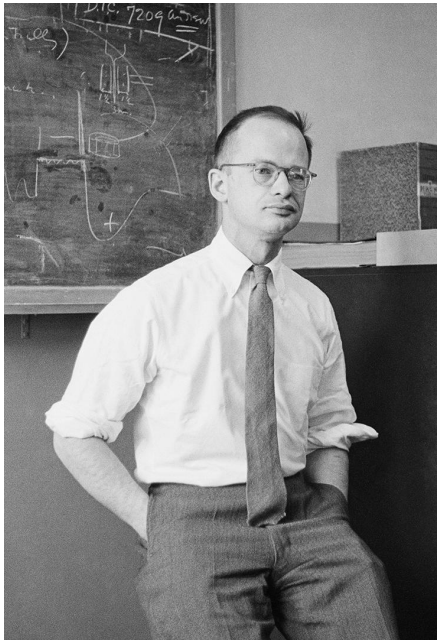

Walter H. Pitts, 23 April 1923 – 14 May 1969, USA

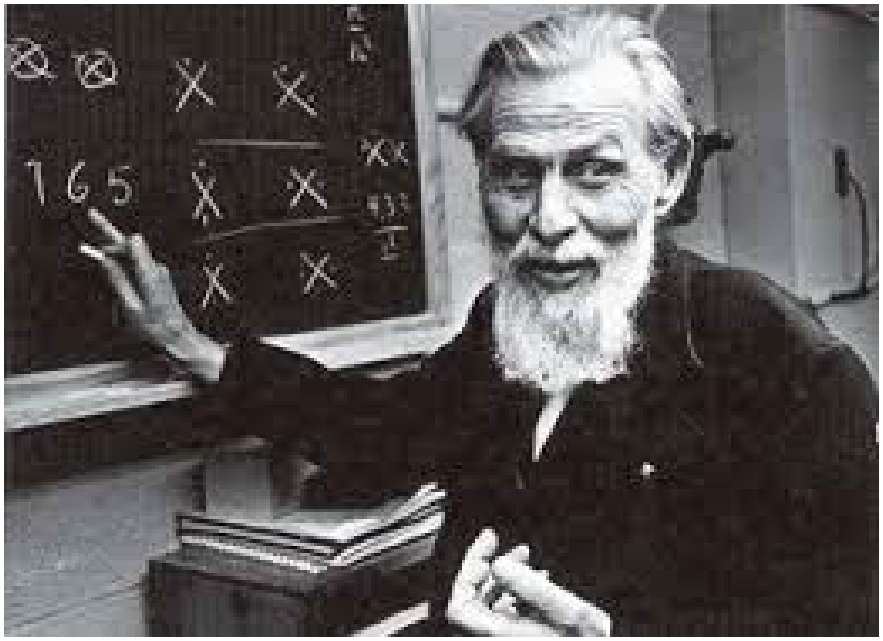

Warren Sturgis McCulloch, Nov. 16, 1898 – Sep. 24, 1969, USA

Neural networks

- In 1948, McCulloch concerned that eventually AI might rule humankind, a topic of much concern in these increasingly automated times. The below article from the September 22, 1948 Brooklyn Daily Eagle records his clarion call about the future

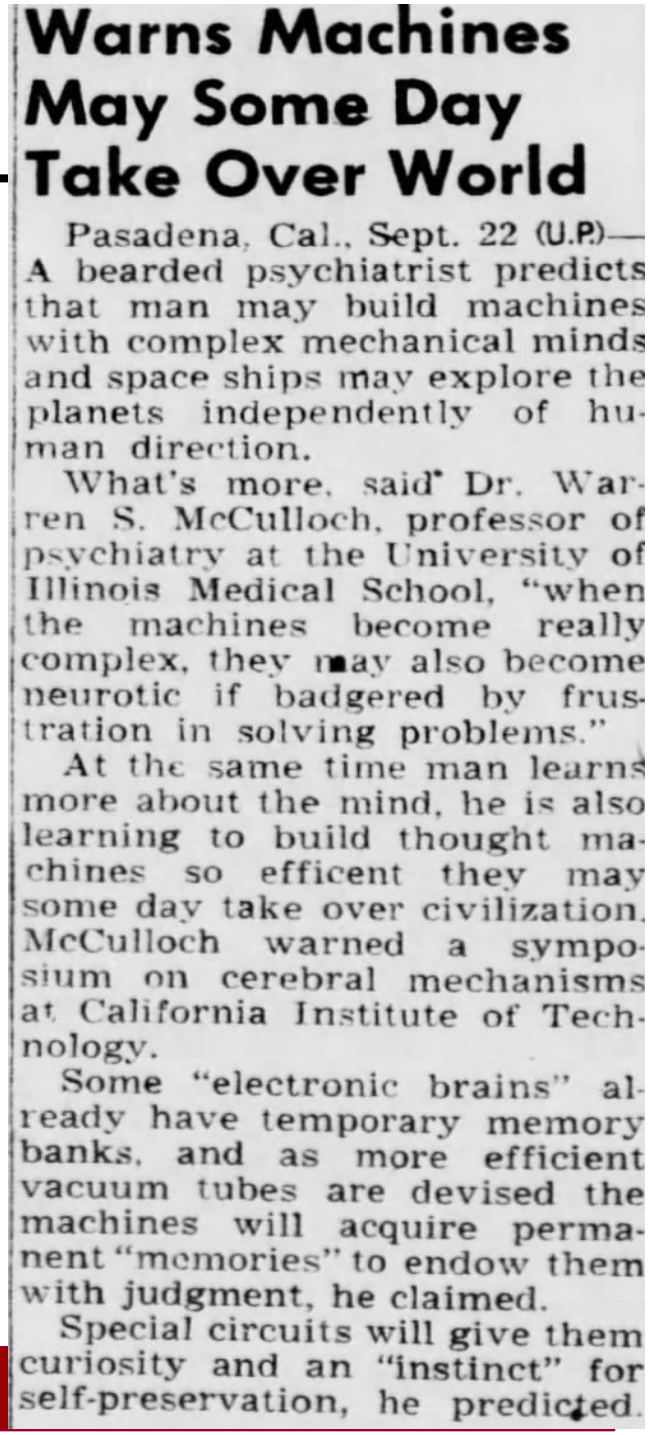

Warns Machines May Some Day Take Over World

Pasadena, Cal., Sept. 22 (U.P.)—A bearded psychiatrist predicts that man may build machines with complex mechanical minds and space ships may explore the planets independently of human direction.

What's more, said Dr. Warren S. McCulloch, professor of psychiatry at the University of Illinois Medical School, "when the machines become really complex, they may also become neurotic if badgered by frustration in solving problems."

At the same time man learns more about the mind, he is also learning to build thought machines so efficent they may some day take over civilization. McCulloch warned a symposium on cerebral mechanisms at California Institute of Technology.

Some "electronic brains" already have temporary memory banks, and as more efficient vacuum tubes are devised the machines will acquire permanent "memories" to endow them with judgment, he claimed.

Special circuits will give them curiosity and an "instinct" for self-preservation, he predicted.

Neural networks

- It is one way to solve a supervised learning problem given labeled training examples $\{\boldsymbol{x}_i, y_i\}(i = 1, \ldots, m)$
- Neural networks give a way of defining a complex, non-linear form of hypothesis $h_{W,b}(\boldsymbol{x})$, where W and b are the parameters we need to learn from training samples

Neural networks

- A single neuron
  - x_1, x_2, and x_3 are the inputs, +1 is the intercept term, $h_{W,b}(\boldsymbol{x})$ is the output of this neuron

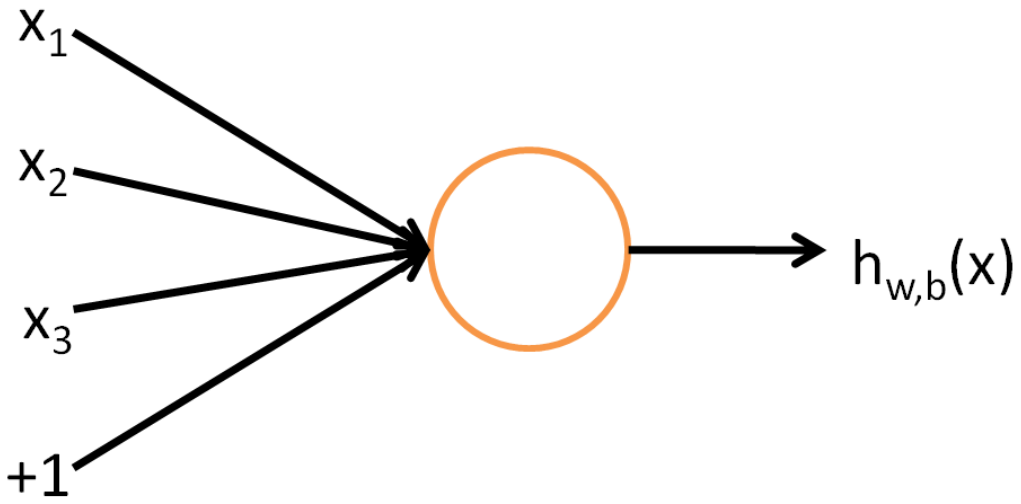

$$h_{W,b}(x) = f\left(W^T \boldsymbol{x}\right) = f\left(\sum_{i=1}^{3} W_i x_i + b\right)$$

where $f(\cdot)$ is the activation function

Neural networks

- Commonly used activation functions
  - Sigmoid function

$$f(z) = \frac{1}{1 + \exp(-z)}$$

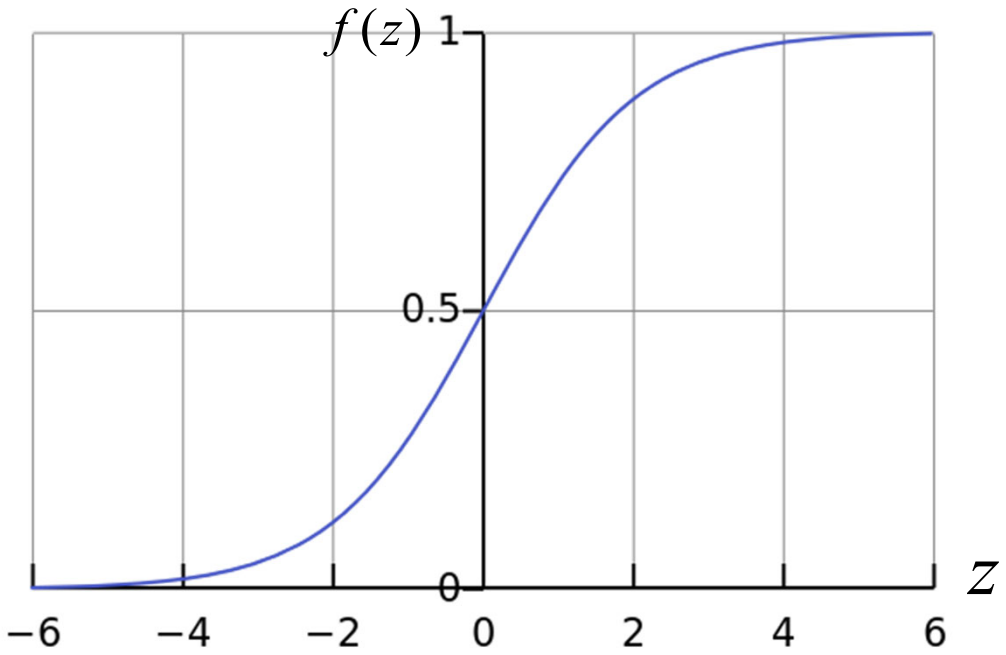

Neural networks

- Commonly used activation functions
  - Tanh function

$$f(z) = \tanh(z) = \frac{e^{z} - e^{-z}}{e^{z} + e^{-z}}$$

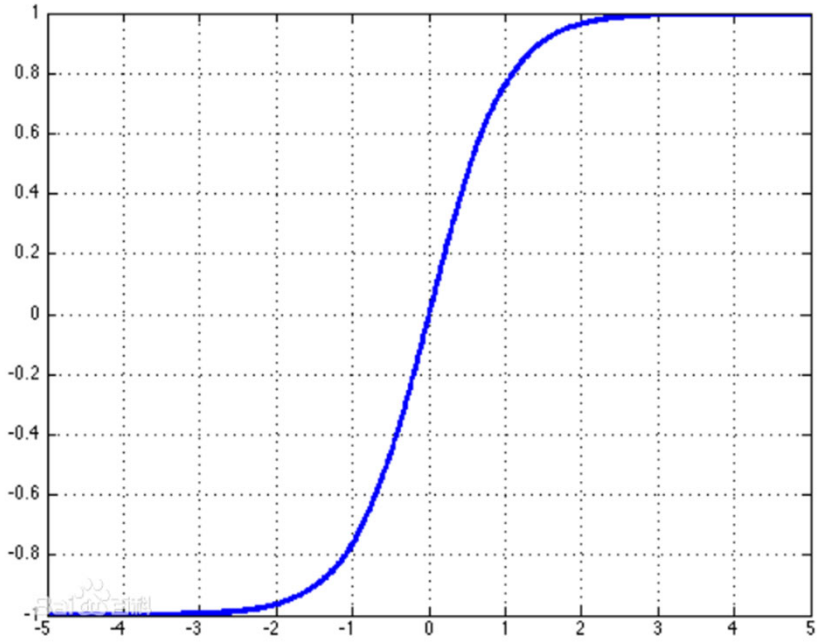

Neural networks

- Commonly used activation functions
  - Rectified linear unit (ReLU)

$$f(z) = \max(0, z)$$

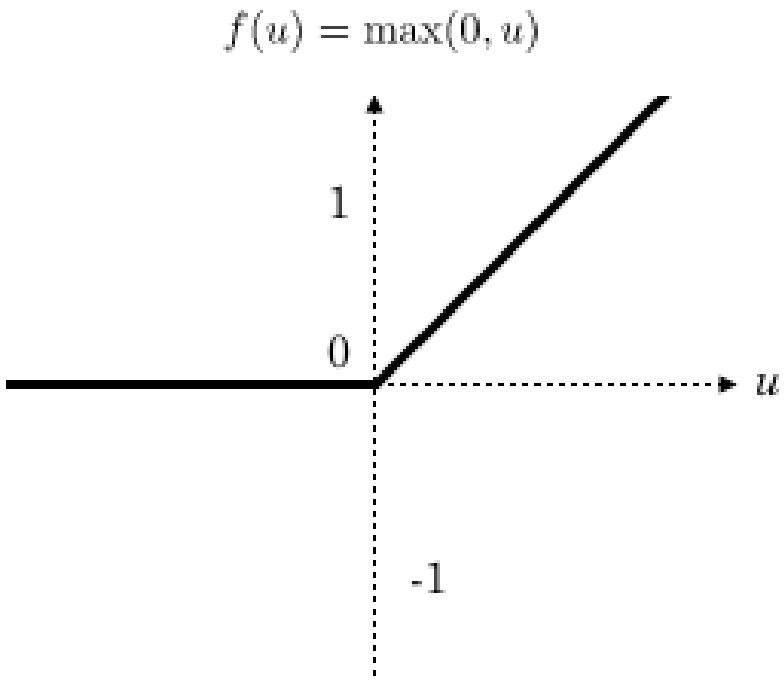

Neural networks

- Commonly used activation functions
  - Leaky Rectified linear unit (ReLU)

$$f(z)=\begin{cases} z, if\ \ z>0 \\ 0.01z, otherwise \end{cases}$$

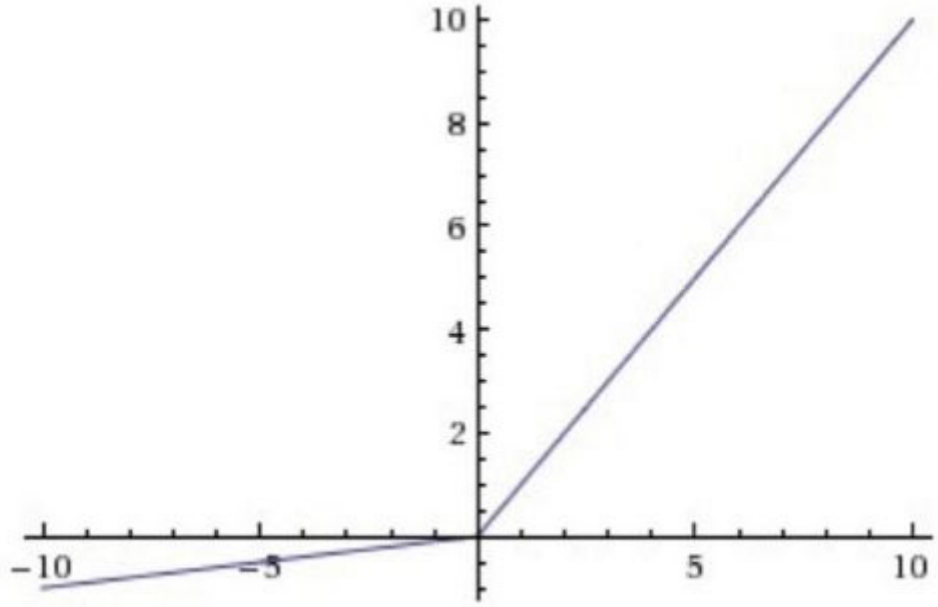

Neural networks

- Commonly used activation functions
  - Softplus (can be regarded as a smooth approximation to ReLU)

$$f(z) = \ln\left(1 + e^{z}\right)$$

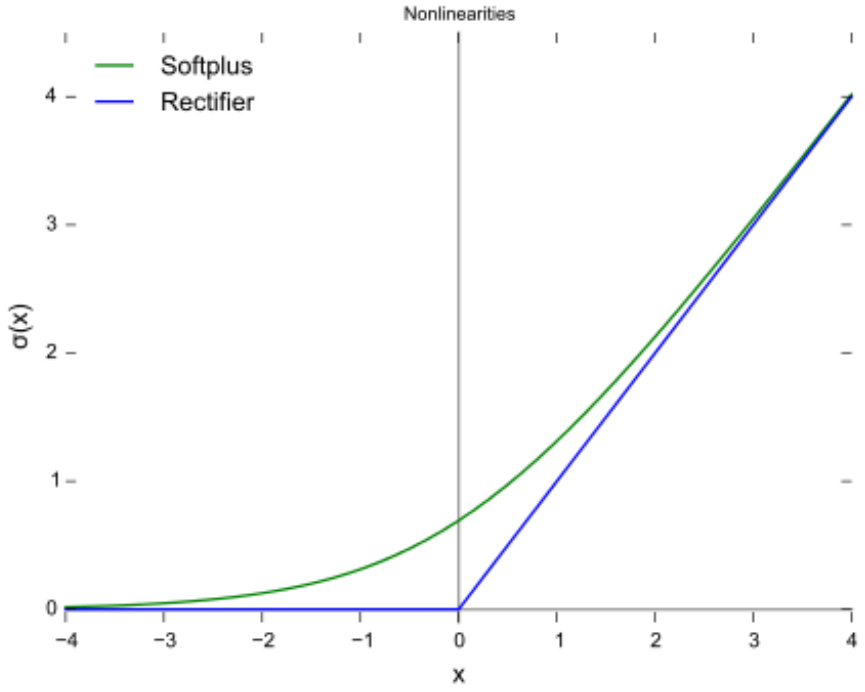

Neural networks

- A neural network is composed by hooking together many simple neurons
- The output of a neuron can be the input of another
- Example, a three layers neural network,

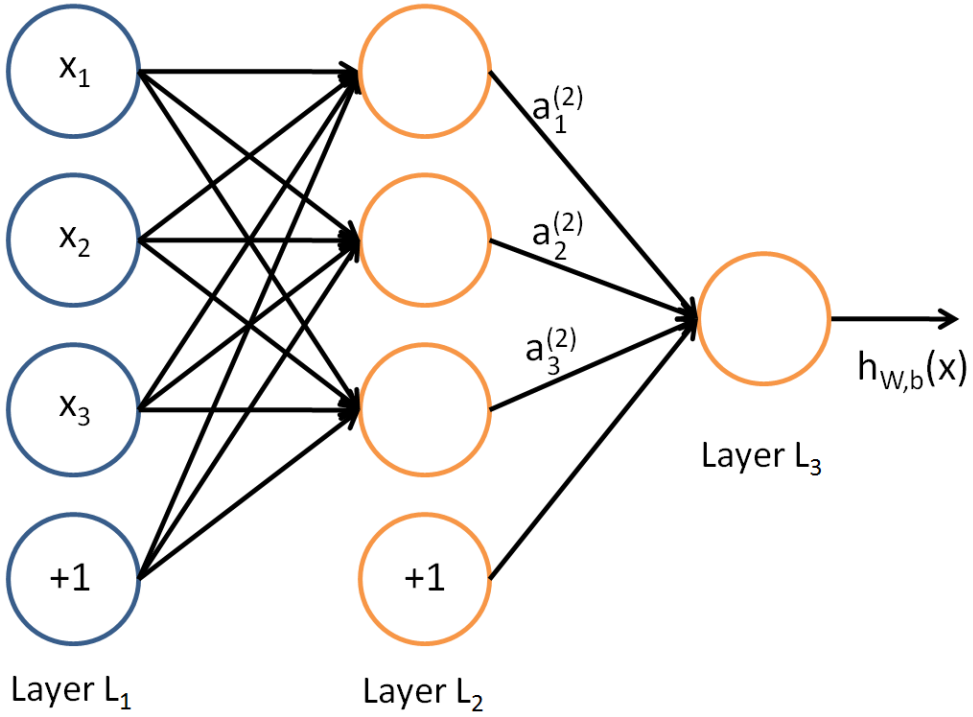

Neural networks

- Terminologies about the neural network
  - The circle labeled +1 are called **bias units**
  - The leftmost layer is called the **input layer**
  - The rightmost layer is the **output layer**
  - The middle layer of nodes is called the **hidden layer**
    - » In our example, there are 3 input units, 3 hidden units, and 1 output unit
  - We denote the activation (output value) of unit i in lay l as $a_i^{(l)}$

Neural networks

$$a_1^{(2)} = f\left(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + b_1^{(1)}\right)$$

$$a_2^{(2)} = f\left(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)}\right)$$

$$a_3^{(2)} = f\left(W_{31}^{(1)}x_1 + W_{32}^{(1)}x_2 + W_{33}^{(1)}x_3 + b_3^{(1)}\right)$$

$$h_{W,b}(\boldsymbol{x}) = a_1^{(3)} = f\left(W_{11}^{(2)}a_1^{(2)} + W_{12}^{(2)}a_2^{(2)} + W_{13}^{1}a_3^{(2)} + b_1^{(2)}\right)$$

Neural networks

- Neural networks can have multiple outputs
- Usually, we can add a softmax layer as the output layer to perform multiclass classification

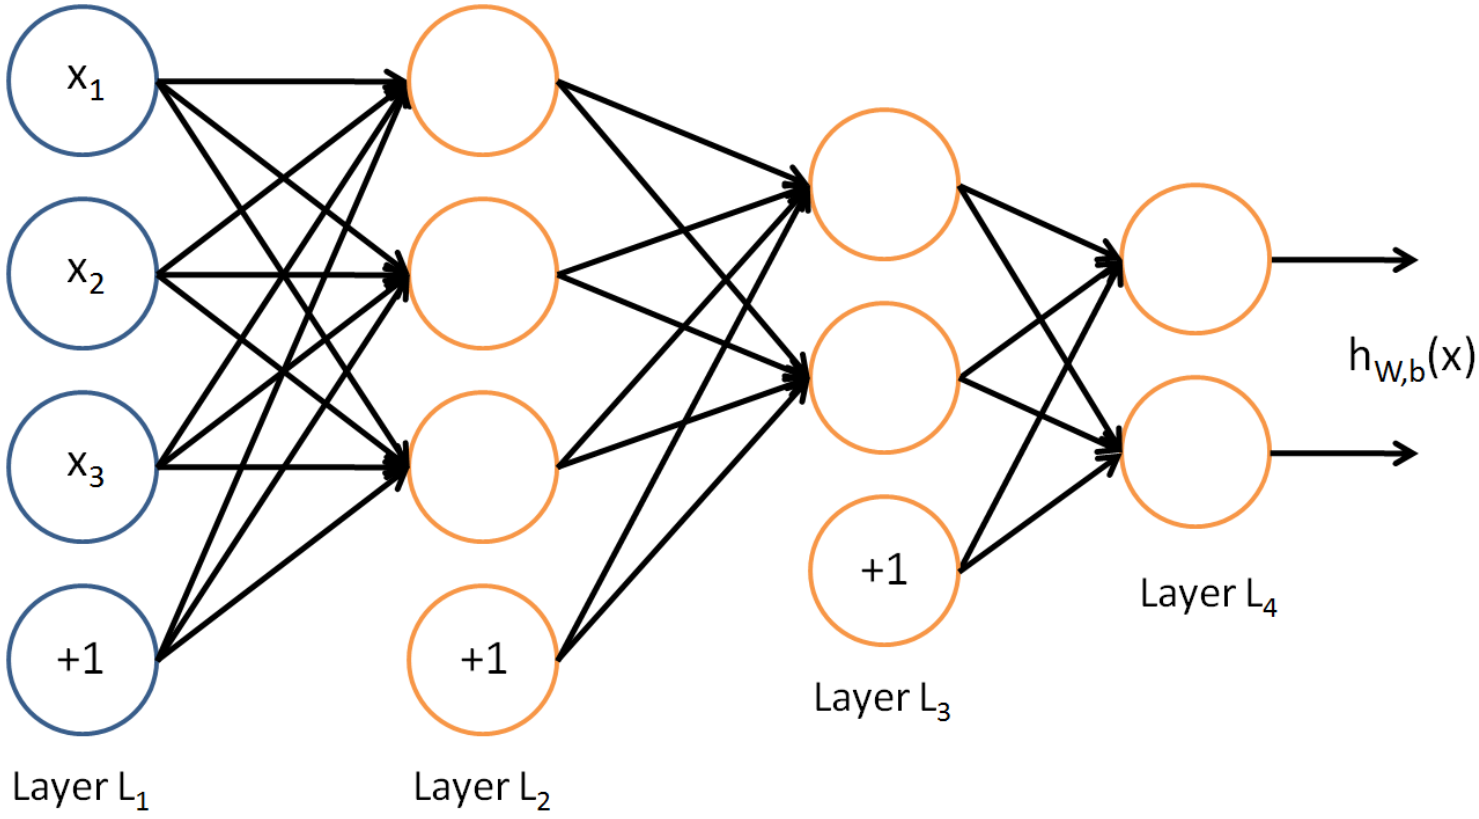

Neural networks

- At the testing stage, given a test input $\boldsymbol{x}$, it is straightforward to evaluate its output
- At the training stage, given a set of training samples, we need to train W and b
  - The key problem is how to compute the gradient
  - Backpropagation algorithm

Neural networks

- Backpropagation
  - A common method of training artificial neural networks and used in conjunction with an optimization method such as gradient descent
  - Its purpose is to compute the partial derivative of the loss to each parameter (weights)
  - neural nets will be very large: impractical to write down gradient formula by hand for all parameters
  - recursive application of the chain rule along a computational graph to compute the gradients of all parameters

Neural networks

- Backpropagation

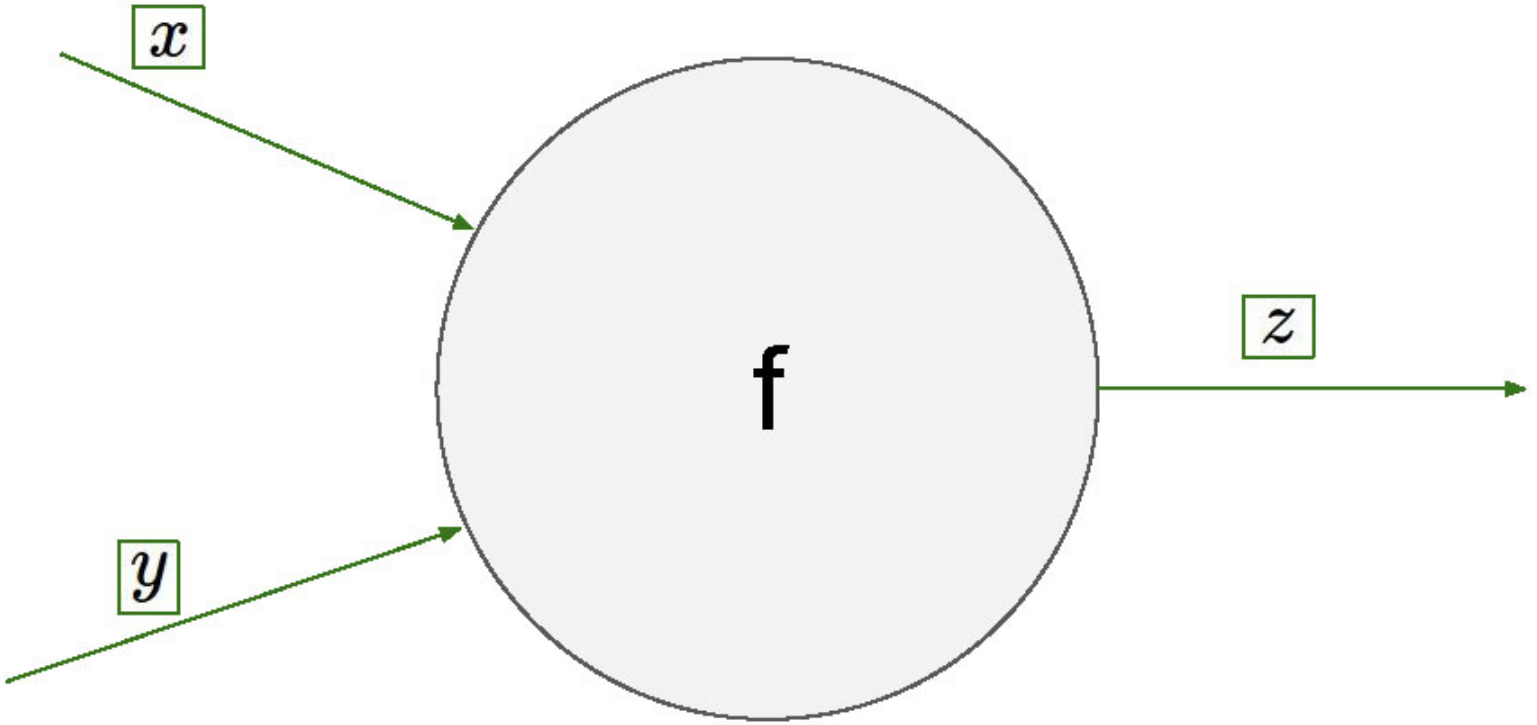

Neural networks

- Backpropagation

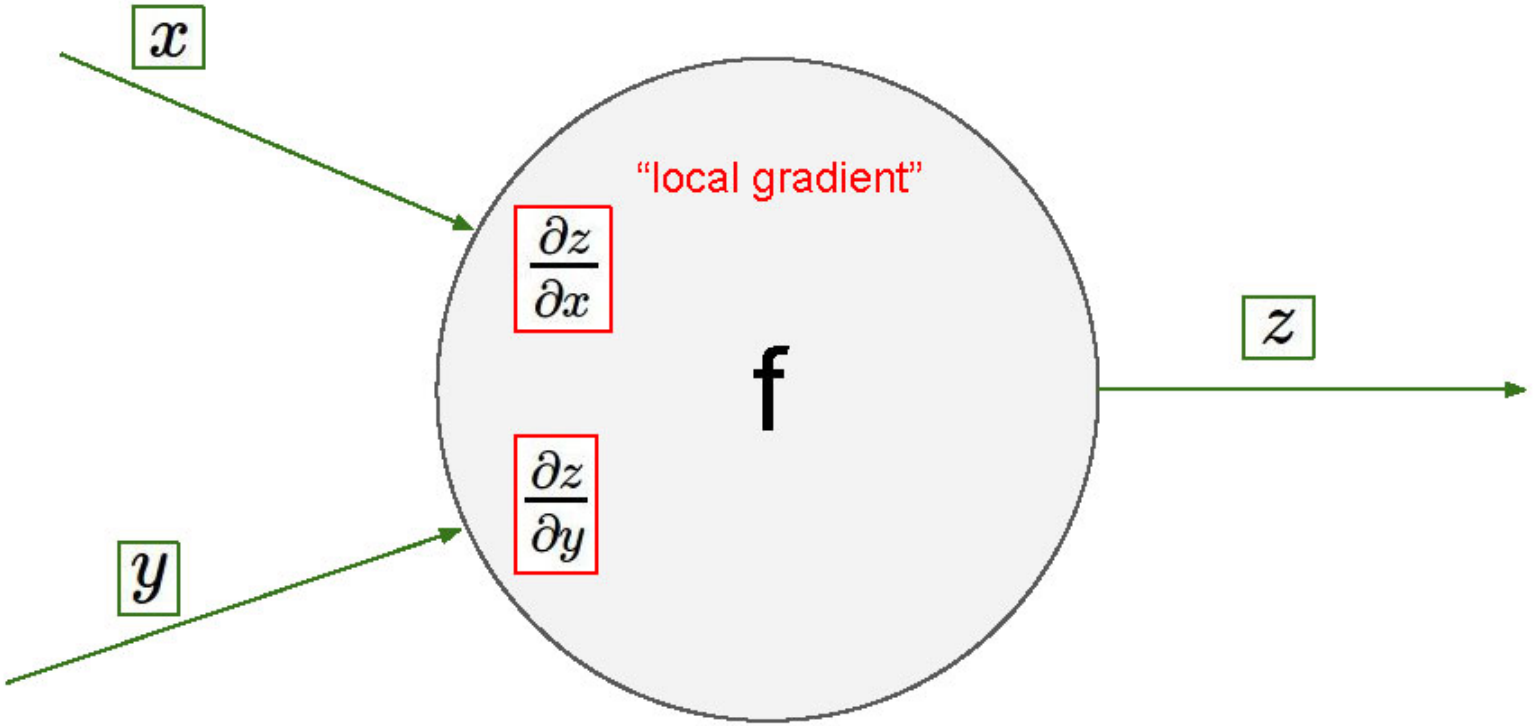

Neural networks

- Backpropagation

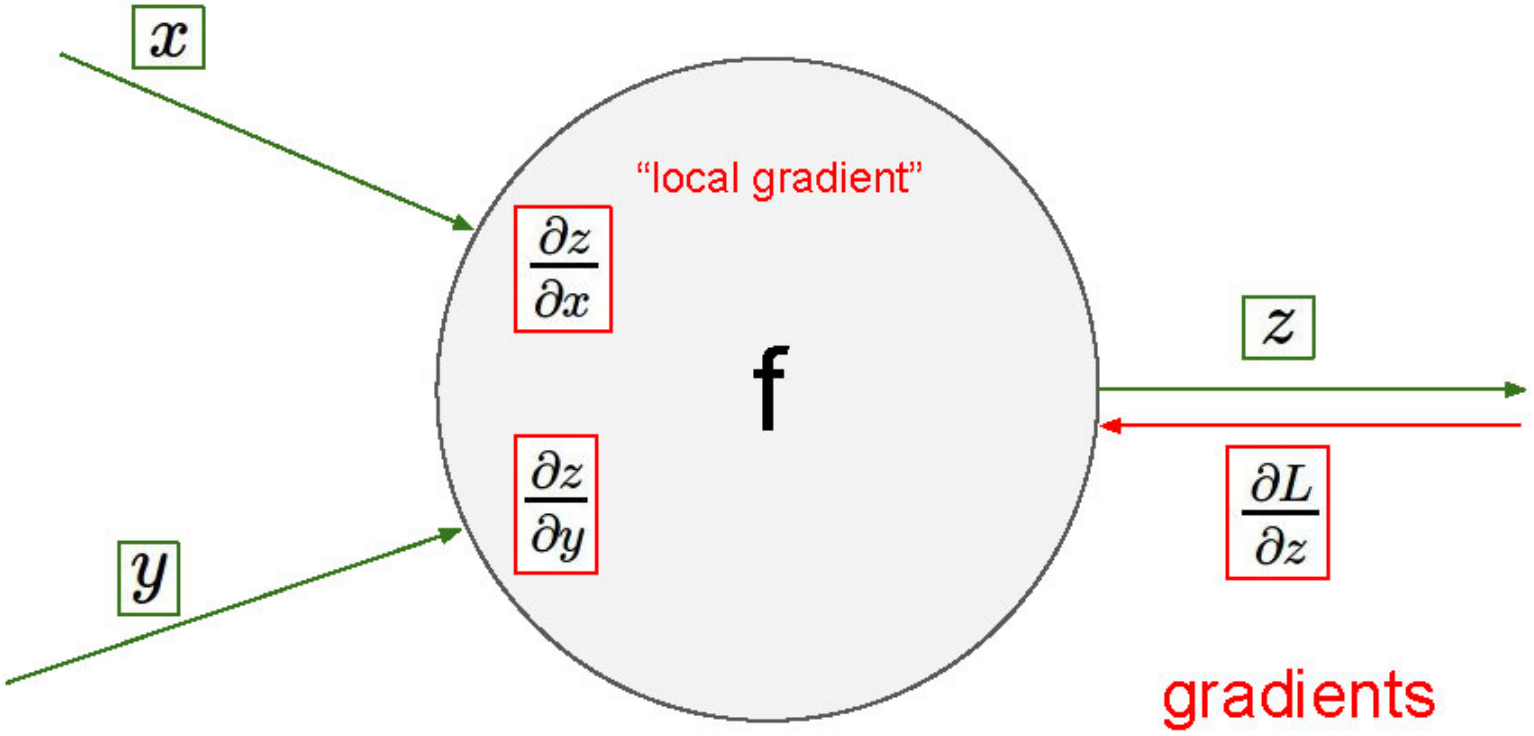

Neural networks

- Backpropagation

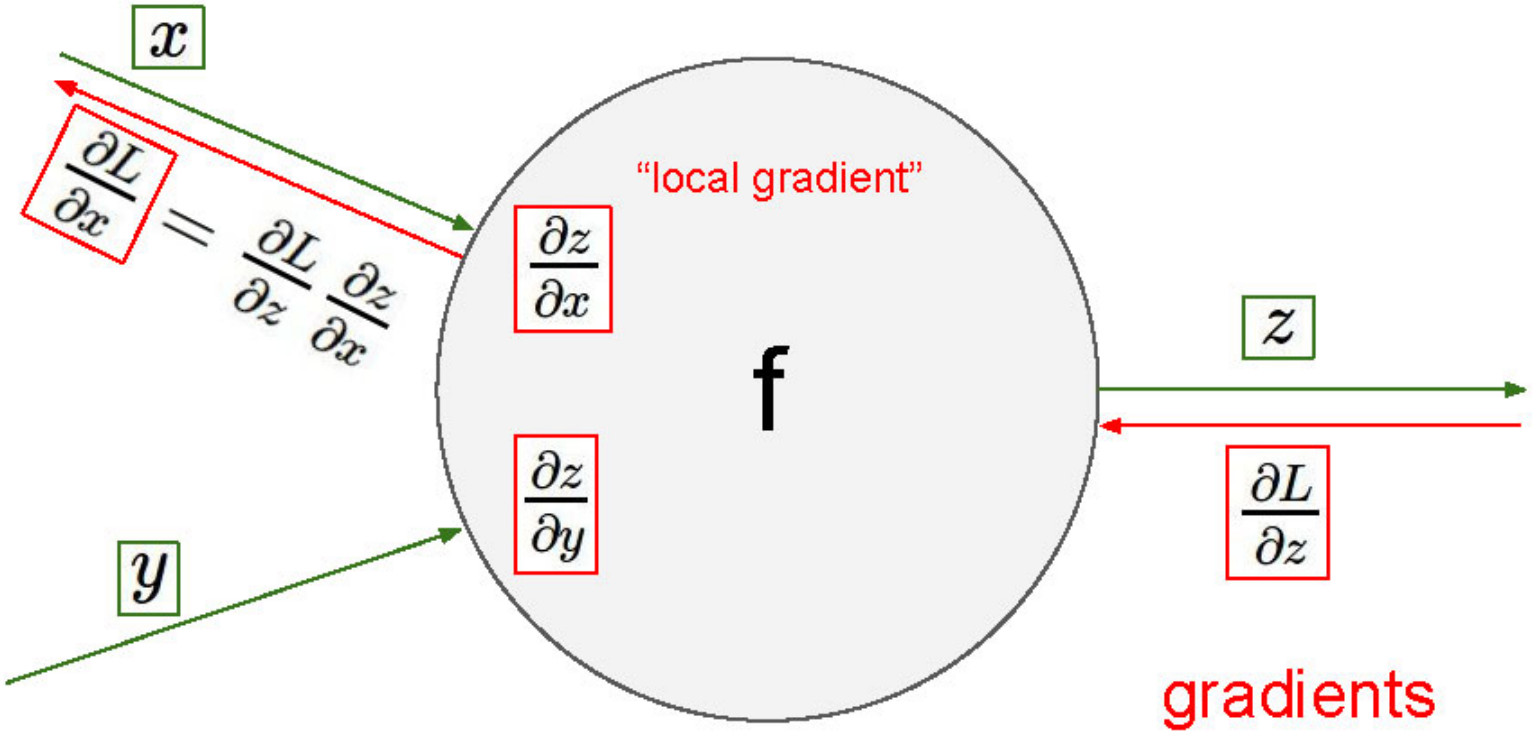

Neural networks

- Backpropagation

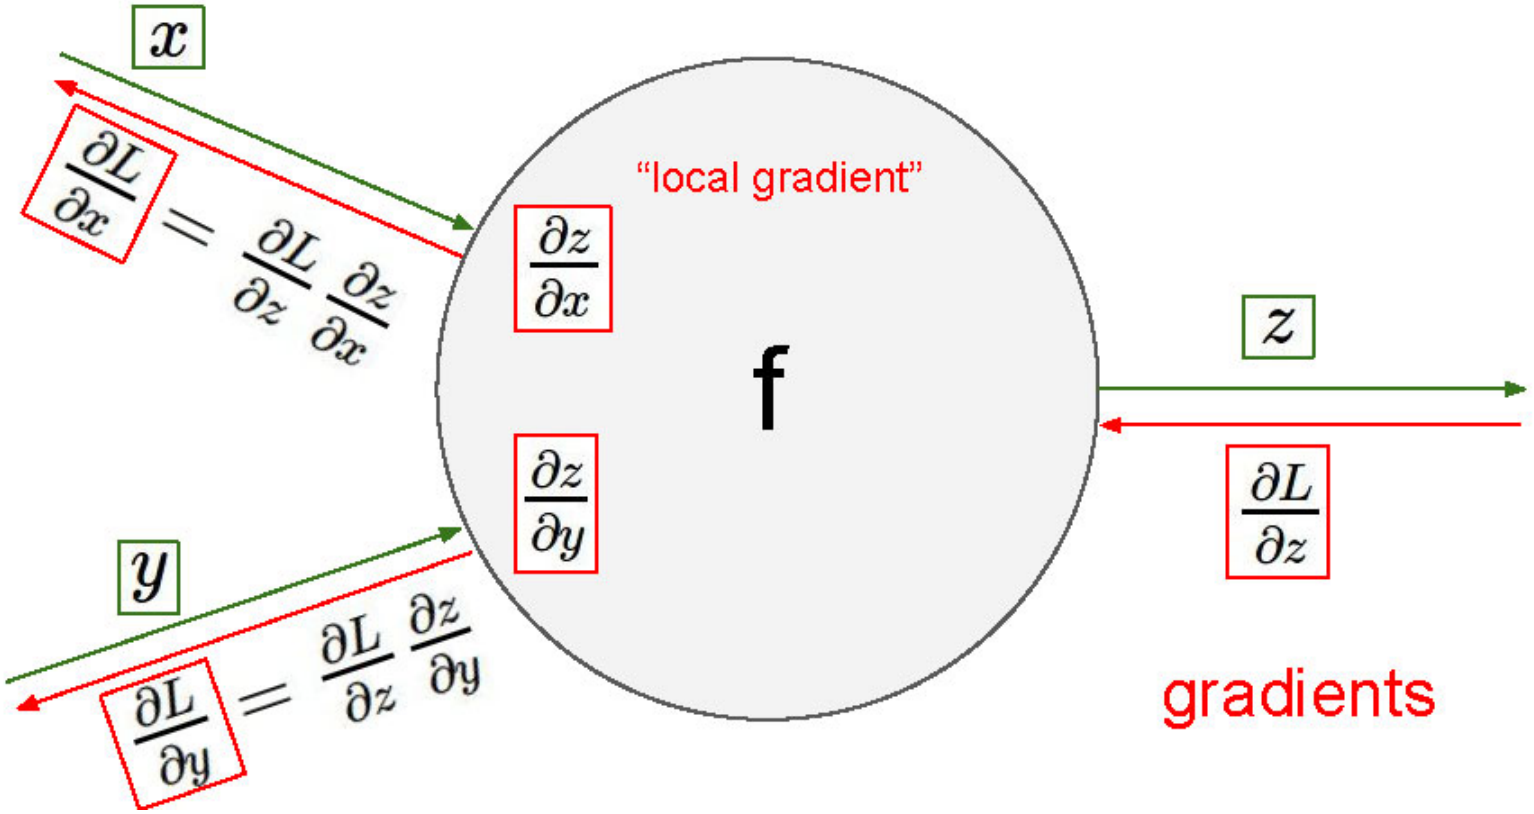

Neural networks

- Backpropagation

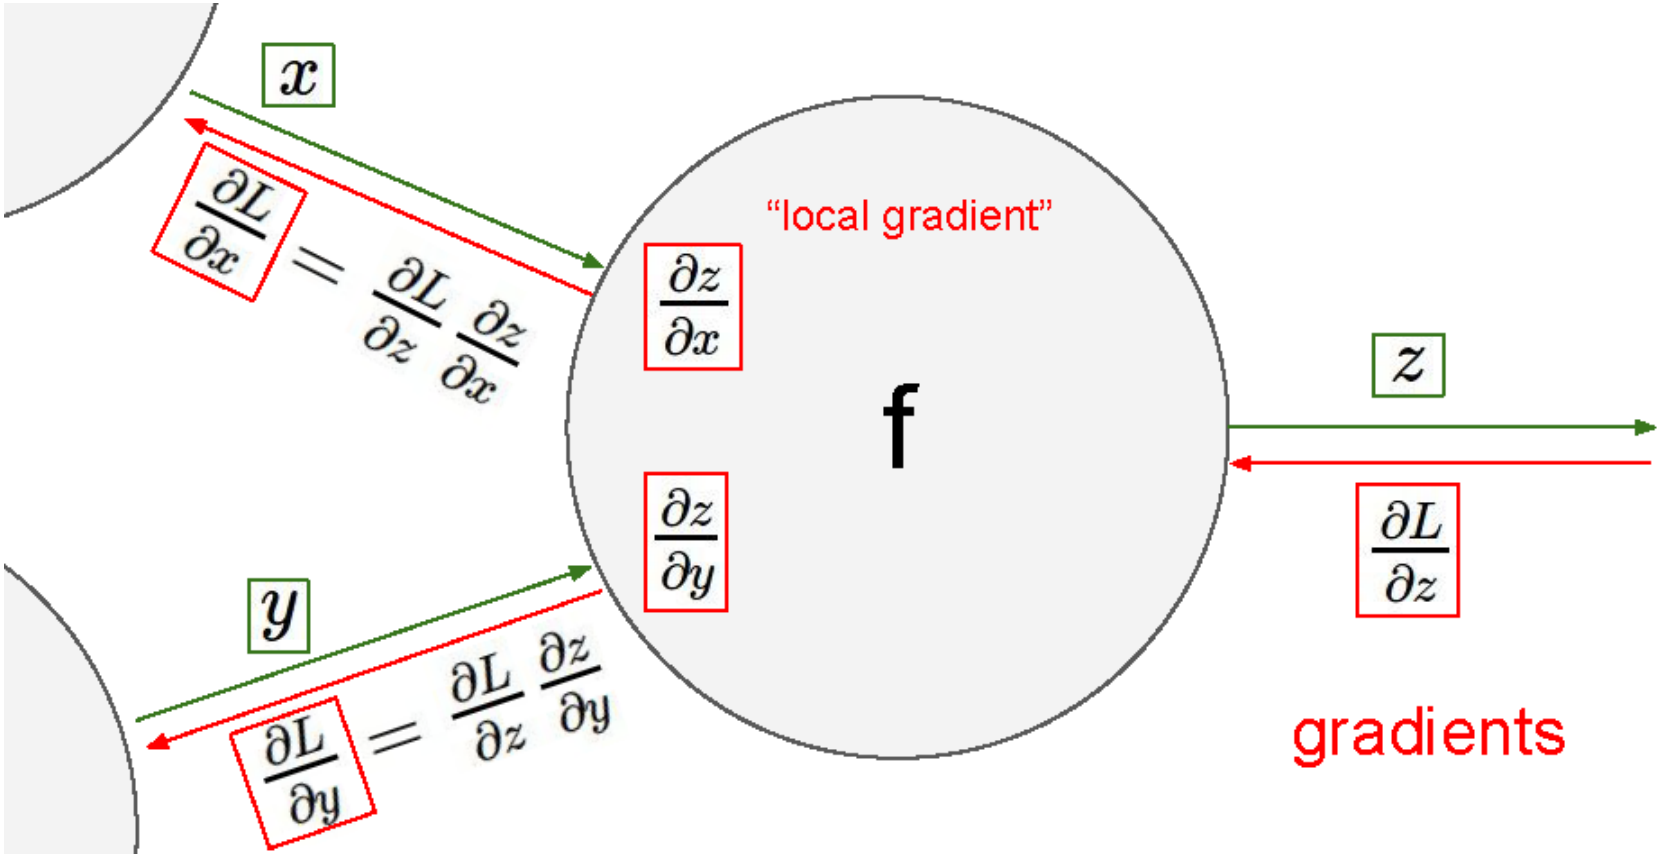

Some techniques for better training neural networks

Neural network is trained by iteratively updating the network parameters $\boldsymbol{\theta}$ using the following basic rule,

$$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \alpha \mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right)$$

where $\mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right) \in \mathbb{R}^{|\boldsymbol{\theta}|}$ is the gradient of the loss function at nth iteration

- Since the problem of neural network training is extremely large-scale, it is not efficient to find the exact optimal α at each iteration step; instead, we need to use some heuristic strategies to appropriately adjust α
- Another issue is the effective estimation of gradients. The training of neural networks generally adopts SGD or mini-batch SGD, which leads to a small number of samples selected each time, resulting in the estimated gradient of each iteration not being consistent with the optimal gradient on the entire training set and having a certain degree of randomness. We can make full use of the average value of the gradients used in the last few iterations to correct the random gradient calculated in the current iteration. This is beneficial for accelerating the convergence of training

Some techniques for better training neural networks

- Heuristic strategies to appropriately adjust α
  - Learning rate decay

    Inverse time decay, $\alpha_t = \alpha_0 \frac{1}{1+\beta \times t}$

    Exponential decay, $\alpha_t = \alpha_0 \beta^t, 0 < \beta < 1$

    Natural Exponential decay, $\alpha_t = \alpha_0 \exp(-\beta \times t)$

  where α_0 is the initial learning rate, t is the iteration index, and β is the decay rate

Some techniques for better training neural networks

- Heuristic strategies to appropriately adjust α
  - Learning rate decay
  - Parameter-adaptive learning rate adjusting policies
    - » AdaGrad[2]

At n-th iteration, first compute the cumulative square of each parameter's partial derivatives,

$$\mathbf{G}_n = \sum_{\tau=1}^{n} \mathbf{g}_\tau\left(\boldsymbol{\theta}_{\tau-1}\right) \odot \mathbf{g}_\tau\left(\boldsymbol{\theta}_{\tau-1}\right)$$

where $\mathbf{g}_\tau \in \mathbb{R}^{|\boldsymbol{\theta}|}$ is the gradient at τ-th iteration, $\odot$ is the elementwise product

Then, $\boldsymbol{\theta}$ is updated as,

$$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \frac{\alpha_0}{\sqrt{\mathbf{G}_n + \varepsilon}} \odot \mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right)$$

Note: all the operations here are elementwise

[2] J. Duchi *et al.*, Adaptive subgradient methods for online learning and stochastic optimization. Journal of machine learning research, 12(7), 2011

Some techniques for better training neural networks

- Heuristic strategies to appropriately adjust α
  - Learning rate decay
  - Parameter-adaptive learning rate adjusting policies
    - » AdaGrad[2]

Properties:

① At this iteration, for a parameter, if its cumulative square values of partial derivates is large, the associated learning rate will be small, otherwise large

② The learning rate for each parameter will continuously decrease as the number of iterations increases; sometimes it is difficult for it to find the optimal solution since the learning rates have already been too small

[2] J. Duchi *et al.*, Adaptive subgradient methods for online learning and stochastic optimization. Journal of machine learning research, 12(7), 2011

Some techniques for better training neural networks

- Heuristic strategies to appropriately adjust α
  - Learning rate decay
  - Parameter-adaptive learning rate adjusting policies
    - » AdaGrad[2]
    - » RMSprop[3] (Root Mean Squared Propagation)

At n-th iteration, first compute the moving average square of each parameter's partial derivatives,

$$\mathbf{G}_n = \beta \mathbf{G}_{n-1} + (1-\beta)\mathbf{g}_n(\boldsymbol{\theta}_{n-1}) \odot \mathbf{g}_n(\boldsymbol{\theta}_{n-1}) = (1-\beta)\sum_{\tau=1}^{n} \beta^{n-\tau} \mathbf{g}_\tau(\boldsymbol{\theta}_{\tau-1}) \odot \mathbf{g}_\tau(\boldsymbol{\theta}_{\tau-1})$$

where β is usually set as 0.9

Then, $\boldsymbol{\theta}$ is updated as,

$$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \frac{\alpha_0}{\sqrt{\mathbf{G}_n + \varepsilon}} \odot \mathbf{g}_n(\boldsymbol{\theta}_{n-1})$$

[3] T. Tieleman and G. Hinton, Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude, COURSERA: Neural networks for machine learning 4 (2), 26-31, 2012

Some techniques for better training neural networks

- Heuristic strategies to appropriately adjust α
  - Learning rate decay
  - Parameter-adaptive learning rate adjusting policies
    - » AdaGrad[2]
    - » RMSprop[3] (Root Mean Squared Propagation)

Properties:

① At this iteration, for a parameter, if its associate element in $\mathbf{G}_n$ is large, the associated learning rate will be small, otherwise large

② Different from AdaGrad, the learning rate for each parameter does not monotonically decrease as the number of iterations increases

[3] T. Tieleman and G. Hinton, Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude, COURSERA: Neural networks for machine learning 4 (2), 26-31, 2012

Some techniques for better training neural networks

- Correction for gradient estimation
  - Momentum method

$$\Delta\boldsymbol{\theta}_n = \rho\Delta\boldsymbol{\theta}_{n-1} - \alpha\mathbf{g}_n\left(\boldsymbol{\theta}_{n-1}\right) = -\alpha\sum_{\tau=1}^{n}\rho^{n-\tau}\mathbf{g}_\tau\left(\boldsymbol{\theta}_{\tau-1}\right)$$

where ρ is usually set as 0.9

($\boldsymbol{\theta}_n$ is updated as $\boldsymbol{\theta}_n = \boldsymbol{\theta}_{n-1} + \Delta\boldsymbol{\theta}_n$)

Properties:

① The gradient used in this iteration is the weighted average of the past gradients

Some techniques for better training neural networks

- Correction for gradient estimation
  - Momentum method
  - NAG (Nesterov Accelerated Gradient)[4]

Moment method

The parameter update is completed as,

$$\begin{cases} \hat{\boldsymbol{\theta}} = \boldsymbol{\theta}_{n-1} + \rho \Delta \boldsymbol{\theta}_{n-1} \\ \boldsymbol{\theta}_n = \hat{\boldsymbol{\theta}} - \alpha \mathbf{g}_n \left(\boldsymbol{\theta}_{n-1} \right) \end{cases}$$

It's more valid if we take gradient at $\hat{\boldsymbol{\theta}}$

improved

NAG

The parameter update is completed as,

$$\begin{cases} \hat{\boldsymbol{\theta}} = \boldsymbol{\theta}_{n-1} + \rho \Delta \boldsymbol{\theta}_{n-1} \\ \boldsymbol{\theta}_n = \hat{\boldsymbol{\theta}} - \alpha \mathbf{g}_n \left(\hat{\boldsymbol{\theta}} \right) = \hat{\boldsymbol{\theta}} - \alpha \mathbf{g}_n \left(\boldsymbol{\theta}_{n-1} + \rho \Delta \boldsymbol{\theta}_{n-1} \right) \end{cases}$$

with NAG

$$\Delta \boldsymbol{\theta}_n = \rho \Delta \boldsymbol{\theta}_{n-1} - \alpha \mathbf{g}_n \left(\boldsymbol{\theta}_{n-1} + \rho \Delta \boldsymbol{\theta}_{n-1} \right)$$

[4] Y. Nesterov, Gradient methods for minimizing composite functions, Mathematical Programming, 140 (1), 125-161, 2013

Some techniques for better training neural networks

- Correction for gradient estimation
  - Momentum method
  - NAG (Nesterov Accelerated Gradient)[4]
  - Adam (Adaptive Moment Estimation Algorithm) [5]

$$\boldsymbol{M}_n = \beta_1 \boldsymbol{M}_{n-1} + (1-\beta_1)\mathbf{g}_n(\boldsymbol{\theta}_{n-1})$$ (similar as the moment method)

$$\mathbf{G}_n = \beta_2 \mathbf{G}_{n-1} + (1-\beta_2)\mathbf{g}_n(\boldsymbol{\theta}_{n-1}) \odot \mathbf{g}_n(\boldsymbol{\theta}_{n-1})$$ (similar as RMSprop)

(β_1 and β_2 are usually set as 0.9 and 0.99, respectively)

When n is small, M_n and G_n are not accurate, so they need to be corrected as,

$$\widehat{\mathbf{M}}_n = \frac{\mathbf{M}_n}{1-\beta_1^n} \qquad \widehat{\mathbf{G}}_n = \frac{\mathbf{G}_n}{1-\beta_2^n}$$

Update is completed as, $$\boldsymbol{\theta}_n := \boldsymbol{\theta}_{n-1} - \frac{\alpha_0}{\sqrt{\widehat{\mathbf{G}}_n + \varepsilon}} \odot \widehat{\mathbf{M}}_n$$

Adam can be deemed as a combination of RMSprop and moment method

[5] D. Kingma and J. Ba, Adam: A method for stochastic optimization, ICLR, 2015

Outline

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- CNN for object detection

A little history about deep learning

深度学习

2006年，Hinton和他的学生Salakhutdinov在《科学》上发表了一篇文章，开启了深度学习在学术界和工业界的浪潮。

Reducing the Dimensionality of Data with Neural Networks

文章提出了深层网络训练中梯度消失问题的解决方案：
无监督预训练对权值进行初始化+有监督训练微调

2012年，Hinton课题组为了证明深度学习的潜力，首次参加ImageNet图像识别比赛，其通过构建的CNN网络AlexNet一举夺得冠军，且碾压第二名（SVM方法）的分类性能。也正是由于该比赛，CNN吸引到了众多研究者的注意

A little history about deep learning

Yann LeCun Geoffrey Hinton Yoshua Bengio

2018 ACM A.M. Turing Award for conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing

A little history about deep learning

深度学习的发展

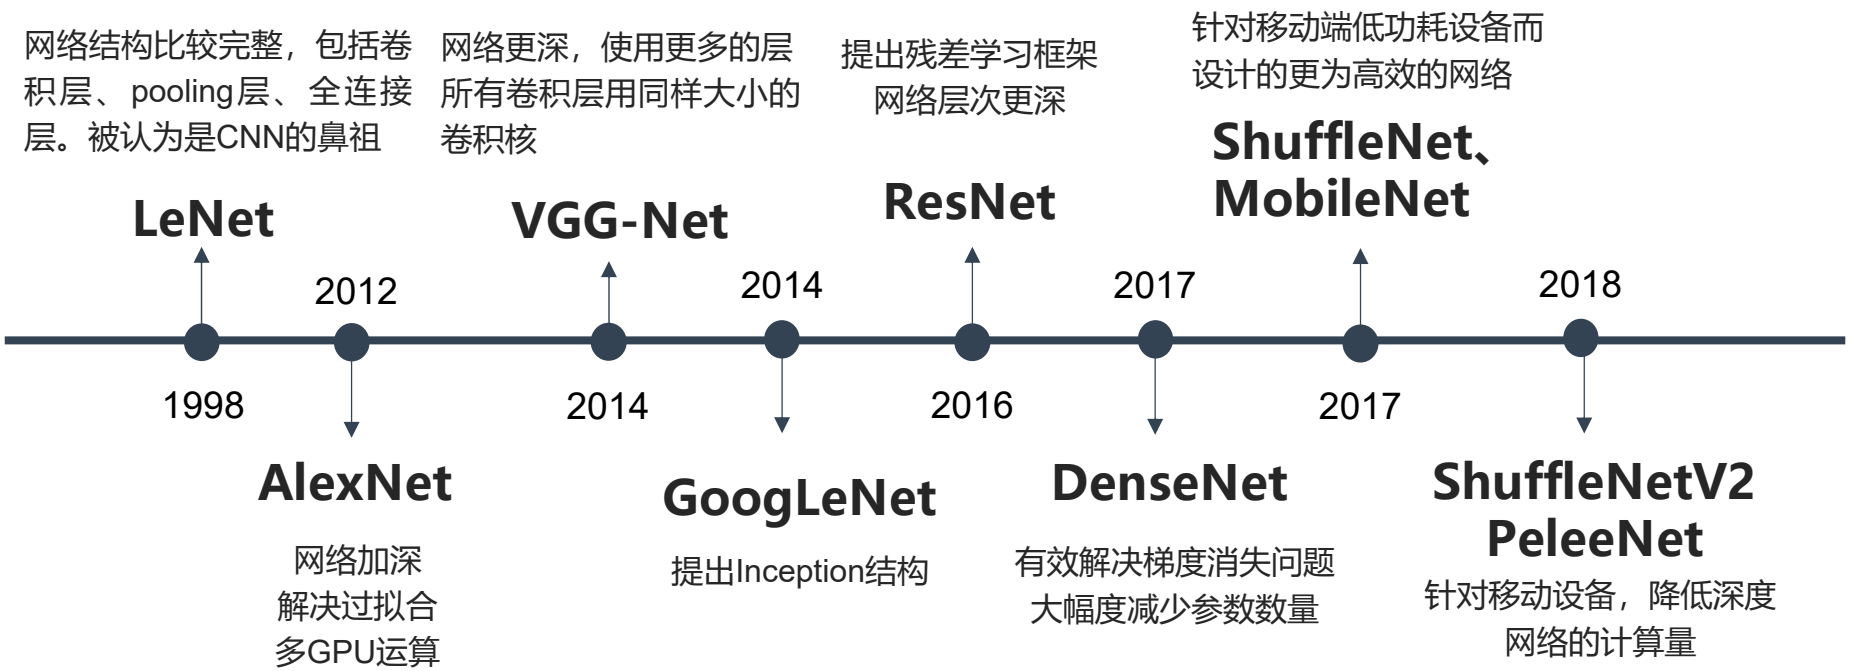

A little history about deep learning

深度学习的应用领域

语音识别

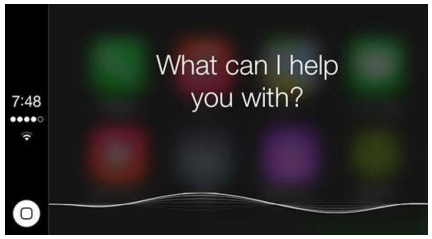

人脸认证

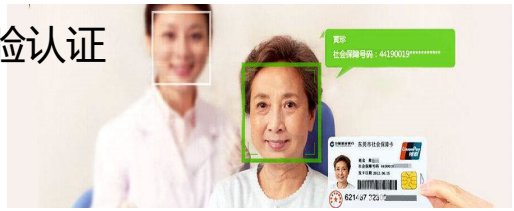

图像分类

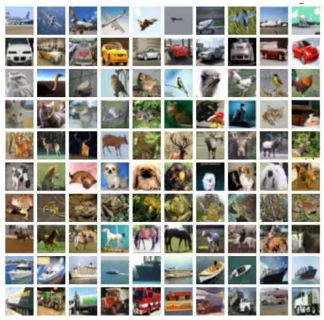

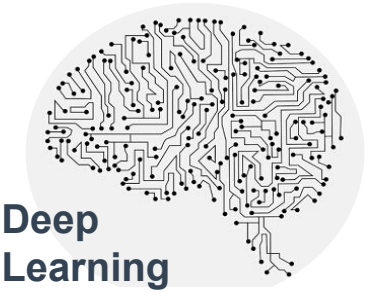

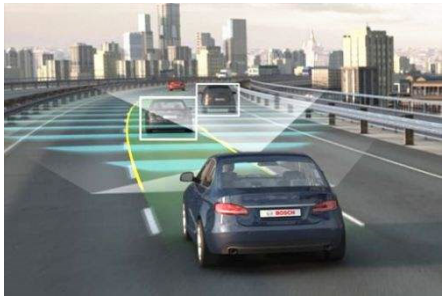

自动驾驶

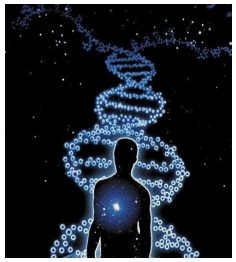

生命科学

游戏博弈

语言翻译

A little history about deep learning

深度学习为什么有效

开源的计算平台

Caffe、Tensorflow、Pytorch 等

计算力是引擎

GPU服务器、集群

数据就是燃料

如ImageNet图像分类数据集等
百万张图片

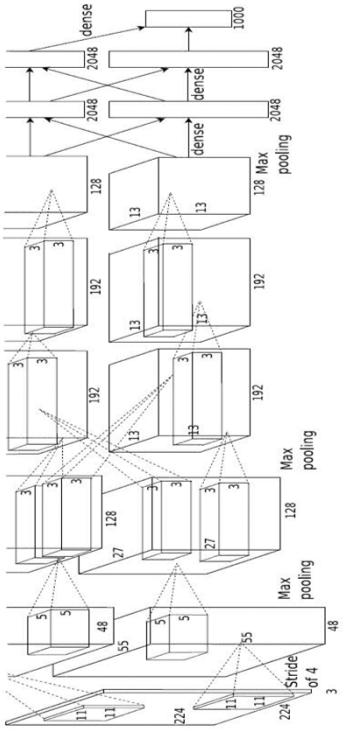

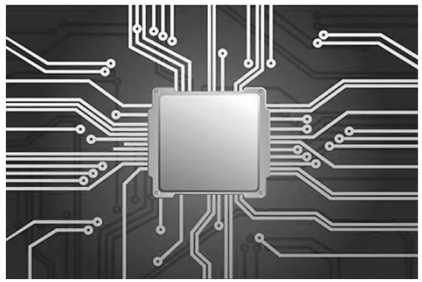

When you know these terminologies, you are an expert of DCNN

Conv layer
Pooling
Padding
Stride
ReLU
Leaky ReLU
SiLU
Res Unit
CBL (conv+BN+L-ReLU)
SPPF
(Spatial Pyramid Pooling Fast)
Fully connection
Skip connection (shortcut)
Concat
Batch normalization
Softmax
l_2-loss
Cross entropy loss
BCE (binary cross entroy)
Backbone+Head
Anchor box (point)

LeNet
AlexNet
NIN
GoogLeNet
VGGNet
ResNet
DenseNet
MobileNet
ShuffleNet
PeleeNet
ShuffleNetV2
EfficientNet

RCNN
Fast-RCNN
Faster-RCNN
SPPNet
MaskNet
SSD
Yolo
YoloV2
YoloV3
….
YoloV8
Pelee-SSD
EfficientDet

Caffe
Caffe2
TensorFlow
MatConvNet
Theano
Torch
PyTorch
Keras
MXNet
DIGITS
TensorRT

When you know these terminologies, you are an expert of DCNN

GTX980
GTX1080
TitanX
TitanXP
Tesla K40
Tesla K80
Jetson TX1
Jetson TX2
RTX 3080
RTX 3090
RTX 4080
RTX 4090

Minibatch SGD
Batch size
Iteration
Epoch
Learning rate
AdaGrad
RMSprop
Momentum
Nesterov
Adam
Finetuning
Pre-training
Early stop
Data augmentation

SYNTHIA
Virtual Kitti
URSA
VIPER
synMT
Grand Theft Auto V

Geoffrey Hinton
Alex Krizhevsky
Yoshua Bengio
Yann LeCun
Ian Goodfellow
Kaiming He
Ross Girshick
Joseph Redmon

ImageNet
VOC2007
VOC2012
MS COCO
Kitti

Convolutional neural network

- Specially designed for data with grid-like structures (LeCun et al. 98)
  - 1D grid: sequential data
  - 2D grid: image
  - 3D grid: video, 3D image volume
- Beat all the existing computer vision technologies on object recognition on ImageNet challenge with a large margin in 2012

Convolutional neural network

- Something you need to know about DCNN
  - Traditional model for PR: fixed/engineered features + trainable classifier
  - For DCNN: it is usually an **end-to-end** architecture; learning data representation and classifier together
  - The learned features from big datasets are **transferable**
  - For training a DCNN, usually we use a **fine-tuning** scheme
  - For training a DCNN, to avoid overfitting, **data augmentation** can be performed

Convolutional neural network

- One simple solution is locally connected neural networks
  - The learned filter is a spatially local pattern
  - A hidden node at a higher layer has a larger receptive field in the input
  - Stacking many such layers leads to “filters” (not anymore linear) which become increasingly “global”

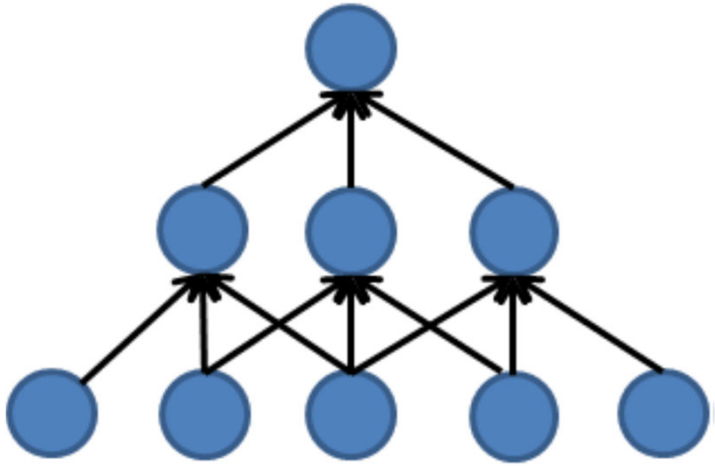

Convolutional neural network

- The first CNN
  - LeNet[6]

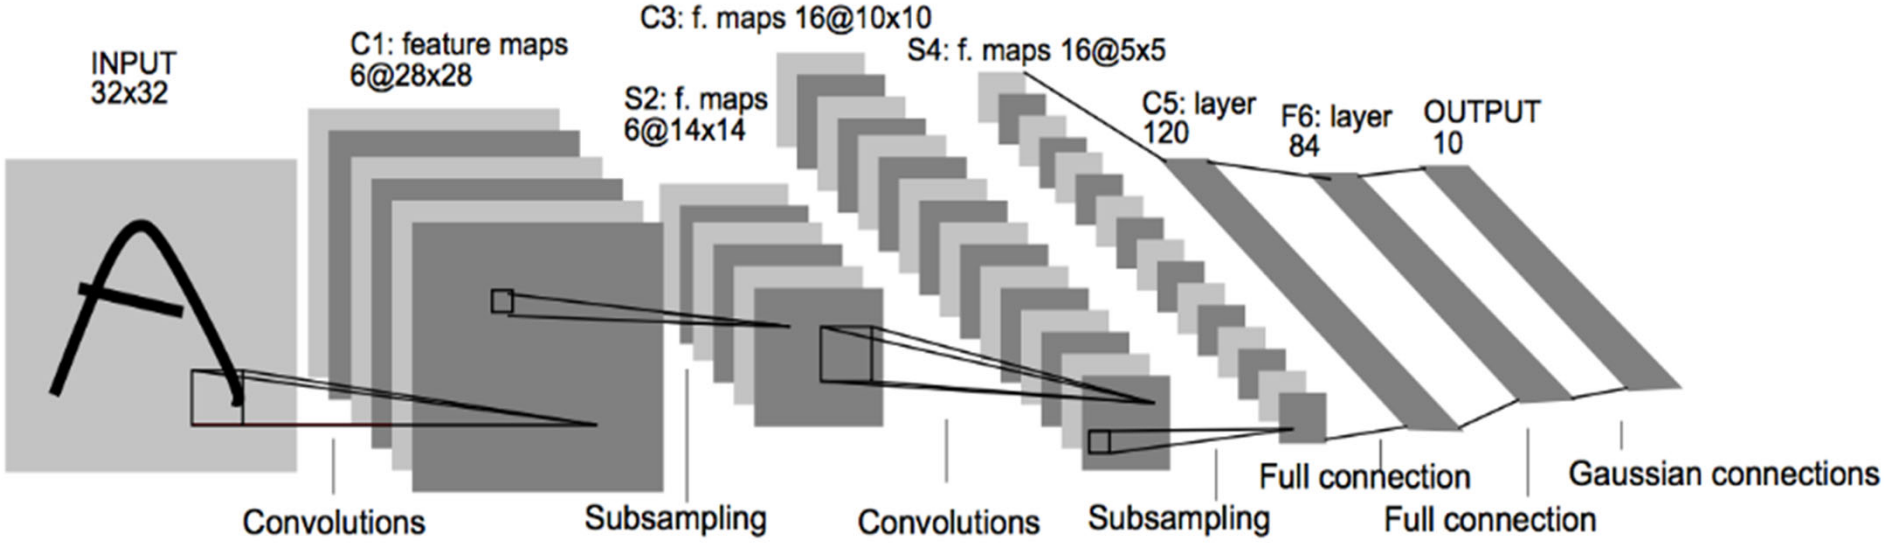

CNN called LeNet by Yann LeCun (1998)

[6] Y. LeCun et al., Gradient-based Learning Applied to Document Recognition, Proceedings of the IEEE, Vol. 86, pp. 2278-2324, 1998

Convolutional neural network

Yann LeCun (1960-)
New York University
Facebook Artificial Intelligence Research
A founding father of convolutional nets

(His name is a little difficult to pronounce)

Convolutional neural network

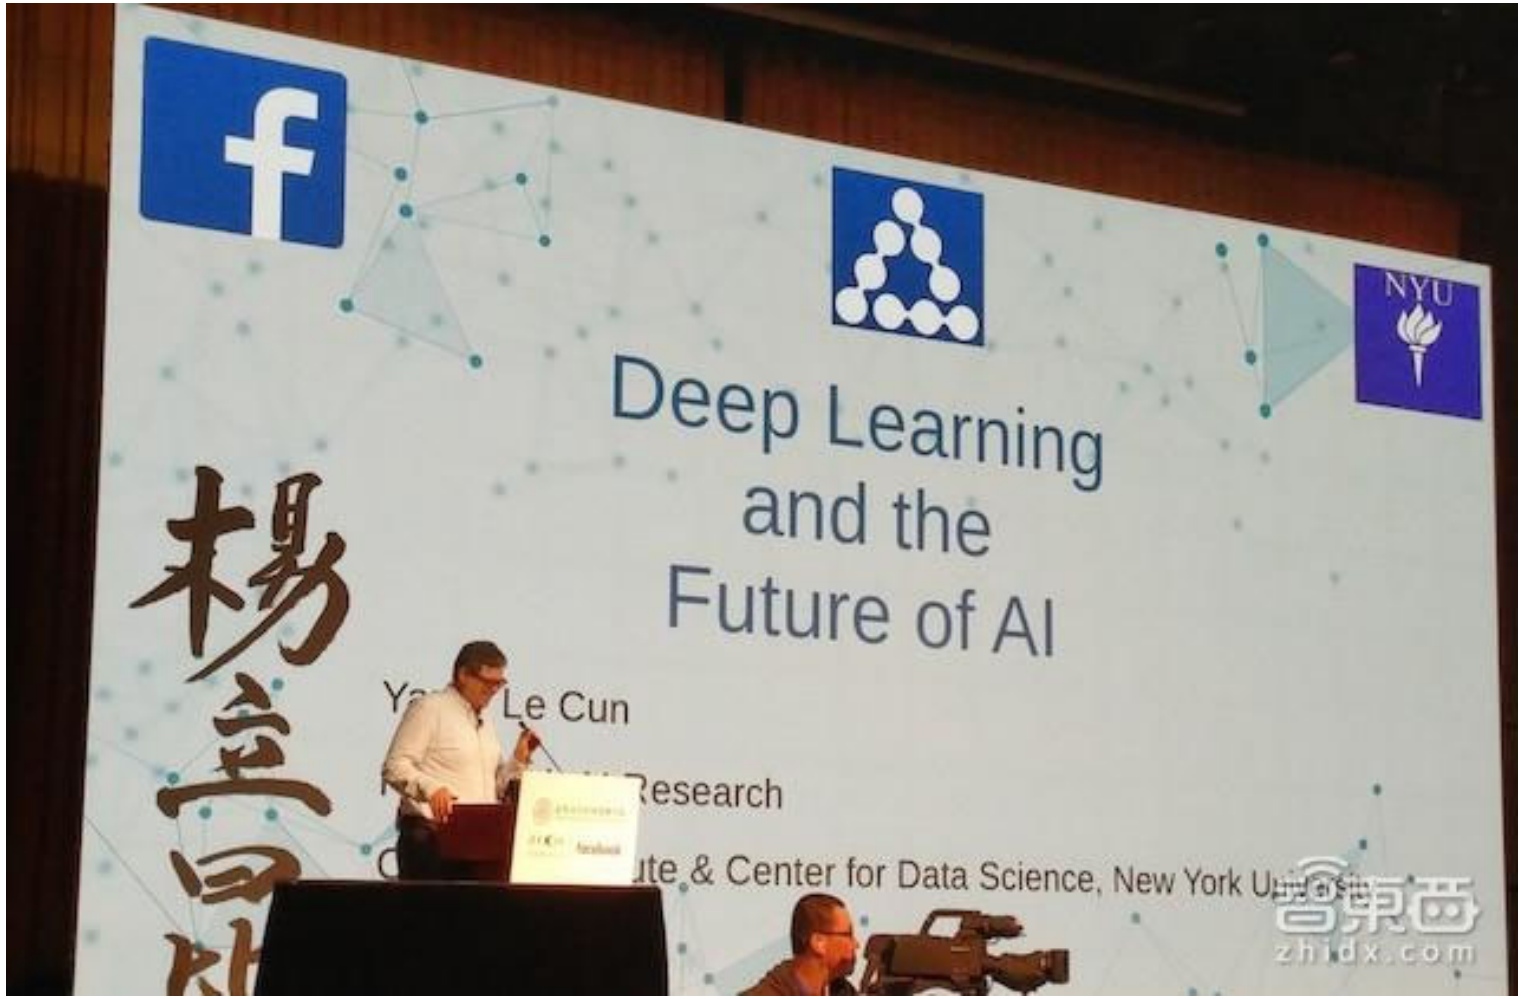

He gives himself a Chinese name, 杨立昆 (Mar. 2017)

Convolutional neural network

- Convolution
  - Computing the responses at hidden nodes is equivalent to convoluting the input image x with a learned filter w

$$net[i,j] = (x * w)[i,j] = \sum_m \sum_n x[m,n]w[i-m, j-n]$$

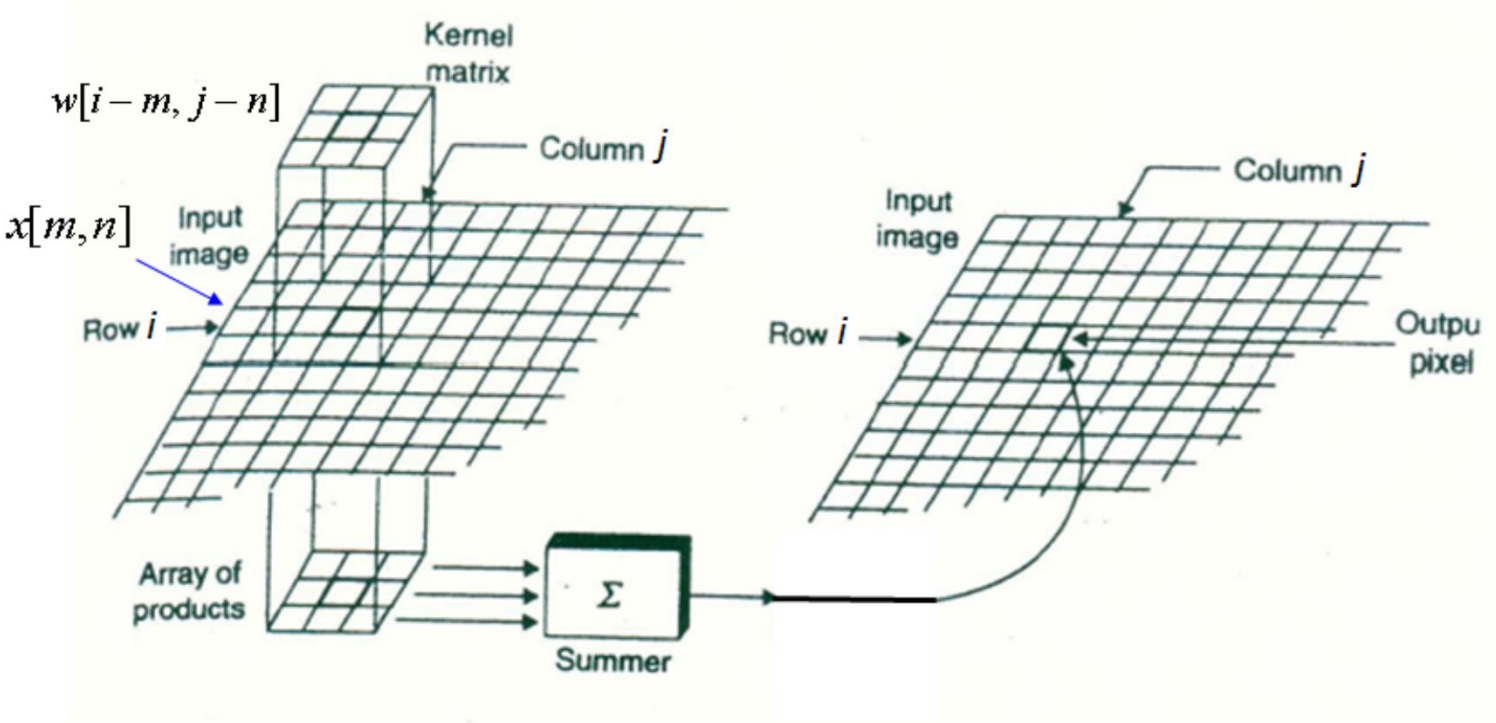

Convolutional neural network

- Downsampled convolution layer (optional)
  - To reduce computational cost, we may want to skip some positions of the filter and sample only every s pixels in each direction. A downsampled convolution function is defined as

$$net(i,j) = (\mathbf{x} * \mathbf{w})[i \times s, j \times s]$$

  - s is referred as the stride of this downsampled convolution
  - Also called as strided convolution

Convolutional neural network

- Multiple filters
  - Multiple filters generate multiple feature maps
  - Detect the spatial distributions of multiple visual patterns

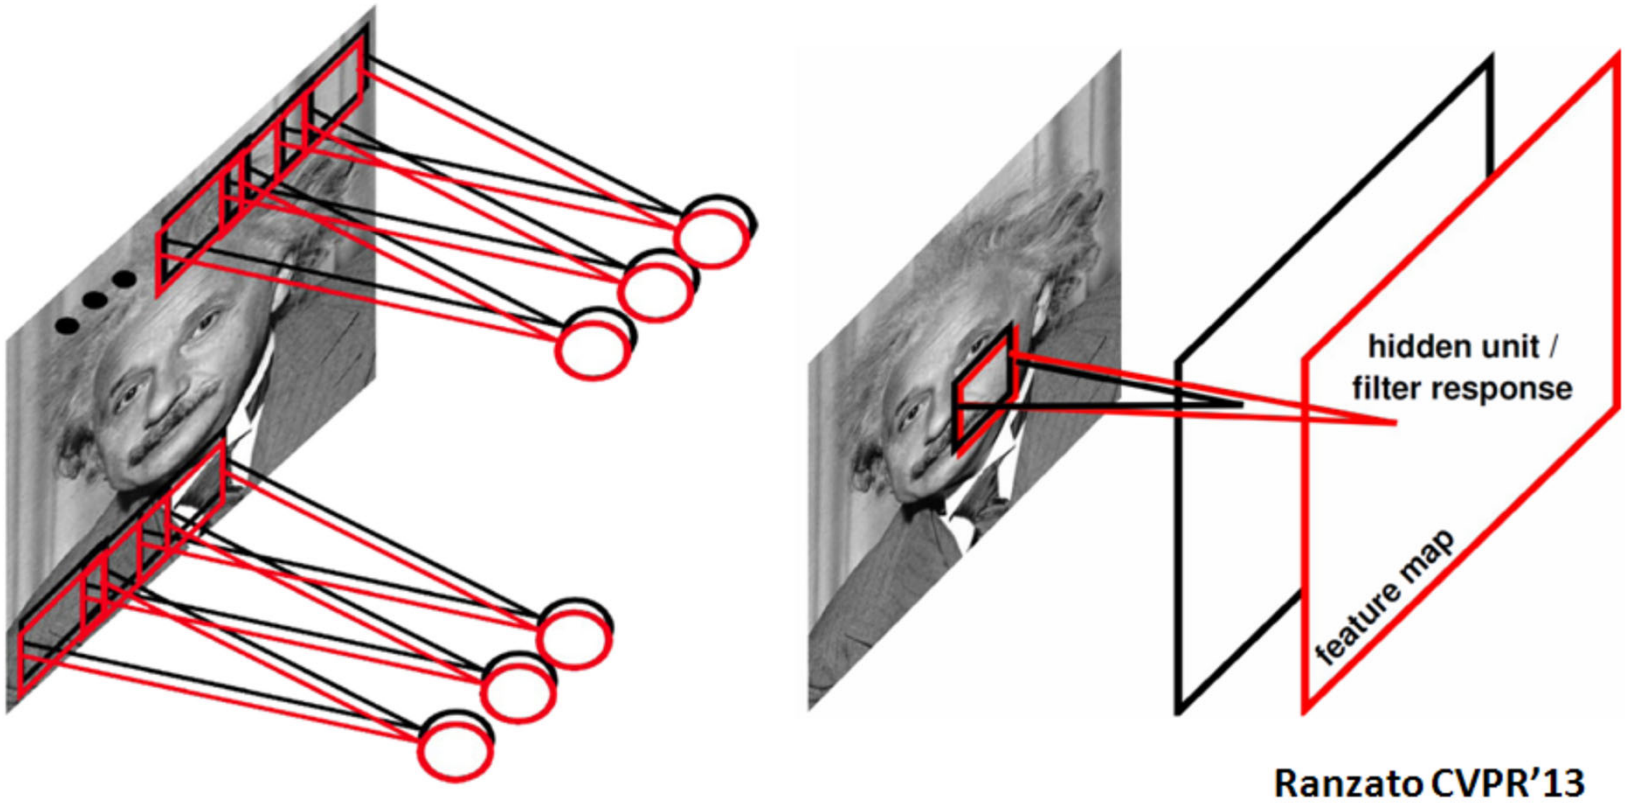

Convolutional neural network

- 3D filtering when input has multiple feature maps

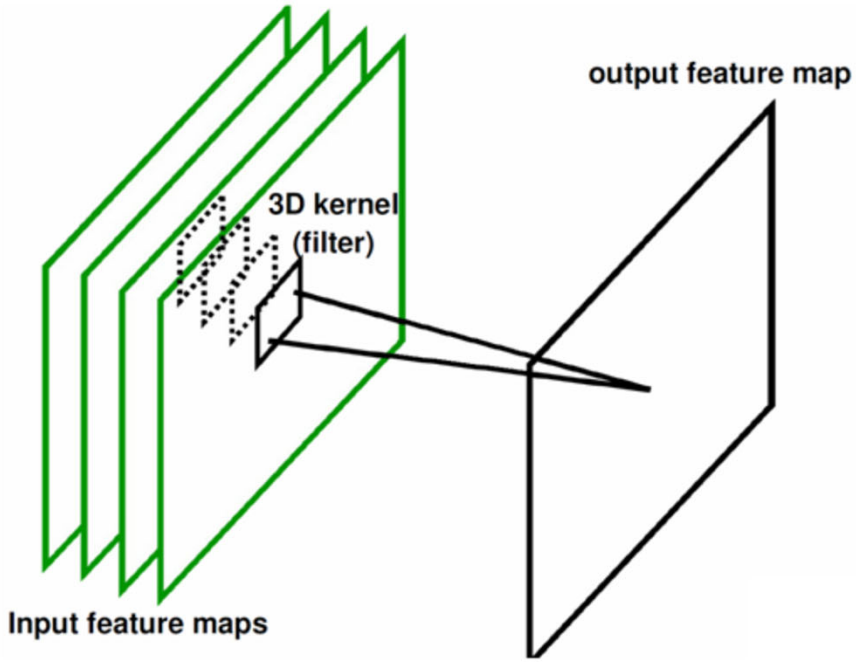

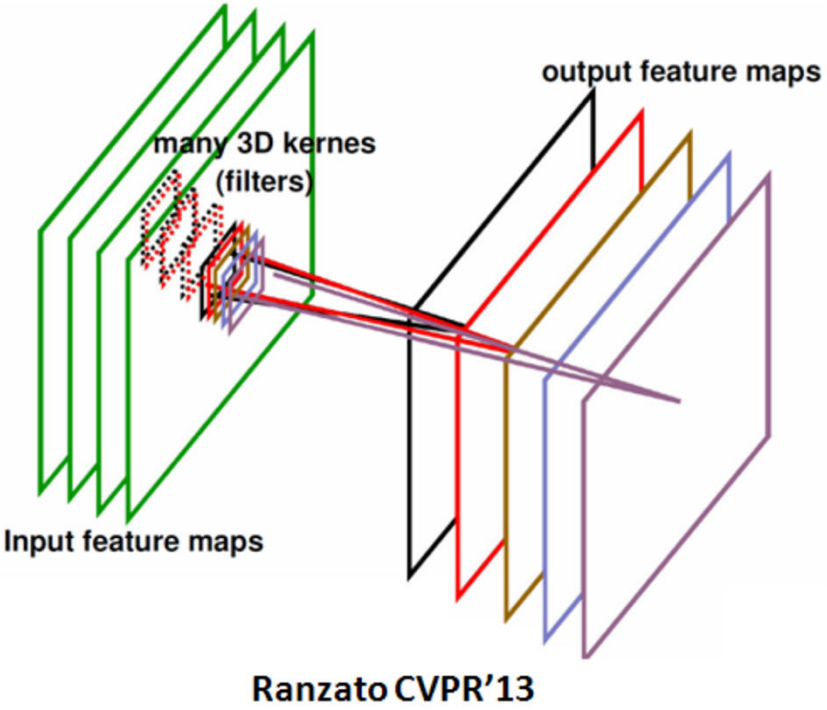

Convolutional neural network

- Convolutional layer

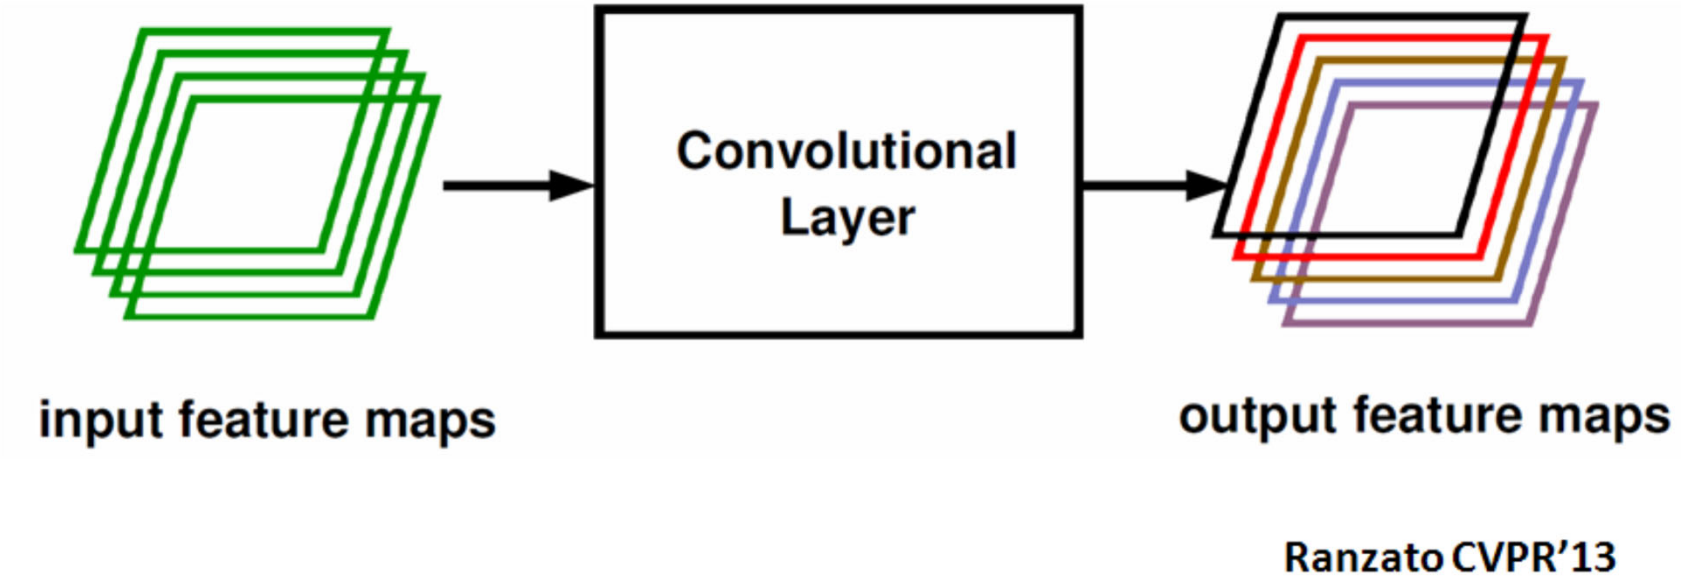

Convolutional neural network

- To the convolution responses, we then perform nonlinear activation
  - ReLU
  - Tanh
  - Sigmoid
  - Leaky ReLU
  - Softplus
  - SiLU

Convolutional neural network

- Local contrast normalization (optional)
  - Normalization can be done within a neighborhood along both spatial and feature dimensions

$$h_{i+1,x,y,k} = \frac{h_{i,x,y,k} - m_{i,N(x,y,k)}}{\sigma_{i,N(x,y,k)}}$$

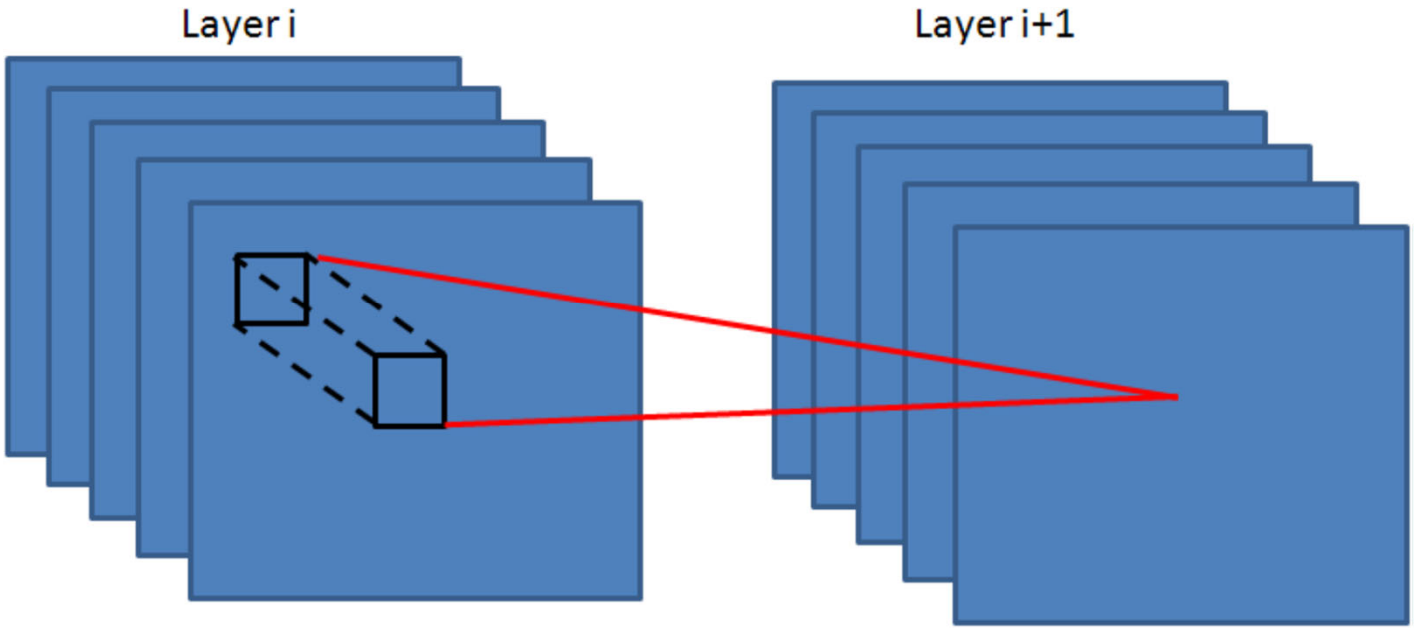

Convolutional neural network

- Then, we perform pooling
  - Max-pooling partitions the input image into a set of rectangles, and for each sub-region, outputs the maximum value
  - Non-linear down-sampling
  - The number of output maps is the same as the number of input maps, but the resolution is reduced
  - Reduce the computational complexity for upper layers and provide a form of translation invariance
  - Average pooling can also be used

Convolutional neural network

- Then, we perform pooling

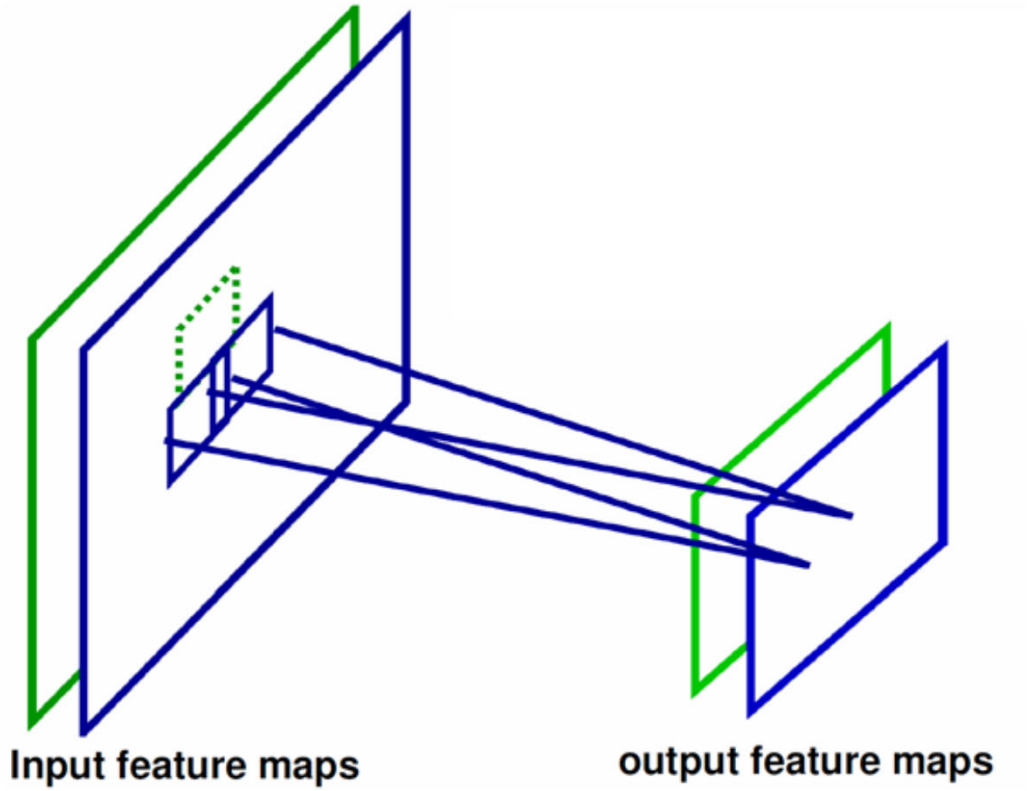

Ranzato CVPR'13

Convolutional neural network

- Typical architecture of CNN
  - Convolutional layer increases the number of feature maps
  - Pooling layer decreases spatial resolution
  - LCN and pooling are optional at each stage

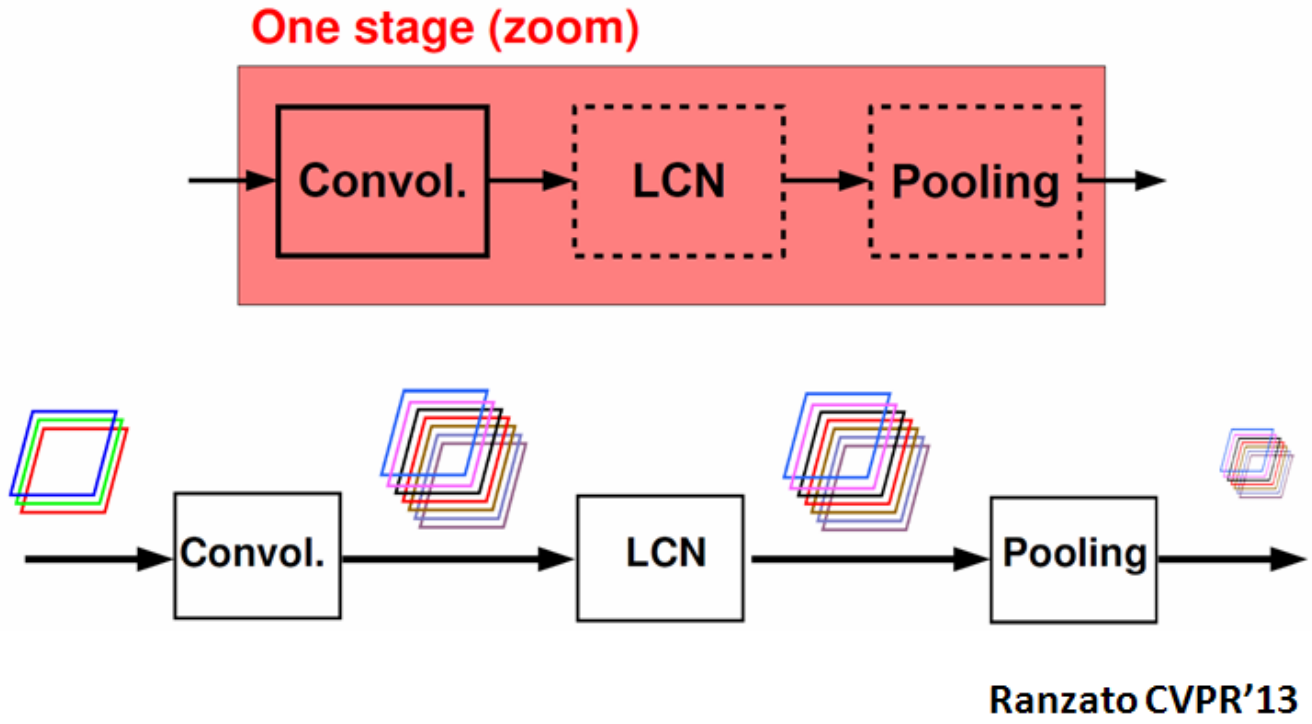

Ranzato CVPR'13

Convolutional neural network

- Typical architecture of CNN

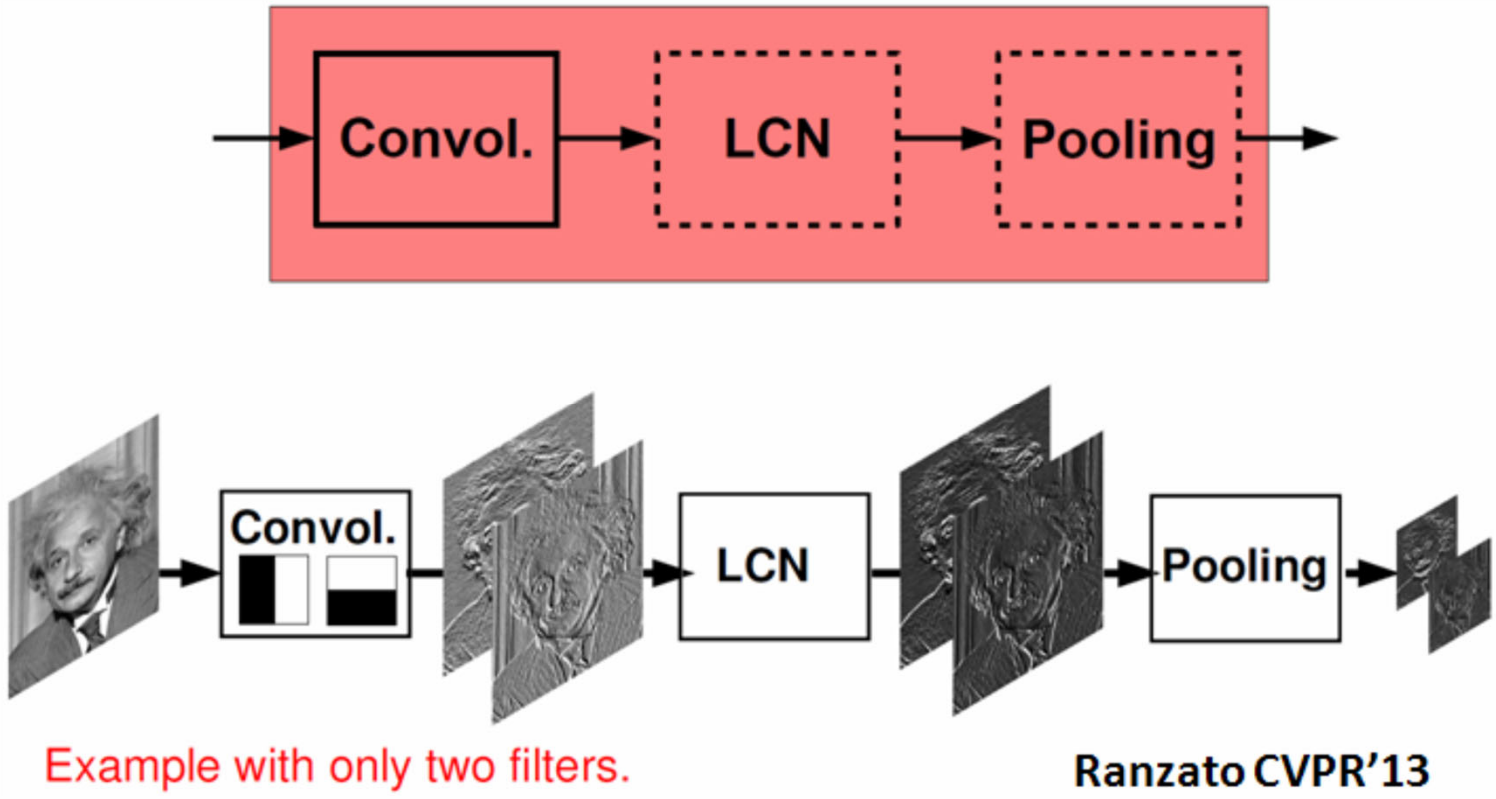

Convolutional neural network

- Typical architecture of CNN

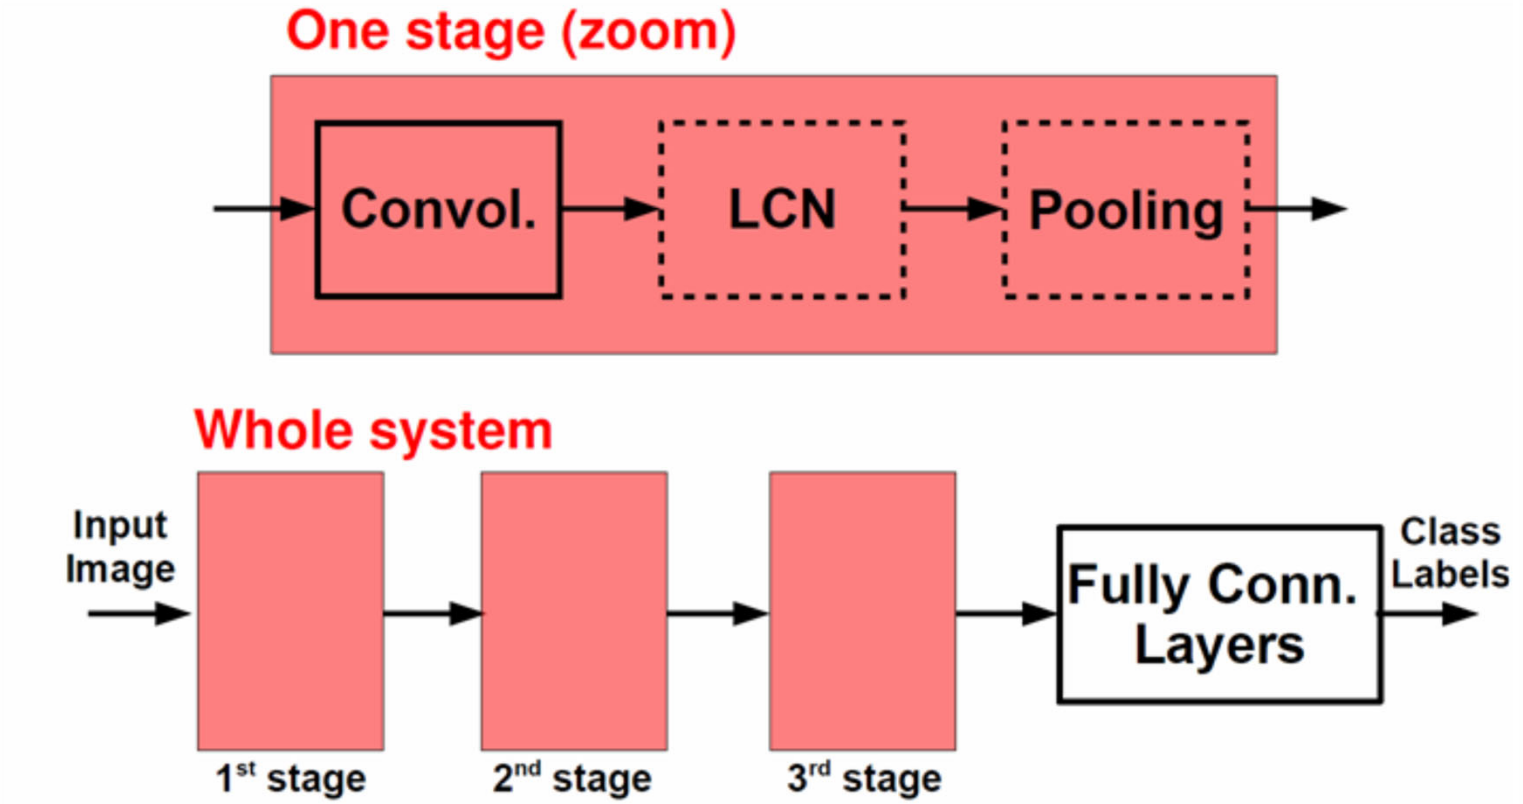

After a few stages, residual spatial resolution is very small.
We have learned a descriptor for the whole image.

Ranzato CVPR'13

Convolutional neural network

- Typical architecture of CNN

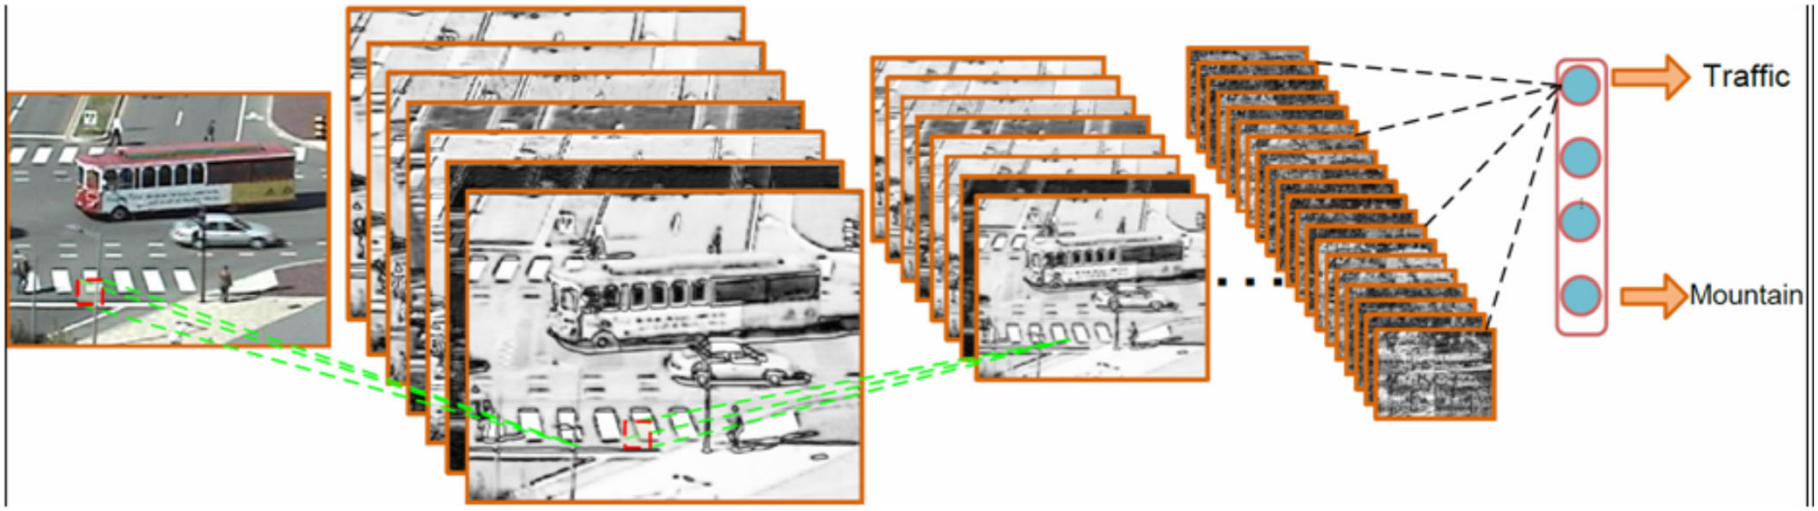

Convolutional neural network

- Some notes about the CNN layers in most recent net architectures
  - Spatial pooling (such as max pooling) is not recommended now. It is usually replaced by a strided convolution, allowing the network to learn its own spatial downsampling
  - Fully connected layers are not recommended now; instead, the last layer is replaced by global average pooling (for classification problems, the number of feature map channels of the last layer should be the same as the number of classes

Convolutional neural network

- Opensource platforms for CNN
  - CAFFE official, http://caffe.berkeleyvision.org/
  - Tensorflow, https://www.tensorflow.org/
  - Pytorch, www.pytorch.org/
  - Theano, http://deeplearning.net/software/theano/

Convolutional neural network

- An online tool for network architecture visualization
  - http://ethereon.github.io/netscope/quickstart.html
  - Network architecture conforms to the CAFFE prototxt format
  - The parameter settings and the output dimension of each layer can be conveniently observed

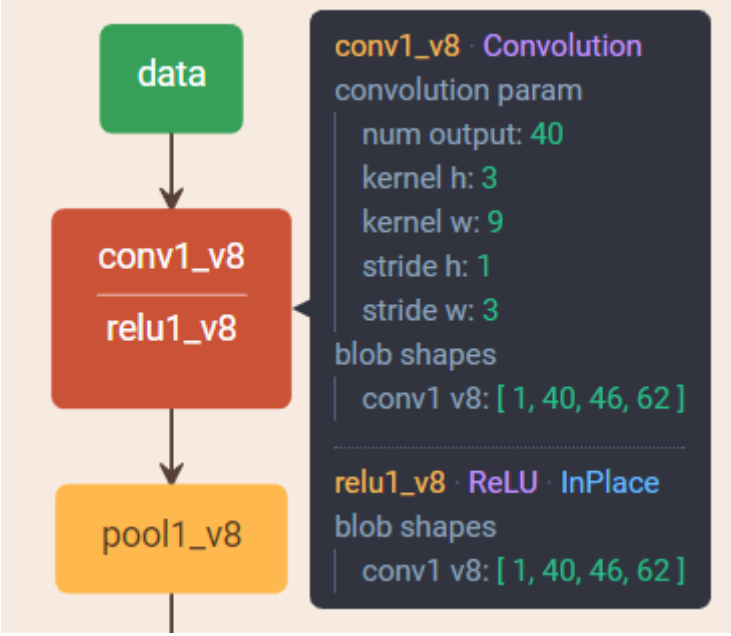

Outline

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
  - AlexNet
  - NIN
  - GoogLeNet
  - ResNet
  - DenseNet
  - PeleeNet
- CNN for object detection

AlexNet (NIPS 2012)

- AlexNet[7]: CNN for object recognition on ImageNet challenge
  - Trained on one million images of 1000 categories collected from the web with two GPU. 2GB RAM on each GPU. 5GB of system memory
  - Training lasts for one week
  - Google and Baidu announced their new visual search engines with the same technology six months after that
  - Google observed that the accuracy of their visual search engine was doubled

[7] A. Krizhevsky *et al*., ImageNet classification with deep convolutional neural networks, in Proc. NIPS, 2012

AlexNet (NIPS 2012)

- ImageNet
  - http://www.image-net.org/

AlexNet (NIPS 2012)

- Architecture of AlexNet
  - 5 convolutional layers and 2 fully connected layers for learning features
  - Max-pooling layers follow first, second, and fifth convolutional layers

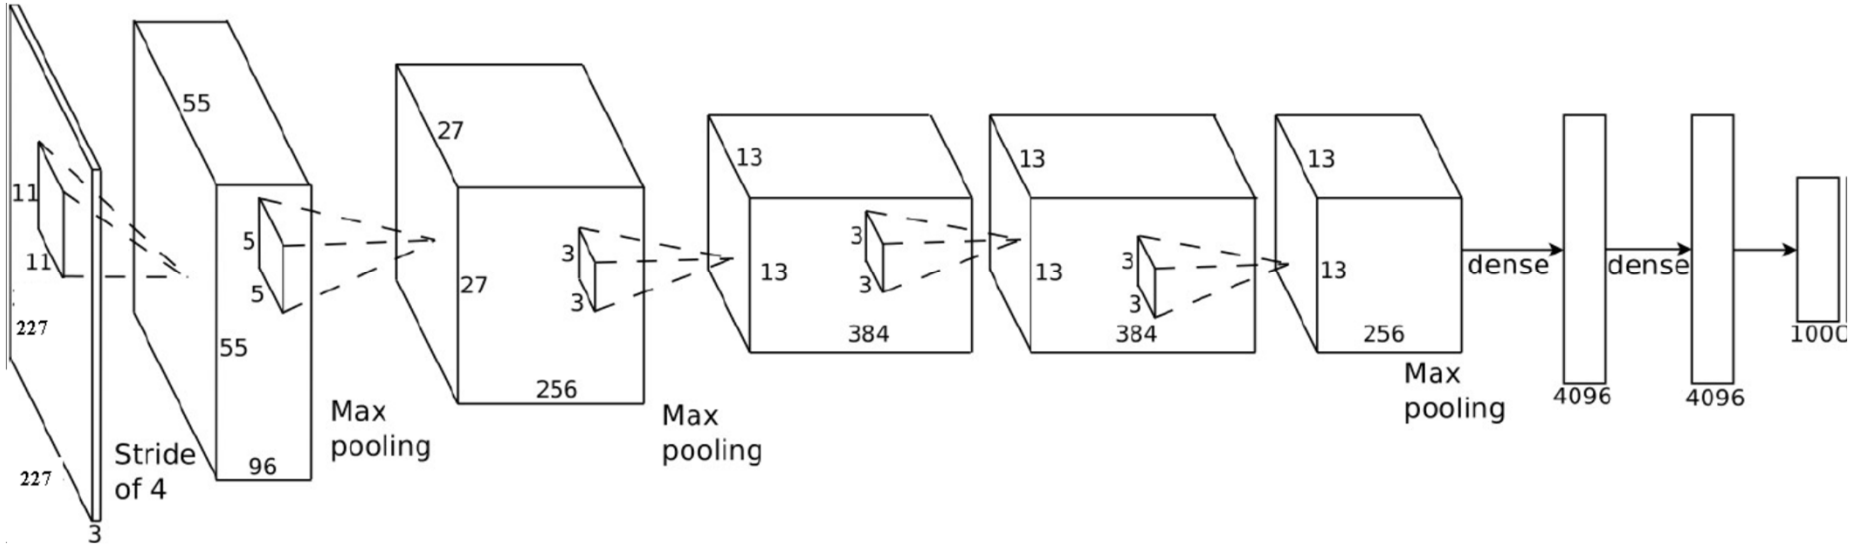

AlexNet (NIPS 2012)

- Architecture of AlexNet
  - The first time deep model is shown to be effective on large scale computer vision task
  - The first time a very large scale deep model is adopted
  - GPU is shown to be very effective on this large deep model

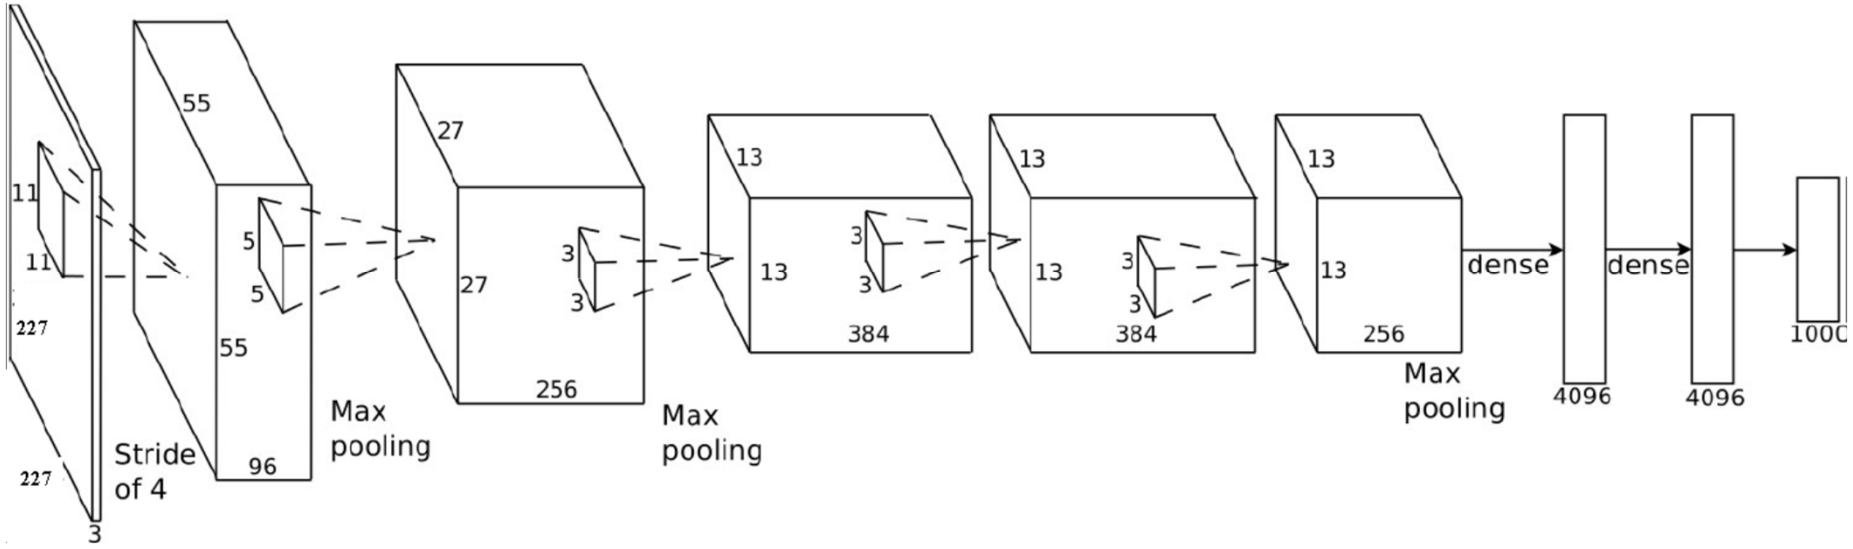

Network In Network (NIN, ICLR 2014)

- Main idea of NIN[8]
  - Conventional convolutional layers uses linear filters followed by a nonlinear activation function to abstract the information within a receptive field
  - Instead, NIN uses micro neural networks with more complex structures to abstract the data within the receptive field
  - The feature maps are obtained by sliding the micro network over the input in a similar manner as CNN
  - Moreover, they use global average pooling over feature maps in the classification layer, which is easier to interpret and less prone to overfitting than traditional fully connected layers

[8] M. Liu *et al*., Network in network, in Proc. ICLR, 2014

Network In Network (NIN, ICLR 2014)

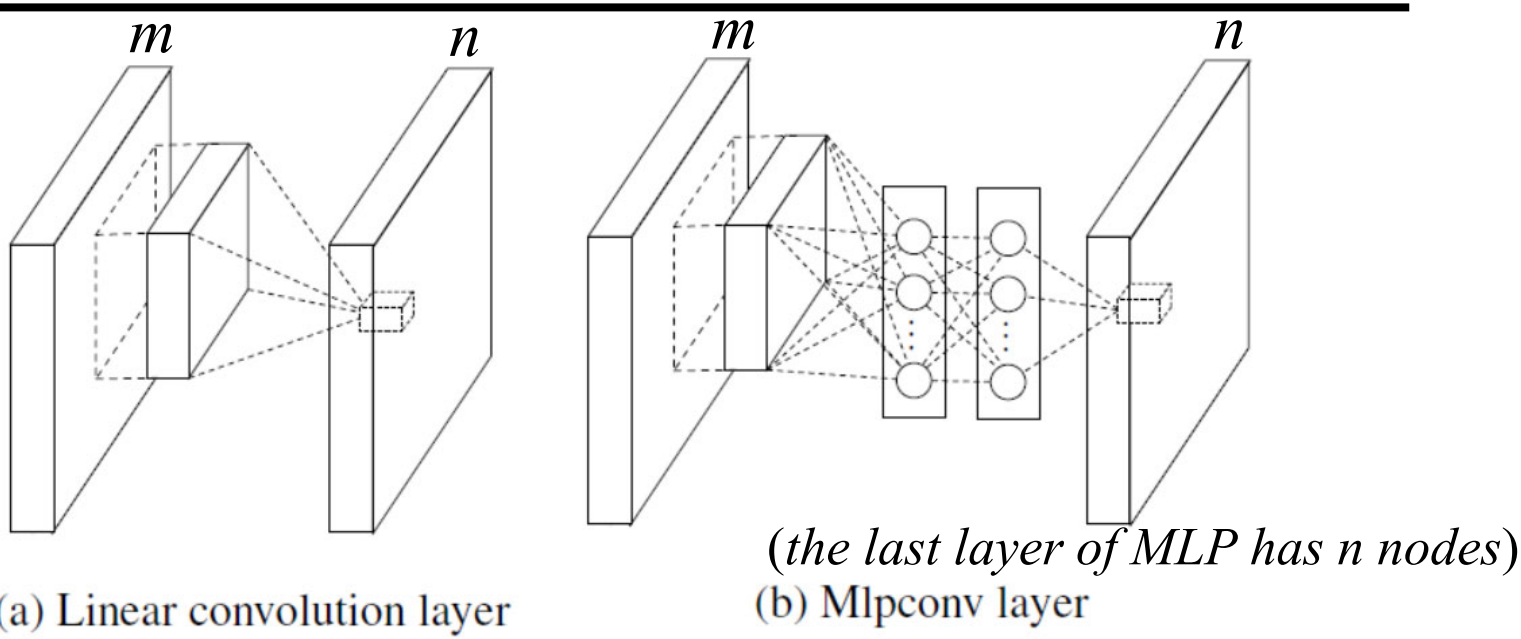

- Comparison of linear convolution layer and mlpconv layer
  - Both the layers map the local receptive field to an output feature vector
  - The mlpconv layer maps the input local patch to the output feature vector with a multilayer perceptron (MLP) consisting of multiple fully connected layers with nonlinear activation functions

Network In Network (NIN, ICLR 2014)

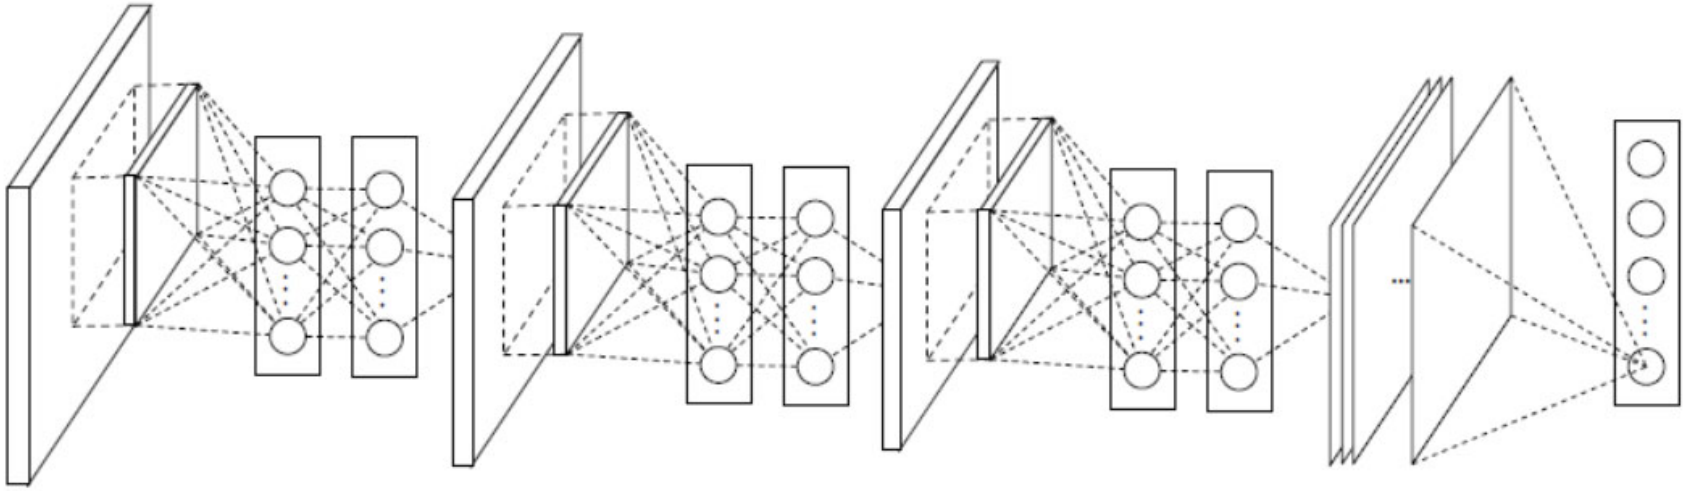

The overall structure of NIN. The last layer is the global average pooling

- More about global average pooling
  - Fully connected layers are prone to overfitting
  - If there are c classes, the last MLP layer should output c feature maps, one feature map for each corresponding category of the classification task
  - Take the average of each feature map to get a c dimensional vector for softmax classification

Network In Network (NIN, ICLR 2014)

- NIN can be implemented with conventional convolutional layers

For a mlpconv layer, suppose that the input feature map is of the size $m \times m \times 32$, the expected output feature map is of the size $m \times m \times 64$, the receptive field is 5×5 ; the mlpconv layer has 2 hidden layers, whose node numbers are 16 and 32, respectively.

How to implement this mlpconv layer with convolutional layers?

GoogLeNet[9] (CVPR 2015)

- Main idea: make the network deeper and wider, while keeping the number of parameters
- Inception module

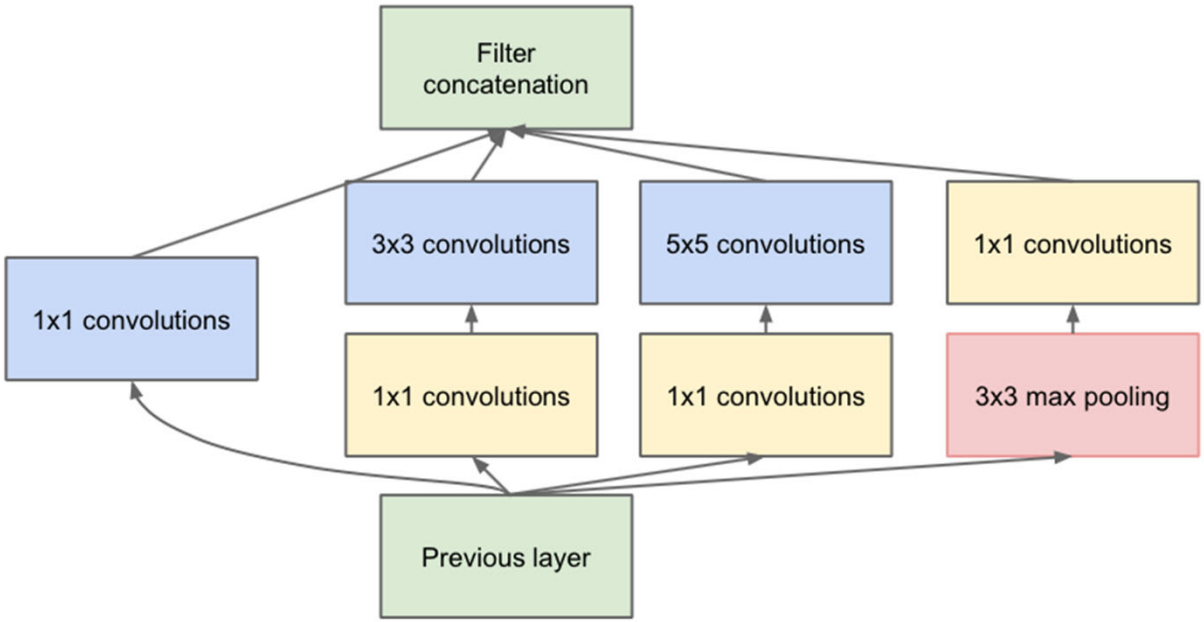

[9] C. Szegedy *et al*., Going deeper with convolutions, in Proc. CVPR, 2015

GoogLeNet (CVPR 2015)

- Many Inception modules can stack together to form a very deep network
- GoogLeNet refers to the version the authors submitted for the ILSVRC 2014 competition
  - This network consists 27 layers (including pooling layers)

ResNet[10] (CVPR 2016 Best Paper)

- What is the problem of stacking more layers using conventional CNNs?
  - Vanishing gradient, which can hamper the convergence
  - Accuracy get saturated, and then degraded

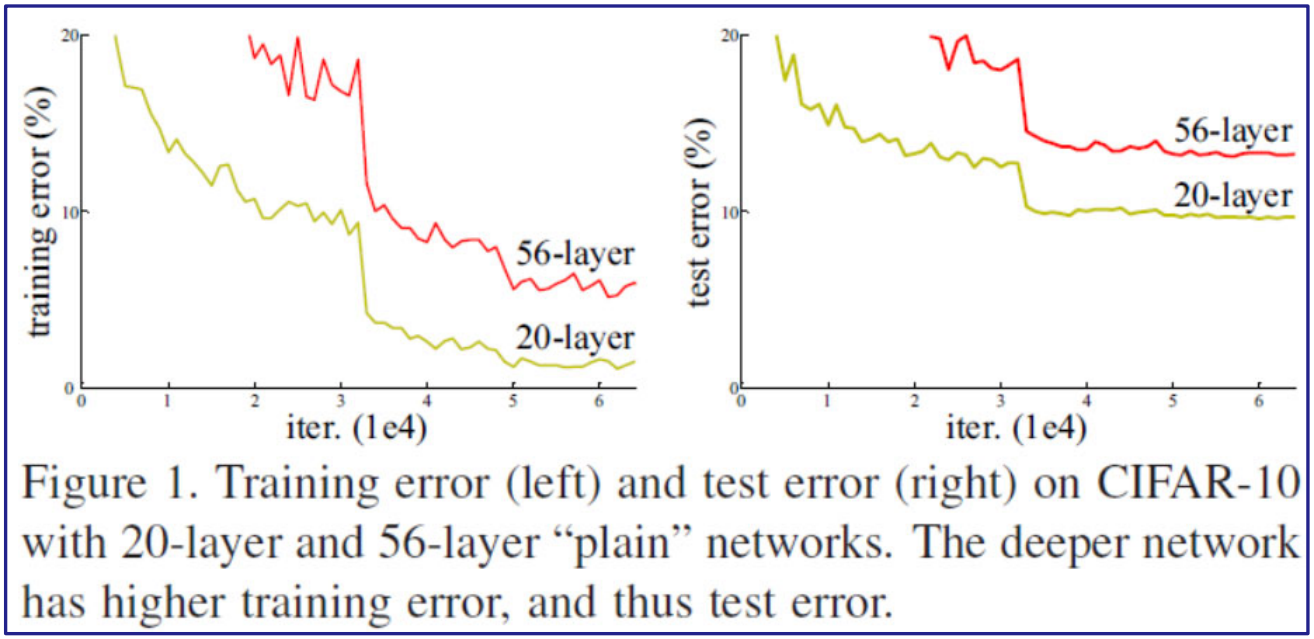

Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer "plain" networks. The deeper network has higher training error, and thus test error.

[10] K. He *et al*., Deep residual learning for image recognition, in Proc. CVPR, 2016

ResNet (CVPR 2016 Best Paper)

Is there any better way to design deeper networks?

Answer: Residual learning

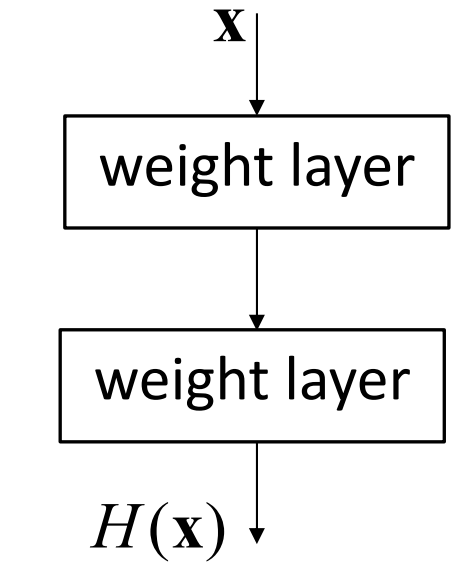

Conventional CNN

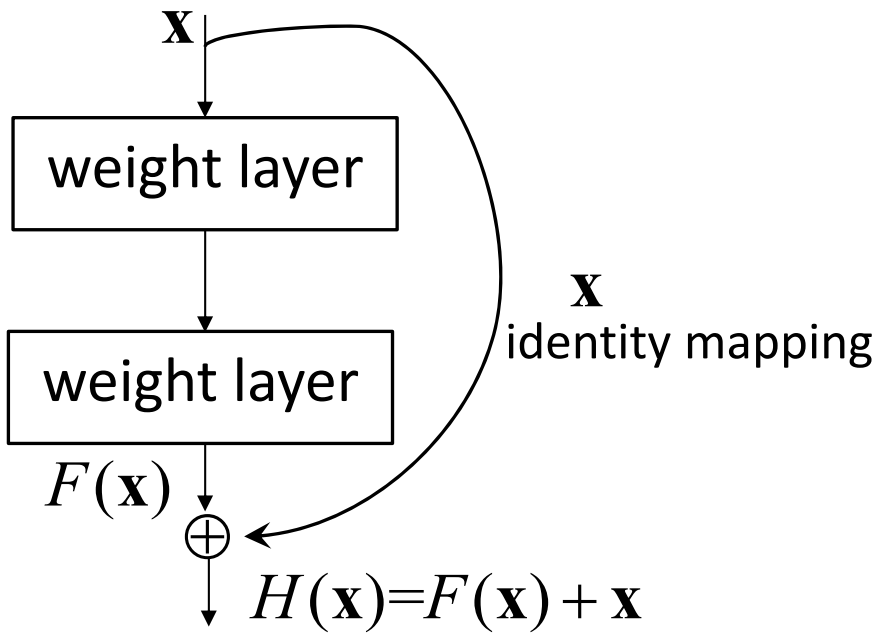

Residual block

ResNet (CVPR 2016 Best Paper)

- It is easier to optimize the residual mapping ($F(\mathbf{x})$) than to optimize the original mapping ($H(\mathbf{x})$)
- Identity mapping is implemented by shortcut
- A residual learning block is defined as,

$$\mathbf{y}=F(\mathbf{x},\{W_i\})+\mathbf{x}$$

where $\mathbf{x}$ and $\mathbf{y}$ are the input and output vectors of the layers $F+\mathbf{x}$ is performed by a shortcut connection and element-wise addition

Note: If the dimensions of F and $\mathbf{x}$ are not equal (usually caused by changing the numbers of input and output channels), a linear projection W_s (implemented with 1*1 convolution) is performed on $\mathbf{x}$ to match the dimensions,

$$\mathbf{y}=F(\mathbf{x},\{W_i\})+W_s\mathbf{x}$$

ResNet (CVPR 2016 Best Paper)

- It is easier to optimize the residual mapping ($F(\mathbf{x})$) than to optimize the original mapping ($H(\mathbf{x})$)
- Identity mapping is implemented by shortcut
- A residual learning block is defined as,

$$\mathbf{y}=F(\mathbf{x},\{W_i\})+\mathbf{x}$$

where $\mathbf{x}$ and $\mathbf{y}$ are the input and output vectors of the layers $F+\mathbf{x}$ is performed by a shortcut connection and element-wise addition

I highly recommend you to take a look the prototxt file of ResNet *(https://github.com/KaimingHe/deep-residual-networks)*

DenseNet[11] (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers
  - Within a dense block, connect all layers with each other in a feed-forward fashion
  - In contrast to ResNet, DenseNet combine features by concatenating them
  - The number of output feature maps of each layer is set as a constant within a dense block and is called as “growth rate”
  - Between two blocks, there is a transition layer, consisting of batch normalization, 1×1 convolution, and average pooling

[11] G. Huang *et al*., Densely connected convolutional networks, in Proc. CVPR, 2017

DenseNet (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers

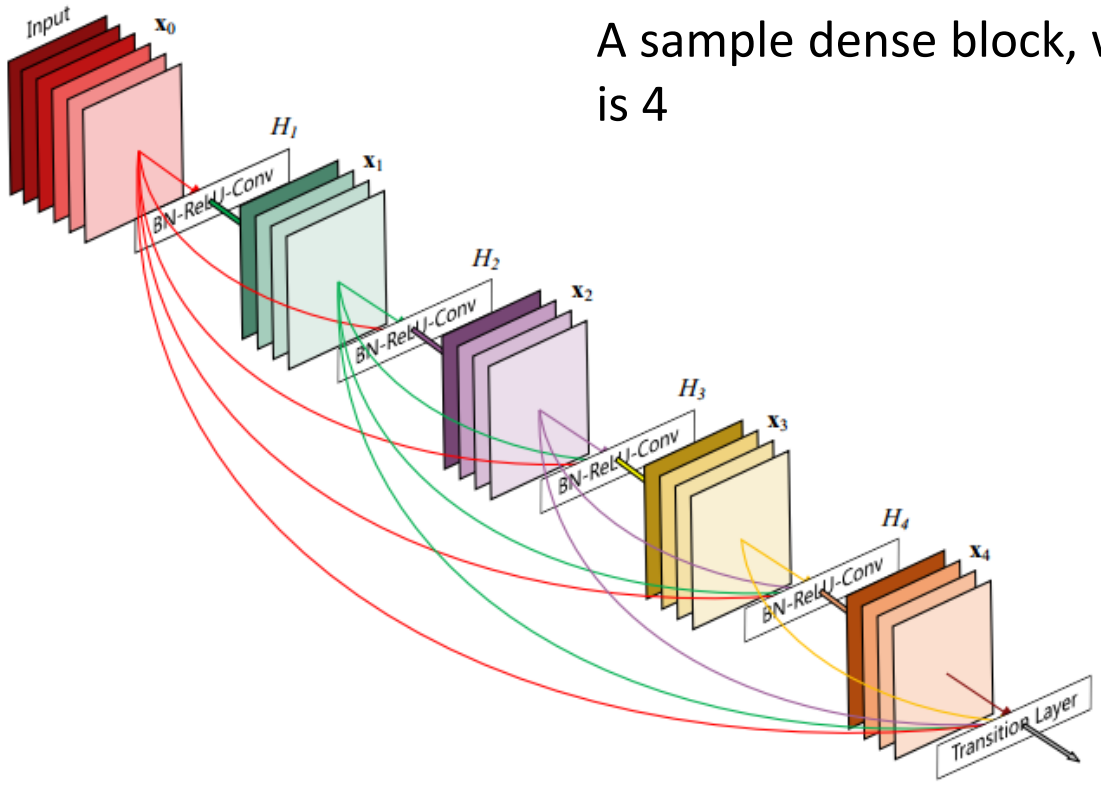

A sample dense block, whose growth rate is 4

DenseNet (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers

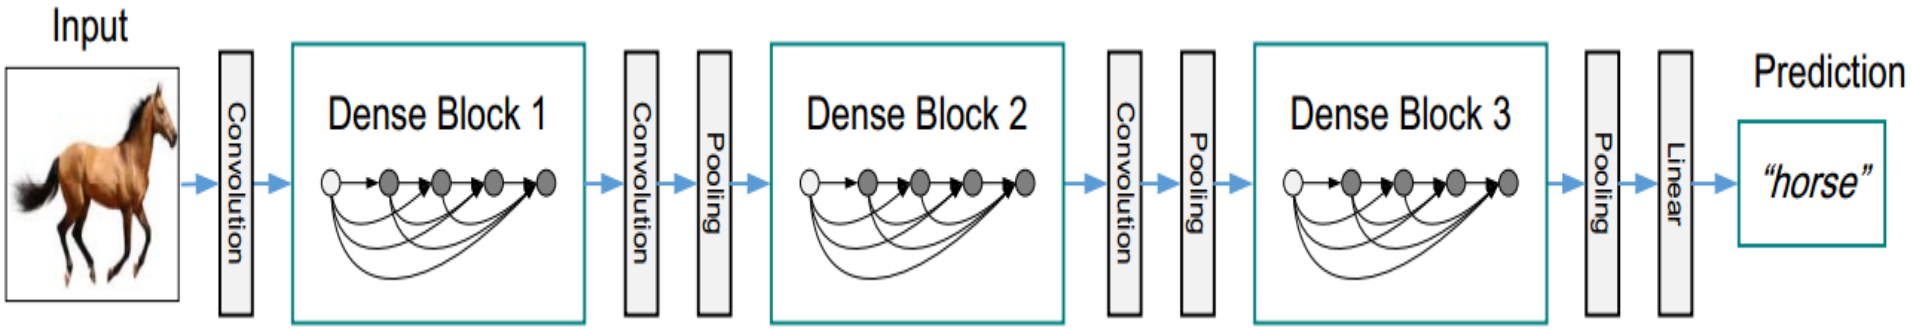

A sample DenseNet with three dense blocks

DenseNet (CVPR 2017 Best Paper)

- Highly motivated by ResNet
- A DenseNet comprises “dense blocks” and transition layers
- More details about DenseNet design
  - Bottleneck layers. A 1×1 convolution layer can be introduced as bottleneck layer before each 3×3 convolution to reduce the number of input feature maps, and thus to improve computational efficiency
  - Compression. If a dense block contains m feature maps, we let the following transition layer generate θm output feature maps where $0<\theta\leq1$ is referred to as the compression factor

PeleeNet[12] (ICLR Workshop 2018)

- MobileNet and ShuffleNet are designed to run on mobile devices with limited computing power and memory resource
- However, they both depend on depthwise separable convolution which lacks efficient implementation in most deep learning frameworks
- PeleeNet
  - based on conventional convolution instead
  - follows the innovate connectivity pattern and some of key design principals of DenseNet
  - achieves a higher accuracy by 0.6% and 11% lower computational cost than MobileNet
  - PeleeNet is only 66% of the model size of MobileNet

[12] J. Wang *et al*., Pelee: A Real-Time Object Detection System on Mobile Devices, in Proc. ICLR Workshop, 2018

PeleeNet (ICLR Workshop 2018)

- PeleeNet's design
  - Two-Way dense layer

  Motivated by GoogLeNet, use a 2-way dense layer to get different scales of receptive fields. One way of the layer uses a small kernel size (3x3), which is good enough to capture small-size objects. The other way of the layer uses two stacked 3x3 convolution to learn visual patterns for large objects

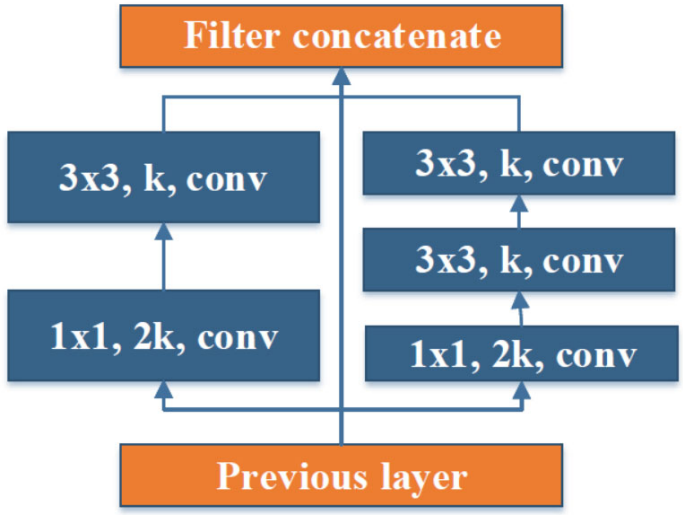

2-way dense layer

PeleeNet (ICLR Workshop 2018)

- PeleeNet's design
  - Two-Way dense layer
  - Stem block

  The stem block can effectively improve the feature expression ability without adding computational cost too much - better than other more expensive methods, e.g., increasing channels of the first convolution layer

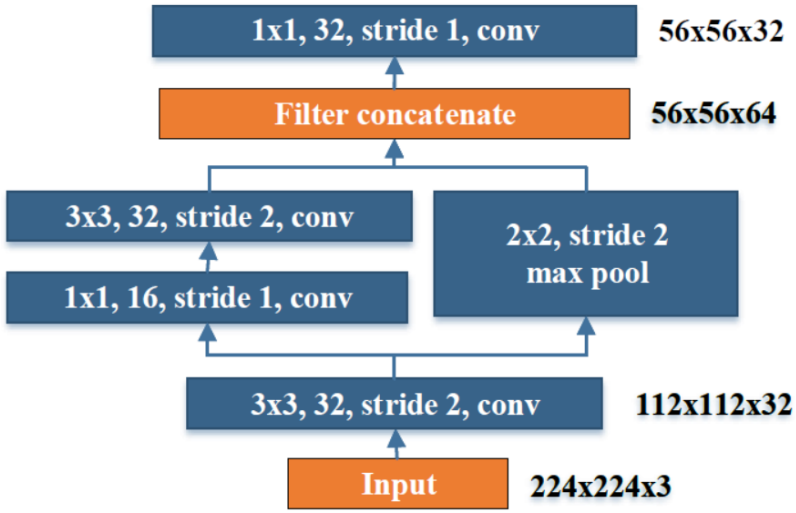

stem block

PeleeNet (ICLR Workshop 2018)

- PeleeNet's design
  - Two-Way dense layer
  - Stem block
  - Dynamic number of channels in bottleneck layer
  - Transition layer without compression
  - Composite Function

  To improve actual speed, they use the conventional wisdom of post activation (Convolution - Batch Normalization - Relu) as their composite function instead of pre-activation used in DenseNet

Outline

- Basic concepts
- Linear model
- Neural network
- Convolutional neural network (CNN)
- Modern CNN architectures
- CNN for object detection

Background

- Detection is different from classification
  - An image classification problem is predicting the label of an image among the predefined labels; It assumes that there is **single** object of interest in the image and it covers a significant portion of image
  - Detection is about not only finding the class of object but also localizing the extent of an object in the image; the object can be lying anywhere in the image and can be of any size (scale)

Background

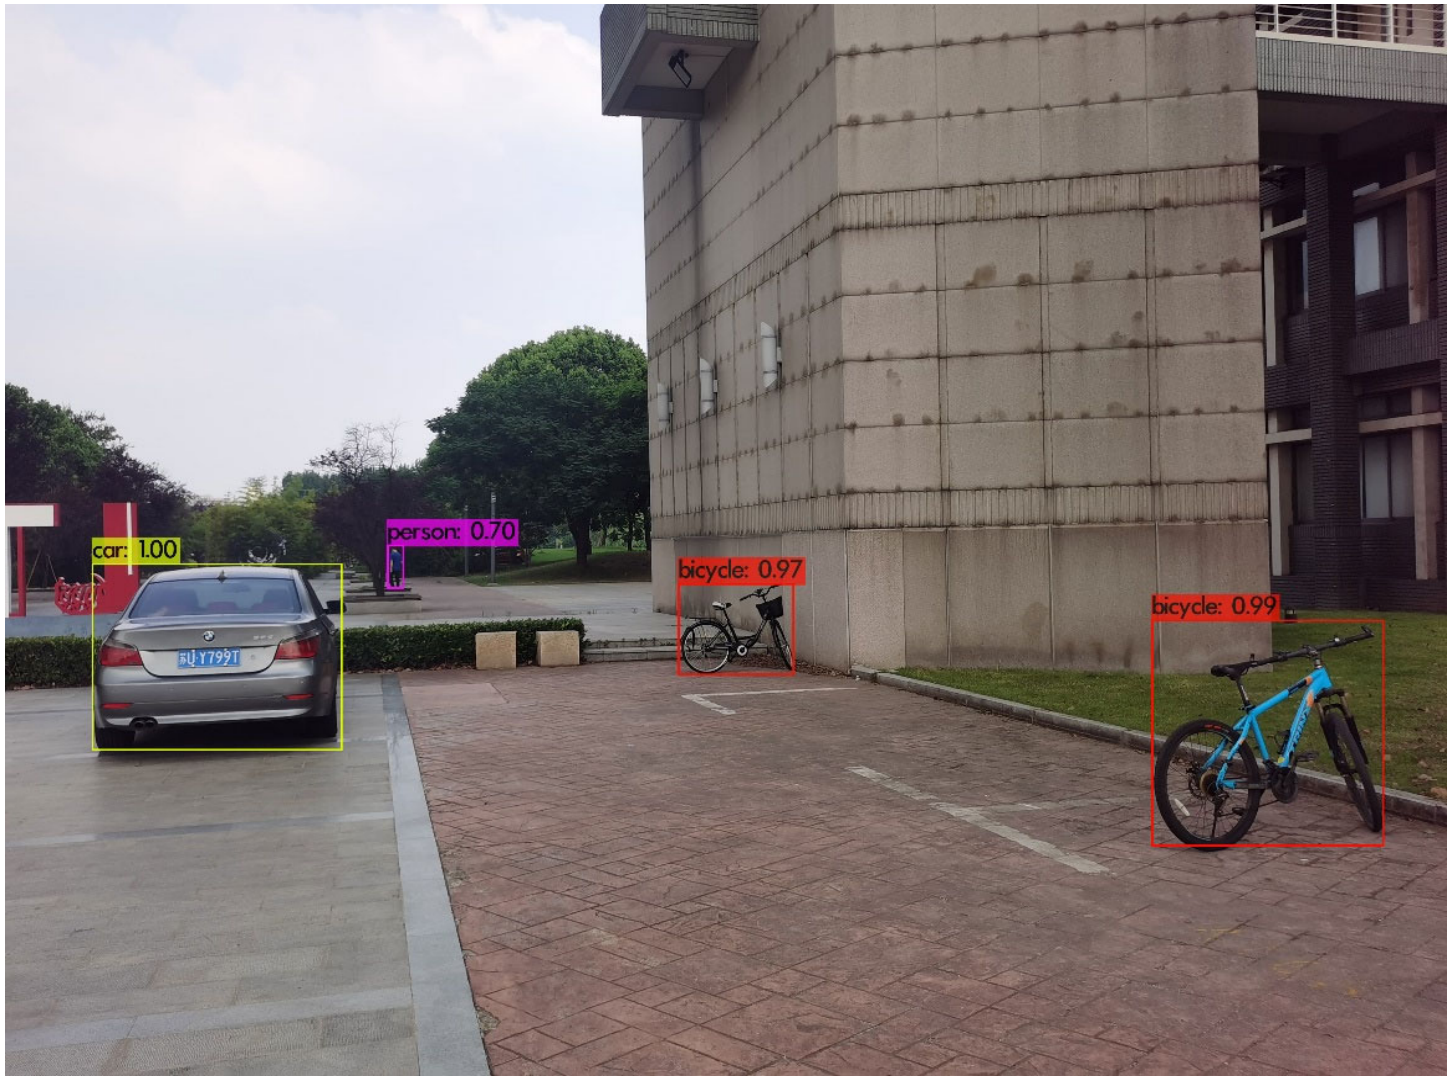

Multiple objects detection

Background

- Traditional methods of detection involved using a block-wise orientation histogram (SIFT or HOG) feature which could not achieve high accuracy in standard datasets such as PASCAL VOC; these methods encode a very low level characteristics of the objects and therefore are not able to distinguish well among the different labels
- Deep learning based methods have become the state-of-the-art in object detection in image; they construct a representation in a hierarchical manner with increasing order of abstraction from lower to higher levels of neural network

Background

- Recent developments of CNN based object detectors
  - R-CNN series
    - » R-CNN (CVPR 2014), FastRCNN (ICCV 2015), FasterRCNN (NIPS 2015), MaskRCNN (ICCV 2017)
  - SSD (single-shot multibox detector) (ECCV 2016)
  - Pelee-SSD (ICLR Workshop 2018)
  - YOLO series
    - » YOLOv1 (CVPR 2016), YOLOv2 (CVPR 2017), …., YOLOv8 (Jan. 2023)
  - EfficientDet (CVPR 2020)

R-CNN

- Brute-force idea
  - One could perform detection by carrying out a classification on different sub-windows or patches or regions extracted from the image. The patch with high probability will not only the class of that region but also implicitly gives its location too in the image
  - One brute force method is to run classification on all the sub-windows formed by sliding different sized patches (to cover each and every location and scale) all through the image

Quite Slow!!!

R-CNN

- R-CNN therefore uses an object proposal algorithm (selective search) in its pipeline which gives out a number (~2000) of TENTATIVE object locations
- These object regions are warped to fixed sized (227X227) regions and are fed to a classification convolutional network which gives the individual probability of the region belonging to background and classes

Region-based Convolution Networks (R-CNNs)

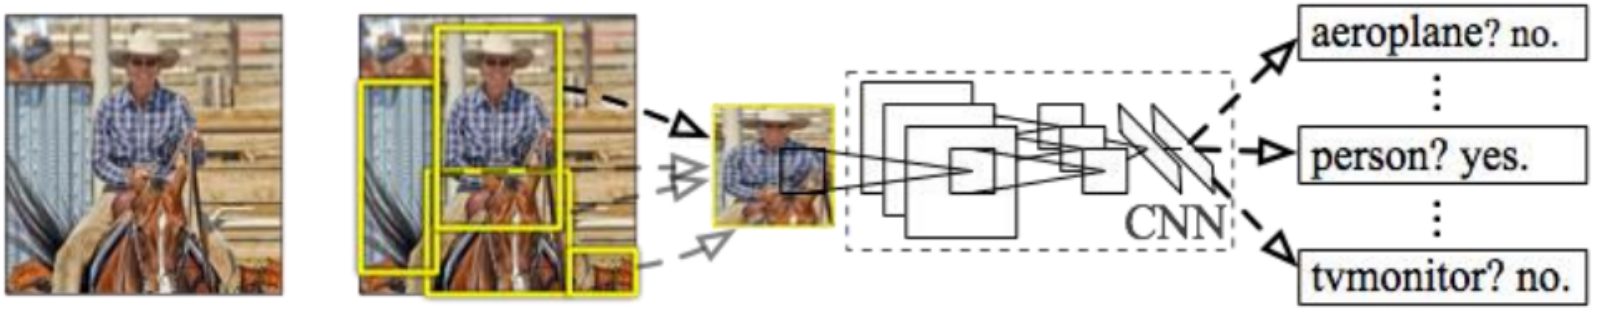

Input image

Extract region proposals (~2k / image)
e.g., selective search
[van de Sande, Uijlings et al.]

Compute CNN features on regions

Classify and refine regions

Fast-RCNN

- Compared to RCNN
  - A major change is a single network with two loss branches pertaining to soft-max classification and bounding box regression
  - This multitask objective is a salient feature of Fast-RCNN as it no longer requires training of the network independently for classification and localization

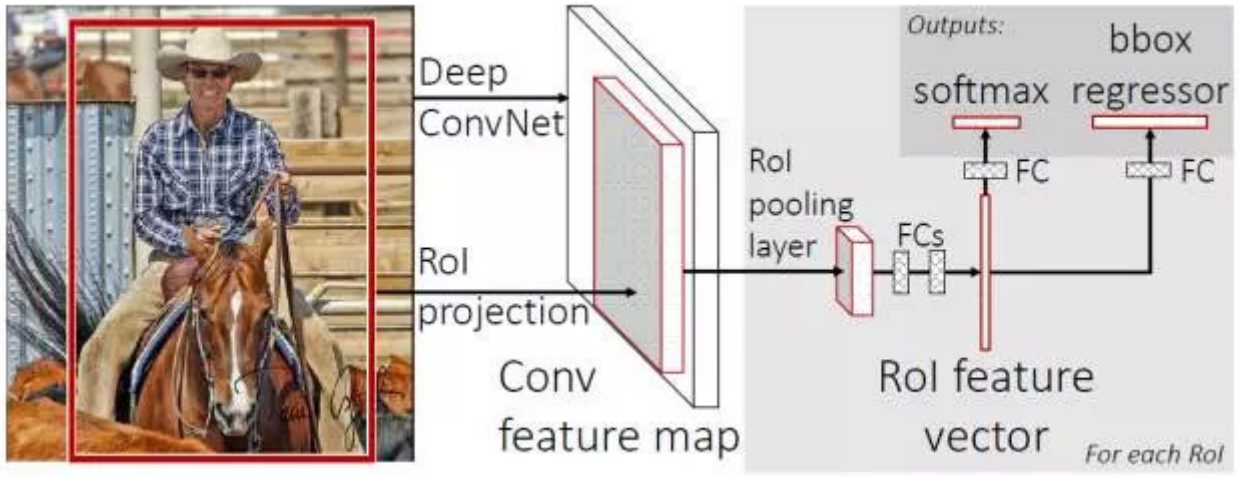

Faster-RCNN

- Compared to Fast-RCNN
  - Faster-RCNN replaces Selective Search with CNN itself for generating the region proposals (called RPN-region proposal network) which gives out tentative regions at almost negligible amount of time

YOLO series

- YOLO (You Only Look Once)
  - The major exceptional idea is that it tackles the object detection problem as a regression problem
  - A single neural network predicts bounding boxes and class probabilities directly from full images in one evaluation
  - The whole detection pipeline is a single network
  - It is extremely fast

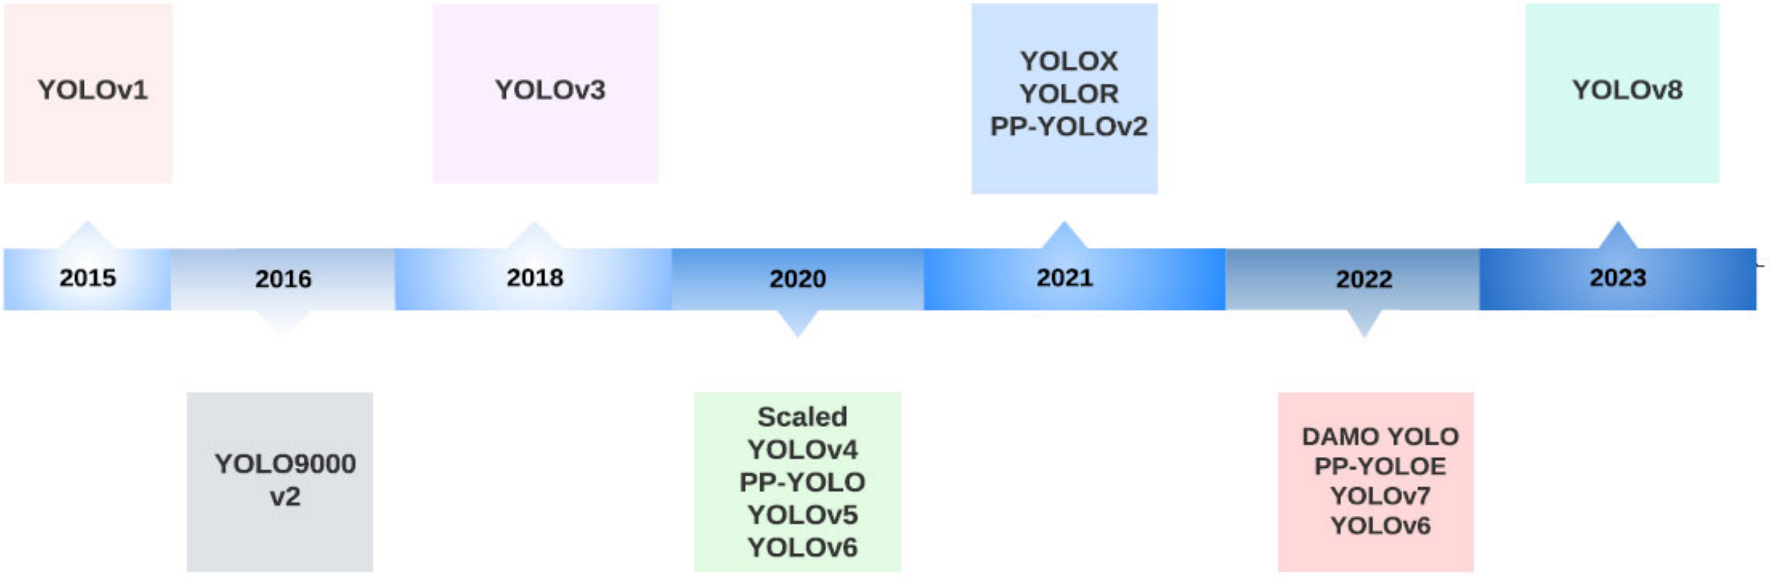

Legends about YOLO

- The original idea of YOLO was proposed by Joseph Redmon in 2016[12]
- Then he improved it as YOLOv2 and YOLOv3 in 2017 and 2018
- However, on Feb. 21, 2020, Redmon announced that he would stop doing research work in CV (*I stopped doing CV research because I saw the impact my work was having. I loved the work but the military applications and privacy concerns eventually became impossible to ignore*), so all the following YOLOs actually do not have "Redmon" in the author lists
- Now from YOLOv1~YOLOv4, the models are trained and inferenced on darknet (a C language based DCNN platform), which is now mainly maintained by a Russian researcher Alexey Bochkovskiy; From YOLOv5, the models are developed by Python and PyTorch
- YOLOv5 was developed by a company, namely Ultralytics, founded by Glenn R. Jocher and on Jan. 2023, they announced YOLOv8

[12] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You only look once: Unified, real-time object detection," *Proceedings of CVPR*, pp. 779–788, 2016

YOLOv1

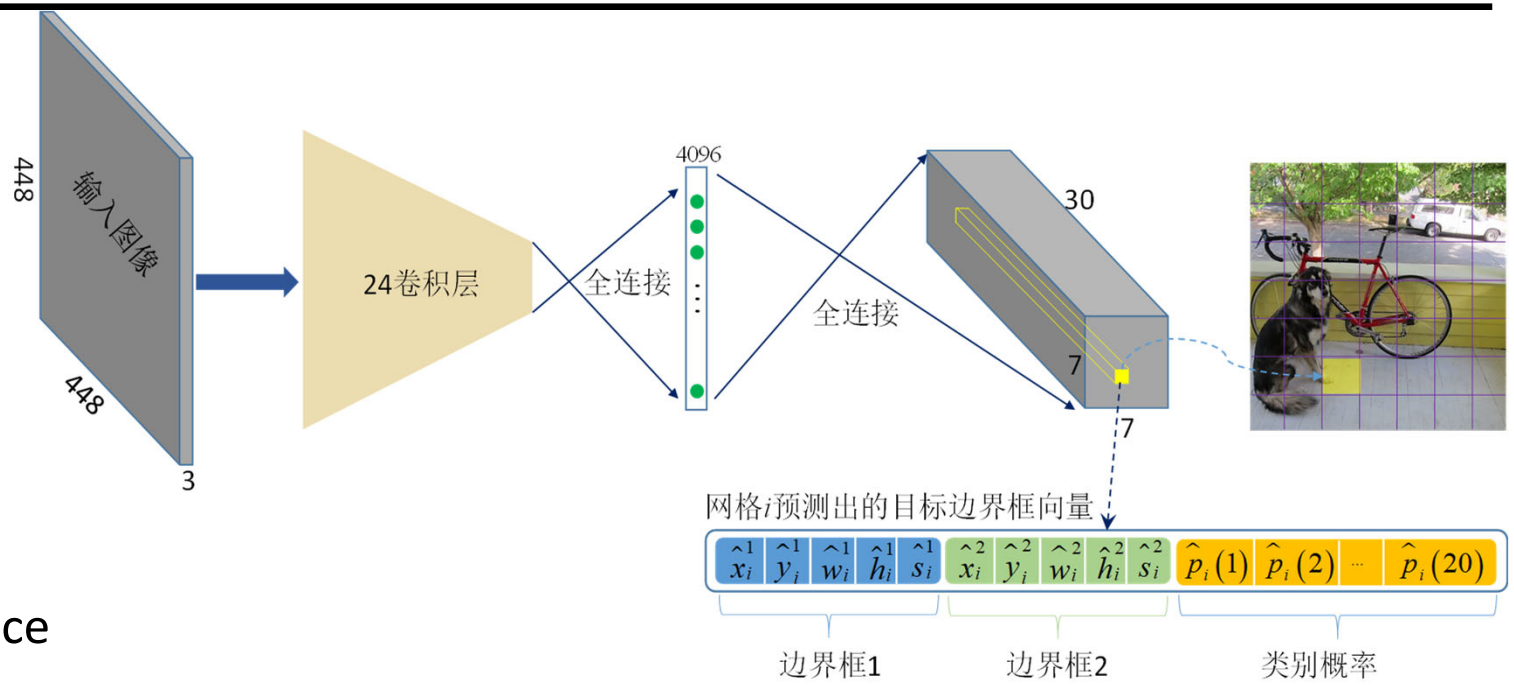

For inference

- The results are derived from the $7 \times 7 \times 30$ output matrix
- Conceptually, the input image is divided in to 7×7 cells and each cell predicts a 30*d* vector
  - Each output vector contains two bounding-boxes (whose centers are relative to the cell's top-left corner and the sizes are relative to the size of the image resolution) and classification probabilities of 20 classes (the default model was trained on VOC which has 20 classes)
  - The final confidence score of the 1st bounding-box in cell i is

$$\hat{s}_i^1 \cdot \hat{p}_i(cls) \quad \text{where} \quad cls = \arg\max_c \left\{ \hat{p}_i(c) : | \, c \in \{20 \, classes\} \right\}$$

- Finally, for each class, perform NMS for the bounding-boxes belonging to that class

YOLOv1

算法 15-1：目标边界框初始集合 $\mathcal{B}$ 的非极大值抑制

输入：目标边界框初始集合$\mathcal{B}$，分类检测置信度集合$\mathcal{S}$，IoU 阈值 τ，分类检测置信度阈值 T

输出：经过 NMS 操作之后的目标边界框集合$\mathcal{F}$

1) $\mathcal{F} \leftarrow \phi$
2) Filter the bounding boxes: $\mathcal{B} \leftarrow \{b \in \mathcal{B} \mid S(b) \geq T\}$
3) Sort the bounding boxes in $\mathcal{B}$ by their confidence scores in descending order
4) **while** $\mathcal{B} \neq \phi$ **do**
5) Select the bounding box b from $\mathcal{B}$ with the highest confidence score
6) Add b to the set of final bounding boxes $\mathcal{F}$: $\mathcal{F} \leftarrow \mathcal{F} \cup \{b\}$
7) Remove b from $\mathcal{B}$: $\mathcal{B} \leftarrow \mathcal{B} - \{b\}$
8) **for** every remaining bounding box r in $\mathcal{B}$ **do**
9) Calculate the IoU between b and r: $iou \leftarrow IoU(b, r)$
10) **if** $iou \geq \tau$
11) Remove r from $\mathcal{B}$: $\mathcal{B} \leftarrow \mathcal{B} - \{r\}$
12) **end if**
13) **end for**
14) **end while**

$$IoU(A,B) = \frac{\text{intersection area between } A \text{ and } B}{\text{union area of } A \text{ and } B} =$$

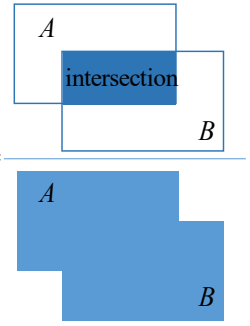

YOLOv1

For training

- Conceptually, the input image is divided in to 7×7 cells
- For cell i, if one ground-truth object's center falls in i, we say cell i links to a ground-truth object and is responsible for the loss computation related to that ground-truth object

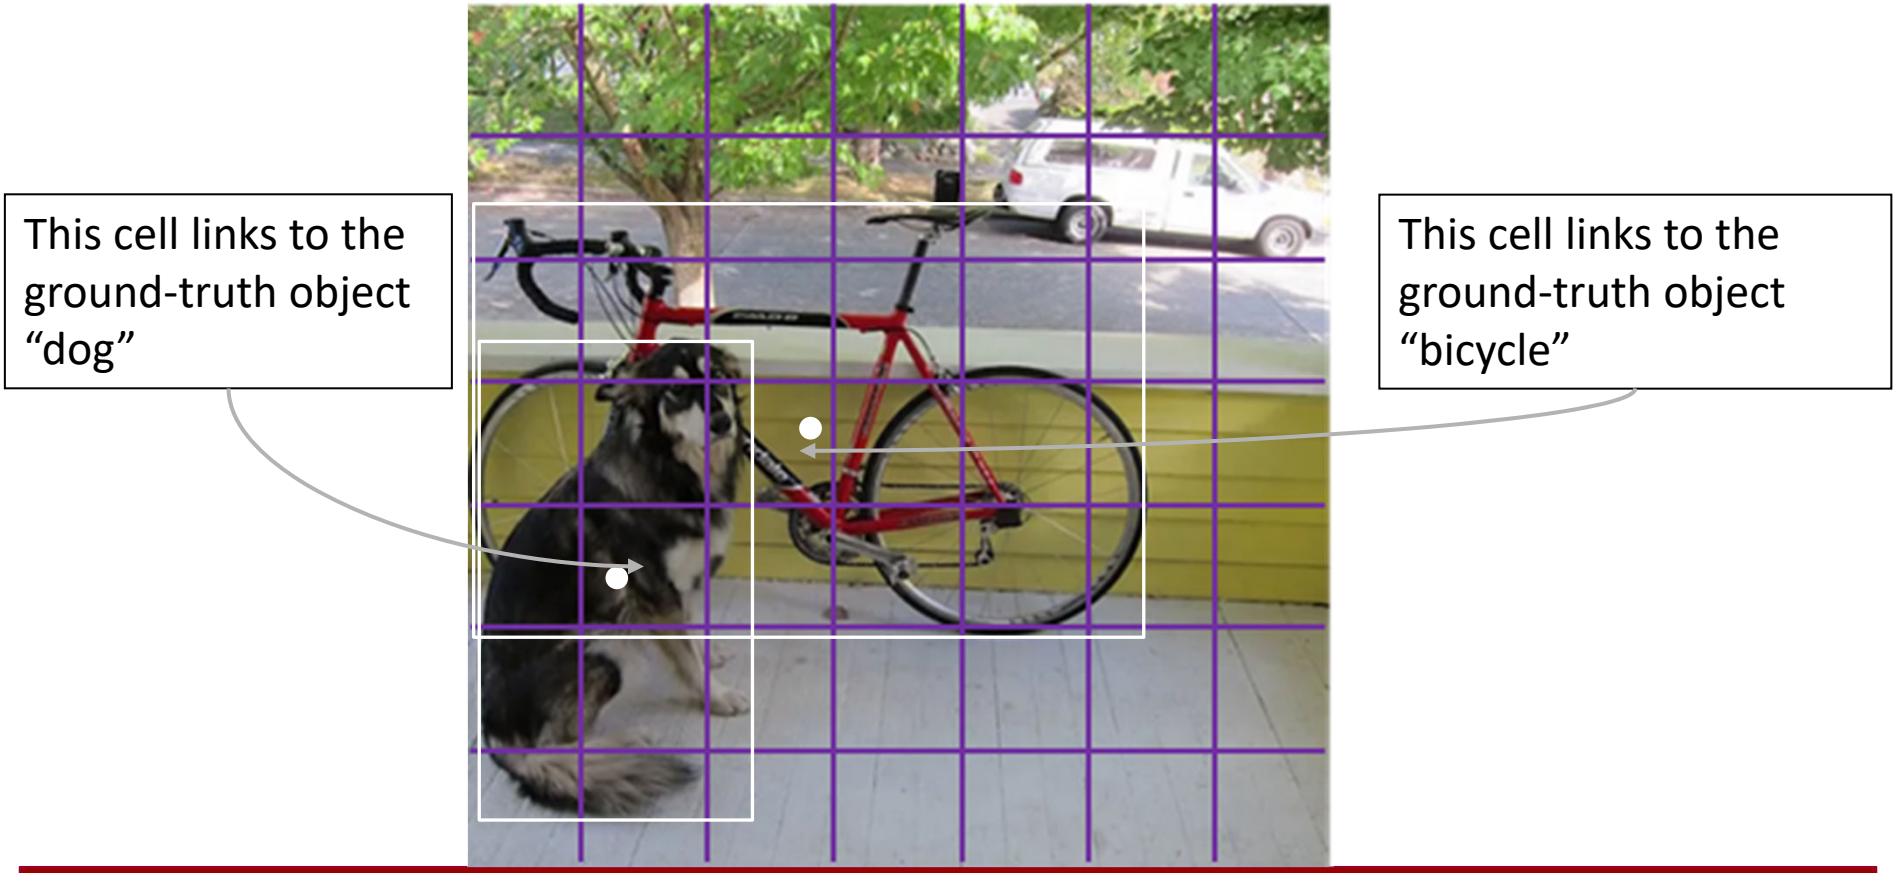

YOLOv1

For training

- Conceptually, the input image is divided in to 7×7 cells
- For cell i, if one ground-truth object's center falls in i, we say cell i links to a ground-truth object and is responsible for the loss computation related to that ground-truth object
- For a cell linking to a ground-truth object, though each time two bounding-boxes are predicted, only the one with the larger IoU with the ground-truth will participate in the loss calculation raised by this ground-truth object
- The loss is composed of three kinds of terms, related to position, objectiveness, and classification

cell i links to a ground-truth object and box j is responsible for predicting it

position

$$\lambda_{coord} \sum_{i=1}^{49} \sum_{j=1}^{2} \mathbf{1}_{ij}^{obj} \left[\left(x_i - \hat{x}_i \right)^2 + \left(y_i - \hat{y}_i \right)^2 + \left(\sqrt{w_i} - \sqrt{\hat{w}_i} \right)^2 + \left(\sqrt{h_i} - \sqrt{\hat{h}_i} \right) \right]$$

objectiveness

$$+ \sum_{i=1}^{49} \sum_{j=1}^{2} \mathbf{1}_{ij}^{obj} \left(s_i - \hat{s}_i \right)^2$$

The ground-truth value for the objectiveness is the IoU of this predicted bb with the ground-truth

$$+ \lambda_{noobj} \sum_{i=1}^{49} \sum_{j=1}^{2} \mathbf{1}_{ij}^{noobj} \left(0 - \hat{s}_i^j \right)^2$$

box j of cell i is not responsible for predicting any ground-truth

classification

$$+ \sum_{i=1}^{49} \left(\mathbf{1}_i^{obj} \sum_{c \in classes} \left(p_i(c) - \hat{p}_i(c) \right)^2 \right)$$

cell i links to a ground-truth

The values with ^ are predicted

YOLOv3

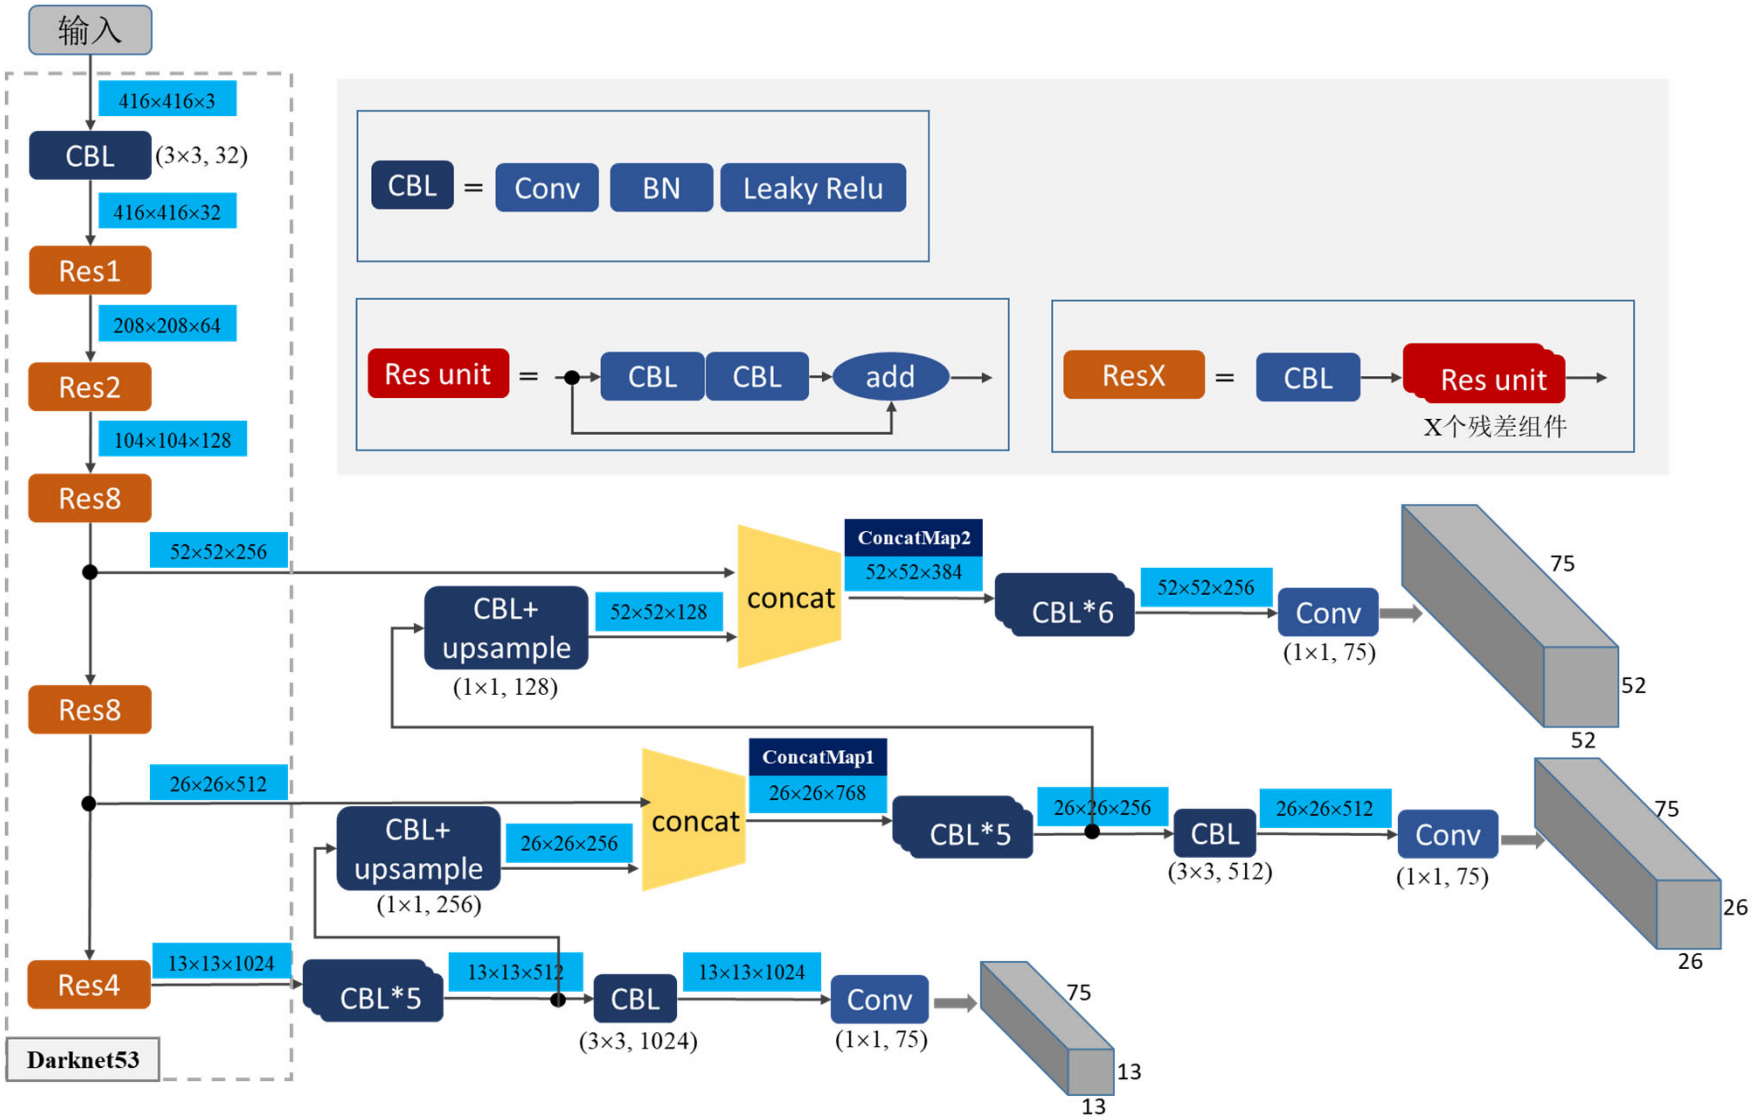
输入
416×416×3
CBL (3×3, 32)
416×416×32
Res1
208×208×64
Res2
104×104×128
Res8
52×52×256
Res8
26×26×512
Res4
13×13×1024
Darknet53
CBL = Conv BN Leaky Relu
Res unit = CBL CBL add
ResX = CBL Res unit
X个残差组件
CBL+ upsample
(1×1, 128)
52×52×128
concat
ConcatMap2
52×52×384
CBL*6
52×52×256
Conv
(1×1, 75)
75
52
52
CBL+ upsample
(1×1, 256)
26×26×256
concat
ConcatMap1
26×26×768
CBL*5
26×26×256
CBL
(3×3, 512)
26×26×512
Conv
(1×1, 75)
75
26
26
CBL*5
13×13×512
CBL
(3×3, 1024)
13×13×1024
Conv
(1×1, 75)
75
13
13

YOLOv3

- It incorporates a multi-scale detection scheme that can better detect objects with various sizes
  - It outputs three matrices with different resolutions, 13×13, 26×26 and 52×52
  - Conceptually, the input image is divided into 13×13, 26×26, and 52×52 cells

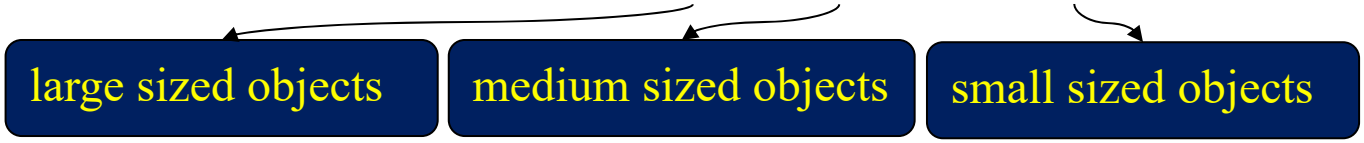

YOLOv3

- It incorporates a multi-scale detection scheme that can better detect objects with various sizes
- It uses the anchor-box mechanism to more accurately predict objects with various aspect ratios
  - An anchor-box is a rectangle with pre-defined size
  - Using anchor-boxes, the predicted box's size related values are relative to the corresponding anchor-box
  - In YOLOv3, three anchor-boxes are used for each cell
  - For the three different output resolutions, three different groups of anchor-boxes are used

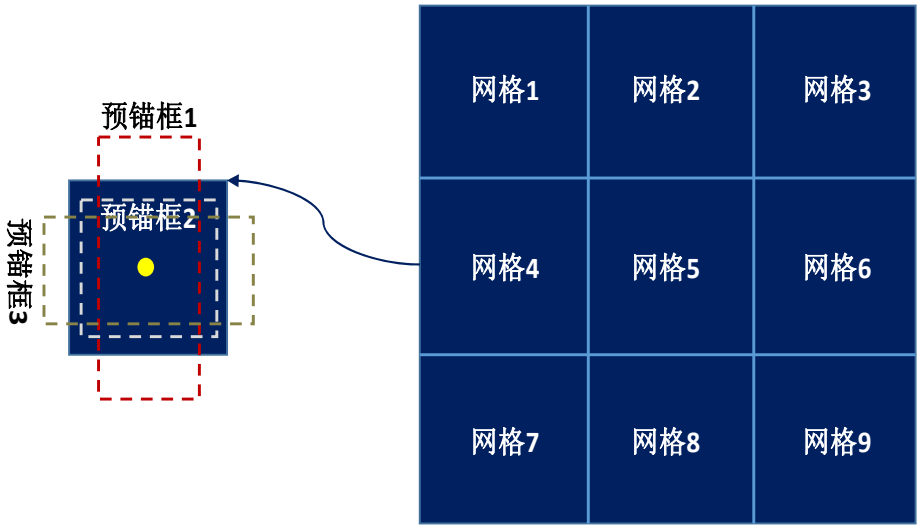

output resolution	Anchor-boxes
13×13	(116, 90), (156, 198), (373, 326)
26×26	(30, 61), (62, 45), (59, 119)
52×52	(10, 13), (16, 30), (33, 23)

YOLOv3

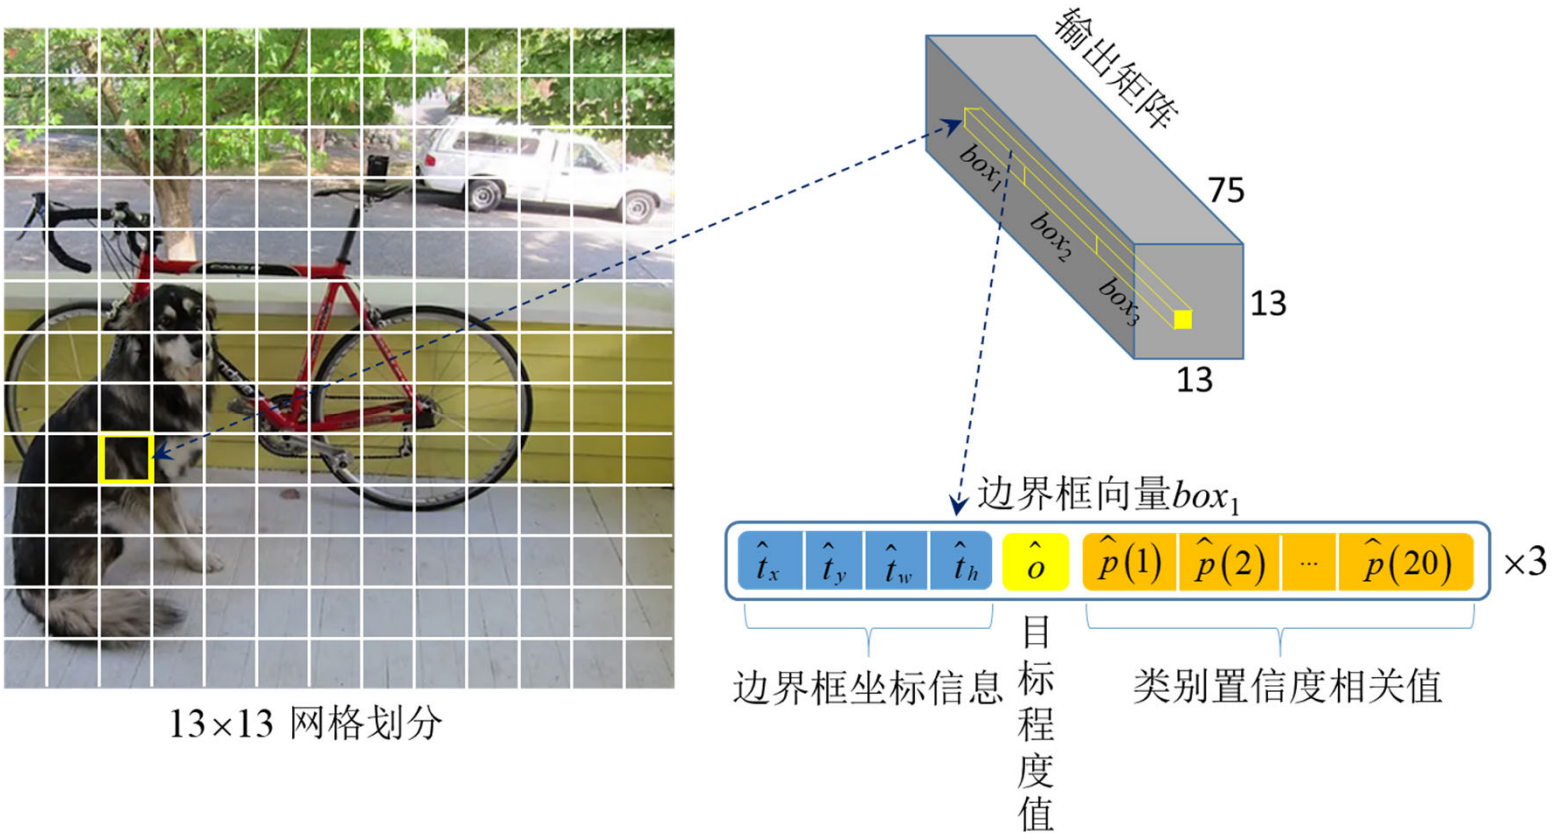

The detection results are derived from the three output matrices with different resolutions

YOLOv3

For inference

- The results are derived from the three output matrices
- Take the feature map with the resolution 13×13 as an example
  - ✓ Each cell predicts a $75d$ vector, which is actually composed of three predicted bounding-boxes, each for one anchor-box
  - ✓ For cell i, suppose one of its predicted boxes is

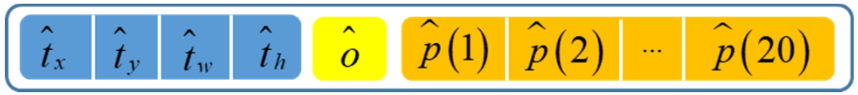

This bounding-box's center is $\begin{cases} b_x = \sigma(\hat{t}_x) + c_x \\ b_y = \sigma(\hat{t}_y) + c_y \end{cases}$, where (c_x, c_y) is cell i's position (values are 0~12)

This bounding-box's size is $\begin{cases} b_w = p_w e^{\hat{t}_w} \\ b_h = p_h e^{\hat{t}_h} \end{cases}$, where (p_w, p_h) is the associated anchor-box's size

This bounding-box's class-dependent confidence score is $\sigma(\hat{o}) \cdot \sigma(\hat{p}(j))$

- Finally, for each class, perform NMS for the bounding-boxes based on their class-dependent confidence scores

YOLOv3

For training

- For the ground-truth object $\mathbf{b}$, it is assigned to an anchor-box according to the IoUs between $\mathbf{b}$ and all the anchor-boxes
- For every anchor-box, it falls in one of the following three cases,
  - It is linked to a ground-truth object; for such an anchor-box, its predicted bounding box will participate in the loss calculation related to position, classification, and objectness
  - It is not linked to any ground-truth object, but the IoU between it and some ground-truth object is over 0.5, then this anchor-box will be “ignored”; such an anchor-box will not participate in any loss term calculation (since in this case, the supervisory information provided by this anchor-box is ambiguous)
  - Others (neither linked to any ground-truth object nor be ignored); for such an anchor-box, its predicted bounding-box only participates in the loss term calculation related to objectness (for this case, the ground-truth for objectness is 0)
- The final loss is a weighted sum of the loss terms related to position, classification, and objectness

YOLOv8

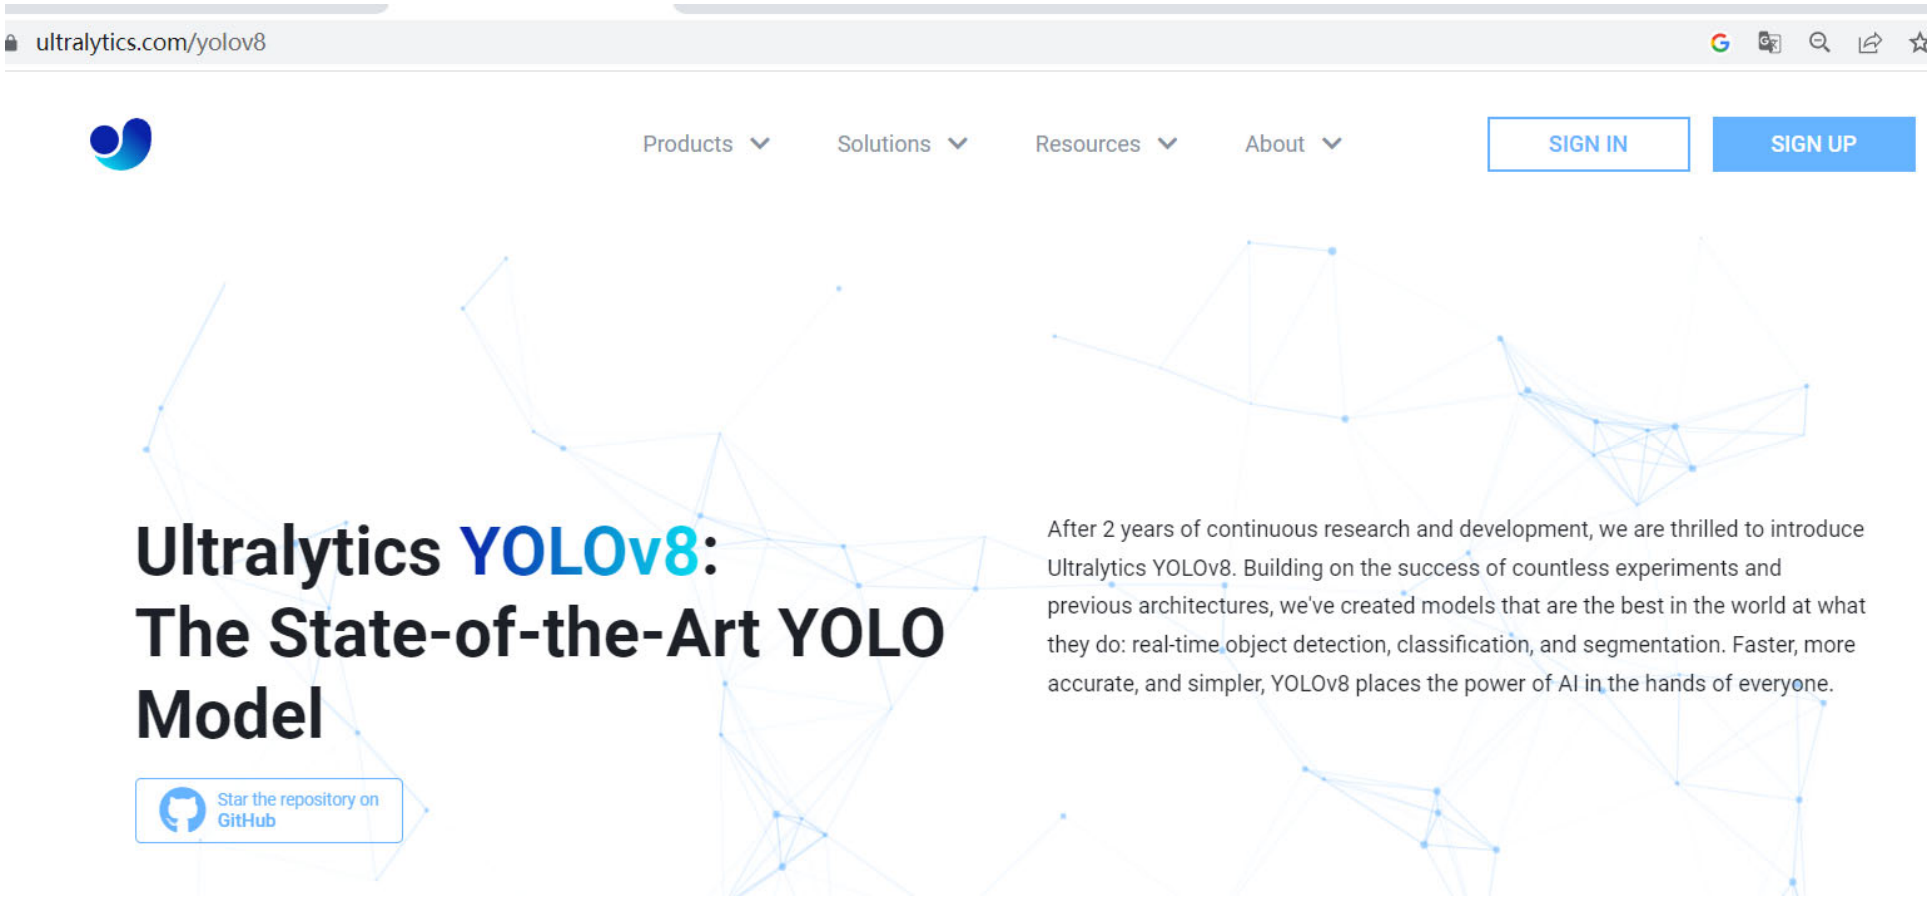
ultralytics.com/yolov8
Products
Solutions
Resources
About
SIGN IN
SIGN UP
Ultralytics YOLOv8:
The State-of-the-Art YOLO Model
Star the repository on GitHub
After 2 years of continuous research and development, we are thrilled to introduce Ultralytics YOLOv8. Building on the success of countless experiments and previous architectures, we've created models that are the best in the world at what they do: real-time object detection, classification, and segmentation. Faster, more accurate, and simpler, YOLOv8 places the power of AI in the hands of everyone.

YOLOv8

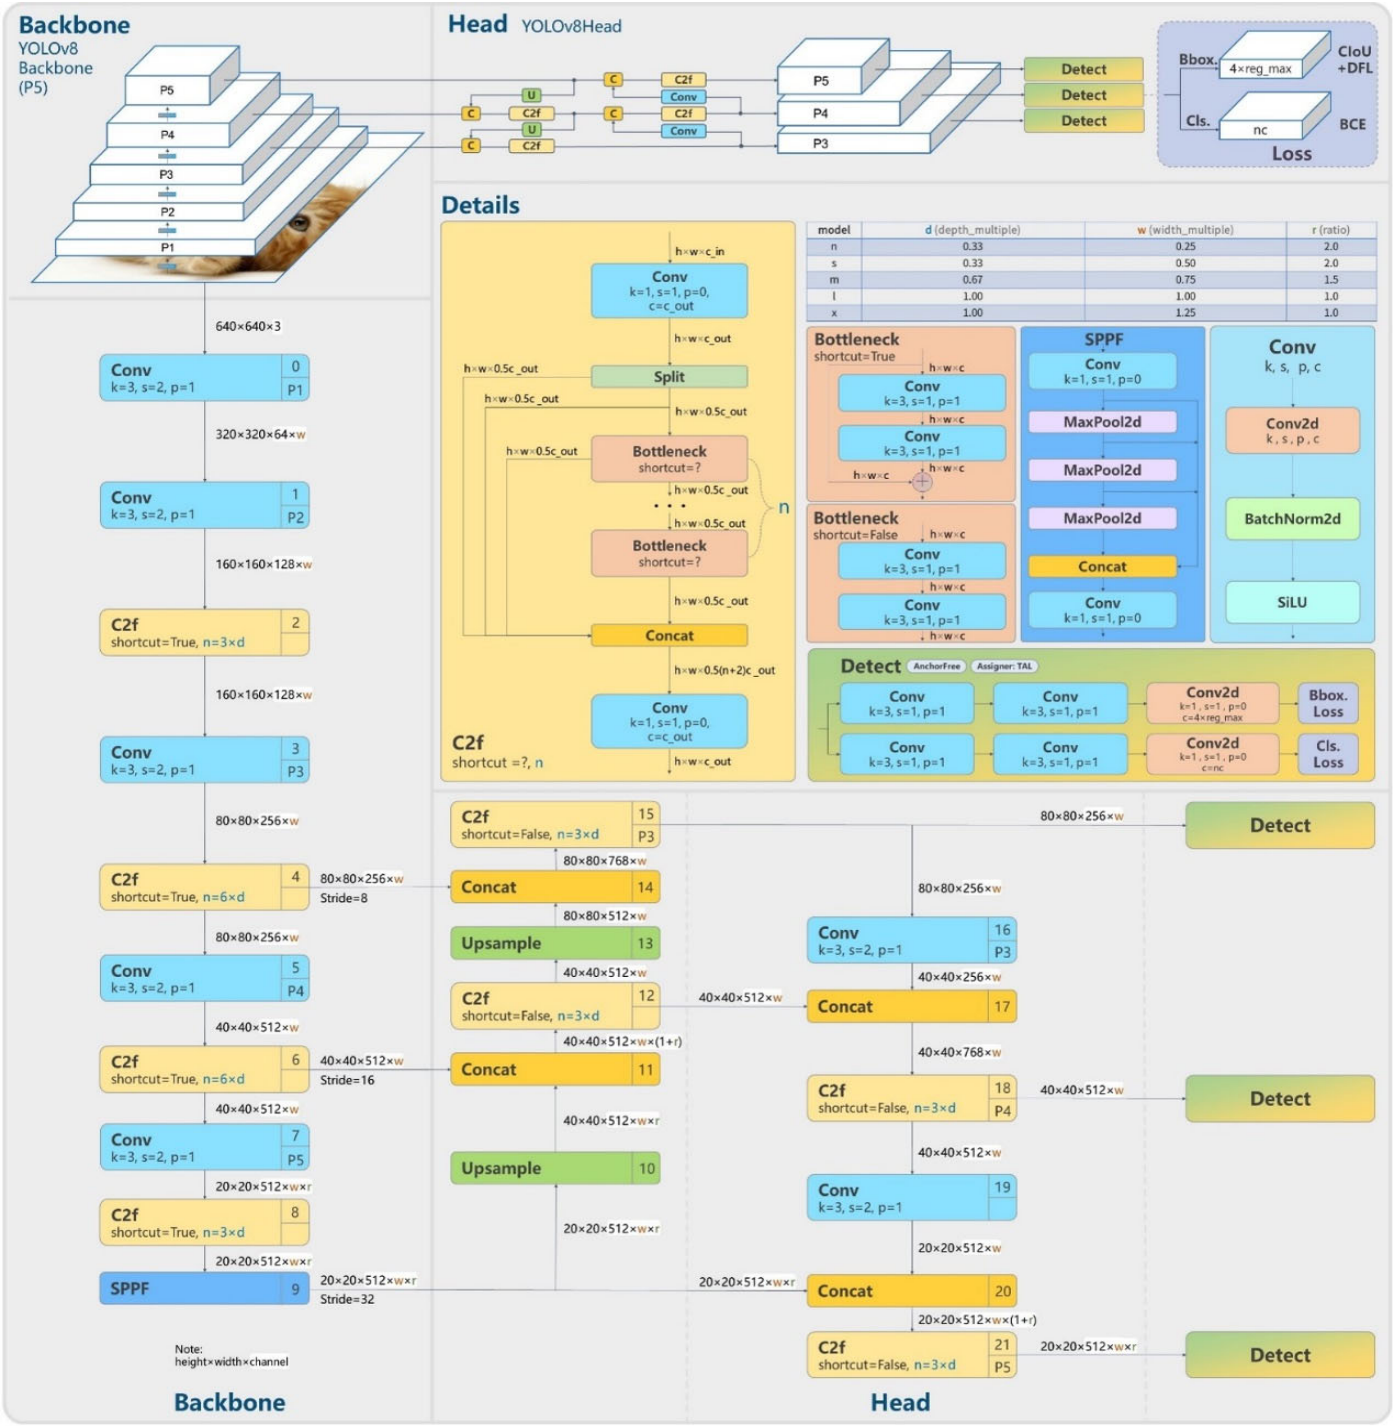
Backbone
YOLOv8
Backbone
(P5)
P5
P4
P3
P2
P1
Head YOLOv8Head
U
C
C2f
Conv
P5
P4
P3
Detect
Detect
Detect
Bbox.
4×reg_max
CIoU
+DFL
Cls.
nc
BCE
Loss
Details
h×w×c_in
Conv
k=1, s=1, p=0,
c=c_out
h×w×c_out
Split
h×w×0.5c_out
Bottleneck
shortcut=?
n
Concat
h×w×0.5(n+2)c_out
C2f
shortcut =?, n
model
d (depth_multiple)
w (width_multiple)
r (ratio)
n
0.33
0.25
2.0
s
0.33
0.50
2.0
m
0.67
0.75
1.5
l
1.00
1.00
1.0
x
1.00
1.25
1.0
Bottleneck
shortcut=True
Conv
k=3, s=1, p=1
Bottleneck
shortcut=False
SPPF
Conv
k=1, s=1, p=0
MaxPool2d
MaxPool2d
MaxPool2d
Concat
Conv
k=1, s=1, p=0
Conv
k, s, p, c
Conv2d
k , s , p , c
BatchNorm2d
SiLU
Detect
AnchorFree
Assigner: TAL
Conv
k=3, s=1, p=1
Conv2d
k=1 , s=1 , p=0
c=4×reg_max
Bbox.
Loss
Cls.
Loss
640×640×3
Conv
k=3, s=2, p=1
0
P1
320×320×64×w
1
P2
160×160×128×w
C2f
shortcut=True, n=3×d
2
160×160×128×w
3
P3
80×80×256×w
4
80×80×256×w
Stride=8
5
P4
40×40×512×w
6
40×40×512×w
Stride=16
7
P5
20×20×512×w×r
8
20×20×512×w×r
SPPF
9
20×20×512×w×r
Stride=32
Note:
height×width×channel
Backbone
C2f
shortcut=False, n=3×d
15
P3
80×80×768×w
Concat
14
80×80×512×w
Upsample
13
40×40×512×w
12
40×40×512×w×(1+r)
11
40×40×512×w×r
Upsample
10
20×20×512×w×r
80×80×256×w
Detect
80×80×256×w
16
P3
40×40×256×w
40×40×512×w
Concat
17
40×40×768×w
18
P4
40×40×512×w
Detect
40×40×512×w
19
20×20×512×w
20×20×512×w×r
Concat
20
20×20×512×w×(1+r)
21
P5
20×20×512×w×r
Detect
Head


person (93%)
person (99%)
person (100%)
speedbump (98%)
speedbump (94%)